

Crash Utility Help

*	files	mod	sbitmapq	tree
alias	foreach	mount	search	union
ascii	fuser	net	set	vm
bpf	gdb	p	sig	vtop
bt	help	ps	struct	waitq
bttop	ipcs	pte	swap	whatis
dev	irq	ptob	sym	wr
dis	kmem	ptov	sys	q
eval	list	rd	task	
exit	log	repeat	timer	
extend	mach	runq	trace	

crash version: 8.0.1++ gdb version: 10.2
For help on any command above, enter "help <command>".
For help on input options, enter "help input".
For help on output options, enter "help output".

```

74 NAME
75     * - pointer-to short-cut
76
77 SYNOPSIS
78     * (struct or union command arguments)
79
80 DESCRIPTION
81     This command is a short-cut command that replaces the requirement to enter
82     "struct" or "union" command names. For details on the arguments to
83     those commands, enter "help struct" or "help union".
84
85 EXAMPLES
86     Dump the page structure at address c02943c0:
87
88     crash> *page c02943c0
89     struct page {
90         next = 0xc0fae740,
91         prev = 0xc0018fb0,
92         inode = 0x0,
93         offset = 0x3f000,
94         next_hash = 0xc02d6310,
95         count = {
96             counter = 0x1
97         },
98         flags = 0x310,
99         wait = 0xc02943d8,
100         pprev_hash = 0x0,
101         buffers = 0x0
102     }
103
104 NAME
105     files - open files
106
107 SYNOPSIS
108     files [-d dentry] | [-p inode] | [-c] [-R reference] [pid | taskp] ...
109
110 DESCRIPTION
111     This command displays information about open files of a context.
112     It prints the context's current root directory and current working
113     directory, and then for each open file descriptor it prints a pointer
114     to its file struct, a pointer to its dentry struct, a pointer to the
115     inode, the file type, and the pathname. If no arguments are entered,
116     the current context is used. The -R option, typically invoked from
117     "foreach files", searches for references to a supplied number, address,
118     or filename argument, and prints only the essential information leading
119     up to and including the reference. The -d option is not context
120     specific, and only shows the data requested.
121
122     -d dentry    given a hexadecimal dentry address, display its inode,
123                  super block, file type, and full pathname.
124     -p inode     given a hexadecimal inode address, dump all of its pages
125                  that are in the page cache.
126     -c           for each open file descriptor, prints a pointer to its
127                  inode, a pointer to the inode's i_mapping address_space
128                  structure, the number of pages of the inode that are in
129                  the page cache, the file type, and the pathname.
130     -R reference search for references to this file descriptor number,
131                  filename, dentry, inode, address_space, or file structure
132                  address.
133     pid         a process PID.
134     taskp       a hexadecimal task_struct pointer.
135
136 EXAMPLES
137     Display the open files of the current context:
138
139     crash> files
140     PID: 720    TASK: c67f2000 CPU: 1    COMMAND: "innd"
141     ROOT: /    CWD: /var/spool/news/articles
142
143     FD   FILE      DENTRY      INODE      TYPE  PATH
144     0    c6b9c740    c7cc45a0    c7c939e0    CHR   /dev/null
145     1    c6b9c800    c537bb20    c54d0000    REG   /var/log/news/news
146     2    c6df9600    c537b420    c5c36360    REG   /var/log/news/errlog

```

```

147      3 c74182c0 c6ede260 c6da3d40 PIPE
148      4 c6df9720 c696c620 c69398c0 SOCK
149      5 c6b9cc20 c68e7000 c6938d80 SOCK
150      6 c6b9c920 c7cc45a0 c7c939e0 CHR /dev/null
151      7 c6b9c680 c58fa5c0 c58a1200 REG /var/lib/news/history
152      8 c6df9f00 c6ede760 c6da3200 PIPE
153      9 c6b9c6e0 c58fa140 c5929560 REG /var/lib/news/history.dir
154     10 c7fa9320 c7fab160 c7fafd40 CHR /dev/console
155     11 c6b9c7a0 c58fa5c0 c58a1200 REG /var/lib/news/history
156     12 c377ec60 c58fa5c0 c58a1200 REG /var/lib/news/history
157     13 c4528aa0 c58fa6c0 c52fbb00 REG /var/lib/news/history.pag
158     14 c6df9420 c68e7700 c6938360 SOCK
159     15 c6df9360 c68e7780 c6938120 SOCK
160     16 c6b9c0e0 c68e7800 c6772000 SOCK
161     17 c6b9c200 c6b5f9c0 c6b5cea0 REG /var/lib/news/active
162     21 c6b9c080 c6ede760 c6da3200 PIPE

```

Display the files opened by the "crond" daemon, which is PID 462:

```

166 crash> files 462
167 PID: 462 TASK: f7220000 CPU: 2 COMMAND: "crond"
168 ROOT: / CWD: /var/spool
169 FD FILE DENTRY INODE TYPE PATH
170 0 f7534ae0 f7538de0 f7518dc0 CHR /dev/console
171 1 f7368f80 f72c7a40 f72f27e0 FIFO pipe:[1456]
172 2 f74f3c80 f72c79c0 f72f2600 FIFO pipe:[1457]
173 3 f7368b60 f72a5be0 f74300c0 REG /var/run/crond.pid
174 4 f7534360 f73408c0 f72c2840 REG /var/log/cron
175 7 f7368ce0 f72c7940 f72f2420 FIFO pipe:[1458]
176 8 f7295de0 f72c7940 f72f2420 FIFO pipe:[1458]
177 21 f74f36e0 f747cdc0 f747e840 CHR /dev/null

```

The -R option is typically invoked from "foreach files". This example shows all tasks that have "/dev/pts/4" open:

```

182 crash> foreach files -R pts/4
183 PID: 18633 TASK: c310a000 CPU: 0 COMMAND: "crash"
184 ROOT: / CWD: /home/CVS_pool/crash
185 FD FILE DENTRY INODE TYPE PATH
186 0 c1412850 c2cb96d0 c2cad430 CHR /dev/pts/4
187 1 c1412850 c2cb96d0 c2cad430 CHR /dev/pts/4
188 2 c1412850 c2cb96d0 c2cad430 CHR /dev/pts/4
189
190 PID: 18664 TASK: c2392000 CPU: 1 COMMAND: "less"
191 ROOT: / CWD: /home/CVS_pool/crash
192 FD FILE DENTRY INODE TYPE PATH
193 1 c1412850 c2cb96d0 c2cad430 CHR /dev/pts/4
194 2 c1412850 c2cb96d0 c2cad430 CHR /dev/pts/4
195
196 PID: 23162 TASK: c5088000 CPU: 1 COMMAND: "bash"
197 ROOT: / CWD: /home/CVS_pool/crash
198 FD FILE DENTRY INODE TYPE PATH
199 0 c1412850 c2cb96d0 c2cad430 CHR /dev/pts/4
200 1 c1412850 c2cb96d0 c2cad430 CHR /dev/pts/4
201 2 c1412850 c2cb96d0 c2cad430 CHR /dev/pts/4
202 255 c1412850 c2cb96d0 c2cad430 CHR /dev/pts/4
203
204 PID: 23159 TASK: c10fc000 CPU: 1 COMMAND: "xterm"
205 ROOT: / CWD: /homes/anderson/
206 FD FILE DENTRY INODE TYPE PATH
207 5 c1560da0 c2cb96d0 c2cad430 CHR /dev/pts/4

```

Display information about the dentry at address f745fd60:

```

211 crash> files -d f745fd60
212 DENTRY INODE SUPERBLK TYPE PATH
213 f745fd60 f7284640 f73a3e00 REG /var/spool/lpd/lpd.lock

```

For each open file, display the number of pages that are in the page cache:

```

217 crash> files -c 1954
218 PID: 1954 TASK: f7a28000 CPU: 1 COMMAND: "syslogd"
219 ROOT: / CWD: /

```

FD	INODE	I_MAPPING	NRPAGES	TYPE	PATH
0	cb3ae868	cb3ae910	0	SOCK	socket:[4690]
2	f2721c5c	f2721d04	461	REG	/var/log/messages
3	cbda4884	cbda492c	47	REG	/var/log/secure
4	e48092c0	e4809368	58	REG	/var/log/maillog
5	f65192c0	f6519368	48	REG	/var/log/cron
6	e4809e48	e4809ef0	0	REG	/var/log/spooler
7	d9c43884	d9c4392c	0	REG	/var/log/boot.log

For the inode at address f59b90fc, display all of its pages that are in the page cache:

```
crash> files -p f59b90fc
```

INODE	NRPAGES
f59b90fc	6

PAGE	PHYSICAL	MAPPING	INDEX	CNT	FLAGS
ca3353e0	39a9f000	f59b91ac	0	2	82c referenced,uptodate,lru,private
ca22cb20	31659000	f59b91ac	1	2	82c referenced,uptodate,lru,private
ca220160	3100b000	f59b91ac	2	2	82c referenced,uptodate,lru,private
calddde0	2eeef000	f59b91ac	3	2	82c referenced,uptodate,lru,private
ca36b300	3b598000	f59b91ac	4	2	82c referenced,uptodate,lru,private
ca202680	30134000	f59b91ac	5	2	82c referenced,uptodate,lru,private

NAME

mod - module information and loading of symbols and debugging data

SYNOPSIS

```
mod -s module [objfile] | -d module | -S [directory] [-D|-t|-r|-R|-o|-g]
```

DESCRIPTION

With no arguments, this command displays basic information of the currently installed modules, consisting of the module address, name, base address, size, the object file name (if known), and whether the module was compiled with CONFIG_KALLSYMS.

The arguments are concerned with with the loading or deleting of symbolic and debugging data from a module's object file. A modules's object file always contains symbolic data (symbol names and addresses), but contains debugging data only if the module was compiled with the -g CFLAG. In addition, the module may have compiled with CONFIG_KALLSYMS, which means that the module's symbolic data will have been loaded into the kernel's address space when it was installed. If the module was not compiled with CONFIG_KALLSYMS, then only the module's exported symbols will be loaded into the kernel's address space. Therefore, for the purpose of this command, it should noted that a kernel module may have been compiled in one of following manners:

1. If the module was built without CONFIG_KALLSYMS and without the -g CFLAG, then the loading of the module's additional non-exported symbols can be accomplished with this command.
2. If the module was built with CONFIG_KALLSYMS, but without the -g CFLAG, then there is no benefit in loading the symbols from the module object file, because all of the module's symbols will have been loaded into the kernel's address space when it was installed.
3. If the module was built with CONFIG_KALLSYMS and with the the -g CFLAG, then the loading of the module's debugging data can be accomplished with this command.
4. If the module was built without CONFIG_KALLSYMS but with the -g CFLAG, then the loading of the both module's symbolic and debugging data can be accomplished with this command.

-s module [objfile] Loads symbolic and debugging data from the object file for the module specified. If no objfile argument is appended, a search will be made for an object file consisting of the module name with a .o or .ko suffix, starting at the /lib/modules/<release> directory on the host system, or if not found there, starting at the directory containing the kernel namelist file. If an objfile argument is appended, then that file will be used.

```

293         -d module    Deletes the symbolic and debugging data of the module
294                       specified.
295     -S [directory]   Load symbolic and debugging data from the object file
296                       for all loaded modules. For each module, a search
297                       will be made for an object file consisting of the
298                       module name with a .o or .ko suffix, starting at the
299                       /lib/modules/<release> directory of the host system,
300                       or if not found there, starting at the directory
301                       containing the kernel namelist file. If a directory
302                       argument is appended, then the search will be restricted
303                       to that directory.
304     -D               Deletes the symbolic and debugging data of all modules.
305     -t               Display the contents of the module's "taints" bitmask
306                       if it is non-zero. When possible, the "taints" bits
307                       are translated to symbolic letters of the taint type;
308                       otherwise the hexadecimal value is shown. In older
309                       kernels, the contents of the "license_gplok" field is
310                       displayed in hexadecimal; the field may be either a
311                       bitmask or a boolean, depending upon the kernel version.
312                       The relevant kernel sources should be consulted for the
313                       meaning of the letter(s) or hexadecimal bit value(s).
314                       For modules that have a "gpgsig_ok" field that is zero
315                       (unsigned), the notation "(U)" is shown.
316     -r               Passes the -readnow flag to the embedded gdb module,
317                       which will override the two-stage strategy that it uses
318                       for reading symbol tables from module object files.
319     -R               Reinitialize module data. All currently-loaded symbolic
320                       and debugging data will be deleted, and the installed
321                       module list will be updated (live system only).
322     -g               When used with -s or -S, add a module object's section
323                       start and end addresses to its symbol list.
324     -o               Load module symbols with old mechanism.
325

```

If the crash session was invoked with the "--mod <directory>" option, or a CRASH_MODULE_PATH environment variable exists, then /lib/modules/<release> will be overridden as the default directory tree that is searched for module object files.

After symbolic and debugging data have been loaded, backtraces and text disassembly will be displayed appropriately. Depending upon the processor architecture, data may also printed symbolically with the "p" command; at a minimum, the "rd" command may be used with module data symbols.

If crash can recognize that the set of modules has changed while running a session on a live kernel, the module data will be reinitialized the next time this command is run; the -r option forces the reinitialization.

EXAMPLES

Display the currently-installed modules:

```

341 crash> mod
342
343 MODULE      NAME                BASE      SIZE  OBJECT FILE
344 f7e44c20    dm_mod                  f7e34000  88568 (not loaded)
345 f7e5a8a0    dm_log                    f7e59000   8354 (not loaded)
346 f7e66420    dm_region_hash            f7e65000   9708 (not loaded)
347 f7e76b60    dm_mirror                 f7e74000  12609 (not loaded)
348 f7e8b8e0    ata_piix                  f7e87000  20637 (not loaded)
349 ...
350

```

Display the currently-installed modules on a system where all modules were compiled with CONFIG_KALLSYMS:

```

351 crash> mod
352
353 MODULE      NAME                BASE      SIZE  OBJECT FILE
354 f7e44c20    dm_mod                  f7e34000  88568 (not loaded) [CONFIG_KALLSYMS]
355 f7e5a8a0    dm_log                    f7e59000   8354 (not loaded) [CONFIG_KALLSYMS]
356 f7e66420    dm_region_hash            f7e65000   9708 (not loaded) [CONFIG_KALLSYMS]
357 f7e76b60    dm_mirror                 f7e74000  12609 (not loaded) [CONFIG_KALLSYMS]
358 f7e8b8e0    ata_piix                  f7e87000  20637 (not loaded) [CONFIG_KALLSYMS]
359 ...
360

```

Load the symbolic and debugging data of all modules:

```

366 crash> mod -S
367 MODULE      NAME      BASE      SIZE  OBJECT FILE
368 f7e44c20 dm_mod      f7e34000  88568
    /lib/modules/2.6.32/kernel/drivers/md/dm-mod.ko
369 f7e5a8a0 dm_log      f7e59000   8354
    /lib/modules/2.6.32/kernel/drivers/md/dm-log.ko
370 f7e66420 dm_region_hash f7e65000   9708
    /lib/modules/2.6.32/kernel/drivers/md/dm-region-hash.ko
371 f7e76b60 dm_mirror    f7e74000  12609
    /lib/modules/2.6.32/kernel/drivers/md/dm-mirror.ko
372 f7e8b8e0 ata_piix    f7e87000  20637
    /lib/modules/2.6.32/kernel/drivers/ata/ata_piix.ko
373 ...
374
375 Load the symbolic and debugging data of the dm_mod module from its
376 known location:
377
378 crash> mod -s dm_mod
379 MODULE      NAME      BASE      SIZE  OBJECT FILE
380 f7e44c20 dm_mod      f7e34000  88568
    /lib/modules/2.6.32/kernel/drivers/md/dm-mod.ko
381
382 Delete the current symbolic and debugging data of the dm_mod module,
383 and then re-load it from a specified object file:
384
385 crash> mod -d dm_mod
386 crash> mod -s dm_mod /tmp/dm_mod.ko
387 MODULE      NAME      BASE      SIZE  OBJECT FILE
388 f7e44c20 dm_mod      f7e34000  88568  /tmp/dm-mod.ko
389
390 After installing a new kernel module on a live system, reinitialize the
391 installed module list:
392
393 crash> !modprobe soundcore
394 crash> mod
395 mod: NOTE: modules have changed on this system -- reinitializing
396 MODULE      NAME      BASE      SIZE  OBJECT FILE
397 f7e44c20 dm_mod      f7e34000  88568  (not loaded)
398 f7e5a8a0 dm_log      f7e59000   8354  (not loaded)
399 f7e62e40 soundcore    f7e62000   6390  (not loaded)
400 f7e66420 dm_region_hash f7e65000   9708  (not loaded)
401 f7e76b60 dm_mirror    f7e74000  12609  (not loaded)
402 f7e8b8e0 ata_piix    f7e87000  20637  (not loaded)
403 ...
404
405 Display modules that are "tainted", where in this case
406 where they are proprietary and unsigned:
407
408 crash> mod -t
409 NAME      TAINT
410 vxspec    P(U)
411 vxportal  P(U)
412 fdd       P(U)
413 vxfs      P(U)
414 vxdmp     P(U)
415 vxio      P(U)
416 vxglm     P(U)
417 vxgms     P(U)
418 vxodm     P(U)
419
420
421 NAME
422 sbitmapq - sbitmap_queue struct contents
423
424 SYNOPSIS
425 sbitmapq [-s struct[.member[,member]] -a address [-p] [-v]] [-x|d] address
426
427 DESCRIPTION
428 The command dumps the contents of the sbitmap_queue structure and
429 the used bits in the bitmap. Also, it shows the dump of a structure
430 array associated with the sbitmap_queue.
431
432 The arguments are as follows:

```

```

433
434 -s struct name of a C-code structure, that is stored in an array
435 associated with sbitmap_queue structure. Use the
436 "struct.member" format in order to display a particular
437 member of the structure. -s option requires -a option
438 -a address address of a structure array associated with sbitmap_queue
439 structure. The set bits in sbitmap are used for the index
440 in an associated array.
441 -p associated with sbitmap_queue array contains the points of
442 structure.
443 -x override default output format with hexadecimal format.
444 -d override default output format with decimal format.
445 -v By default, the sbitmap command shows only a used sbitmap
446 index and a structure address in the associated array.
447 This flag says to print a formatted display of the
448 contents of a structure in an associated array. -v option
449 requires of -s.
450

```

451 EXAMPLES

452 All examples are shown on the base of Linux Target system with iSCSI
453 transport.

454 Display the common sbitmap information for target session:

```

455
456 crash> struct -oh se_session 0xc0000000e118c760 | grep sbitmap_queue
457 [c0000000e118c808] struct sbitmap_queue sess_tag_pool;
458 crash>
459 crash> sbitmapq c0000000e118c808
460 depth = 136
461 busy = 4
462 cleared = 26
463 bits_per_word = 32
464 map_nr = 5
465 alloc_hint = {74, 36, 123, 101}
466 wake_batch = 8
467 wake_index = 0
468 ws_active = 0
469 ws = {
470     { .wait_cnt = 8, .wait = inactive },
471     { .wait_cnt = 8, .wait = inactive },
472     { .wait_cnt = 8, .wait = inactive },
473     { .wait_cnt = 8, .wait = inactive },
474     { .wait_cnt = 8, .wait = inactive },
475     { .wait_cnt = 8, .wait = inactive },
476     { .wait_cnt = 8, .wait = inactive },
477     { .wait_cnt = 8, .wait = inactive },
478     { .wait_cnt = 8, .wait = inactive },
479     { .wait_cnt = 8, .wait = inactive },
480 }
481 round_robin = 0
482 min_shallow_depth = 4294967295
483
484 00000000: 0000 0000 0000 0000 0030 0000 0000 0000
485 00000010: 00
486

```

487 Display the addresses of structure are associated with
488 sbitmap_queue (for iscsi it is 'iscsi_cmd' structure):

```

489
490 crash> struct se_session 0xc0000000e118c760 | grep sess_cmd_map
491 sess_cmd_map = 0xc0000000671c0000,
492 crash>
493 crash> sbitmapq -s iscsi_cmd -a 0xc0000000671c0000 c0000000e118c808
494 76: 0xc0000000671d5600
495 77: 0xc0000000671d5a80
496

```

497 Dump of formatted content of structures:

```

498
499 crash> sbitmapq -s iscsi_cmd -a 0xc0000000671c0000 -v c0000000e118c808
500 76 (0xc0000000671d5600):
501 struct iscsi_cmd {
502     dataout_timer_flags = 0,
503     dataout_timeout_retries = 0 '\000',
504     error_recovery_count = 0 '\000',
505     deferred_i_state = ISTATE_NO_STATE,

```

```

506     i_state = ISTATE_SENT_STATUS,
507     ...
508     first_data_sg = 0xc0000000e306b080,
509     first_data_sg_off = 0,
510     kmapped_nents = 1,
511     sense_reason = 0
512 }
513 77 (0xc0000000671d5a80):
514 struct iscsi_cmd {
515     dataout_timer_flags = 0,
516     dataout_timeout_retries = 0 '\000',
517     error_recovery_count = 0 '\000',
518     deferred_i_state = ISTATE_NO_STATE,
519     i_state = ISTATE_NEW_CMD,
520     ...
521     first_data_sg = 0x0,
522     first_data_sg_off = 0,
523     kmapped_nents = 0,
524     sense_reason = 0
525 }

```

NAME

tree - display radix tree, XArray or red-black tree

SYNOPSIS

```

tree [-t [radix|xarray|rbtree]] [-r offset] [-[s|S] struct[.member[,member]]]
    [-[x|d] [-o offset] [-l] [-p] [-N] start

```

DESCRIPTION

This command dumps the contents of a radix tree, an XArray, or a red-black tree. The arguments are as follows:

- t type The type of tree to dump; the type string can be one of "radix", "rbtree", or "xarray", or alternatively, "ra", "rb" or "x" are acceptable. If not specified, rbtree is the default type.
- r offset If the "start" argument is the address of a data structure that contains an `radix_tree_root`, `xarray` or `rb_root` structure, then this is the offset to that structure member. If the offset is non-zero, then this option is required. The offset may be entered in either of two manners:
 1. In "structure.member" format.
 2. A number of bytes.
- o offset For red-black trees only, the offset of the `rb_node` within its containing data structure; if the offset is non-zero, then this option is required. The offset may be entered in either of two manners:
 1. In "structure.member" format.
 2. A number of bytes.
- s struct This option is not applicable to radix trees. For each entry in a tree, format and print it as this type of data structure; use the "struct.member" format in order to display a particular member of the structure. To display multiple members of a structure, use a comma-separated list of members. If any structure member contains an embedded structure or is an array, the the output may be restricted to the embedded structure or an array element by expressing the member argument as "struct.member.member" or "struct.member[index]"; embedded member specifications may extend beyond one level deep by expressing the struct argument as "struct.member.member.member...".
- S struct Similar to -s, but instead of parsing gdb output, member values are read directly from memory, so the command works much faster for 1-, 2-, 4-, and 8-byte members.
- l For red-black trees, dump the tree sorted in linear order starting with the leftmost node and progressing to the right. This option does not apply to radix trees.
- p Display the node's position information, showing the relationship between it and the root. For red-black trees, a position that indicates "root/l/r" means that the node is the right child of the left child of the root node. For radix trees and xarrays, the index, the height, and the slot index values are shown with respect to the root.


```

579         -x Override default output format with hexadecimal format.
580         -d Override default output format with decimal format.
581
582 The meaning of the "start" argument, which can be expressed either in
583 hexadecimal format or symbolically, depends upon whether the -N option
584 is prepended:

```

```

585     start  The address of a radix_tree_root, xarray or rb_root structure, or
586            the address of a structure containing the radix_tree_root, xarray
587            or rb_root structure; if the latter, then the "-r offset" option
588            must be used if the member offset of the root structure is
589            non-zero.
590
591     -N start  The address of a radix_tree_node, xa_node or rb_node structure,
592            bypassing the radix_tree_root, xarray, or rb_root that points
593            to it.
594
595
596

```

EXAMPLES

The `vmap_area_root` is a standalone `rb_root` structure. Display the virtual addresses of each `vmap_area` in its red-black tree:

```

600
601 crash> whatis vmap_area_root
602 struct rb_root vmap_area_root;
603 crash> tree -t rbtree -o vmap_area.rb_node vmap_area_root
604 ffff880128c508c0
605 ffff88012cb68140
606 ffff88012c9afec0
607 ffff88012d65c440
608 ...

```

Display the `vmap_area`'s `va_start` and `va_end` members of each of the entries above expressing the `vmap_area.rb_node` offset as a number of bytes:

```

613
614 crash> tree -t rbtree -o 24 vmap_area_root -s vmap_area.va_start,va_end
615 ffff880128c508c0
616     va_start = 0xfffffc90014900000
617     va_end   = 0xfffffc90014921000
618 ffff88012cb68140
619     va_start = 0xfffffc900110c0000
620     va_end   = 0xfffffc900110d1000
621 ffff88012c9afec0
622     va_start = 0xfffffc900000640000
623     va_end   = 0xfffffc900000642000
624 ffff88012d65c440
625     va_start = 0xfffffc900000620000
626     va_end   = 0xfffffc900000622000
627 ...

```

Alternatively, use the `-N` option with the `rb_node` address contained in the `vmap_area_root` structure:

```

631
632 crash> p vmap_area_root
633 vmap_area_root = $8 = {
634     rb_node = 0xfffff880128c508d8
635 }
636 crash> tree -t rbtree -o vmap_area.rb_node -N 0xfffff880128c508d8
637 ffff880128c508c0
638 ffff88012cb68140
639 ffff88012c9afec0
640 ffff88012d65c440

```

Display the virtual address of each `vm_area_struct` in the red-black tree that has its root inside an `mm_struct` located at `ffff880128b5a300`. The `vm_area_struct.vm_rb` `rb_node` member has an offset of `0x38` bytes:

```

641
642 crash> tree -t rbtree -r mm_struct.mm_rb ffff880128b5a300 -o 0x38
643 ffff88012a0de080
644 ffff880123e3ac78
645 ffff880123e3a700
646 ffff88012b2837c8
647 ...

```

```
652      ffff880128c02ed0
653      ffff8801292e7958
654      ffff880123e3a318
655      ffff880123e3ad40
656
```

657 Add the -p option to the command above to show position information:

```
658
659 crash> tree -t rbtree -r mm_struct.mm_rb ffff880128b5a300 -o 0x38 -p
660 ffff88012a0de080
661     position: root
662 ffff880123e3ac78
663     position: root/l
664 ffff880123e3a700
665     position: root/l/l
666 ffff88012b2837c8
667     position: root/l/l/l
668 ...
669 ffff880128c02ed0
670     position: root/r/r/l/r
671 ffff8801292e7958
672     position: root/r/r/l/r/r
673 ffff880123e3a318
674     position: root/r/r/r
675 ffff880123e3ad40
676     position: root/r/r/r/r
677
```

678 Given an mm_struct address of 0xffff880074b5be80, list the VMA tree in linear
679 order from the leftmost node progressing to the right using the -l option:

```
680
681 crash> tree -ls vm_area_struct.vm_start -o vm_area_struct.vm_rb \
682 -r mm_struct.mm_rb 0xffff880074b5be80 | paste - -
683 ffff88001f2c50e0      vm_start = 0x400000
684 ffff88001f2c5290      vm_start = 0xceb000
685 ffff880074bfc6c0      vm_start = 0xcec000
686 ffff88001f2c4bd0      vm_start = 0xd10000
687 ffff880074bfc948      vm_start = 0x1fe9000
688 ffff880036e54510      vm_start = 0x7ff6aa296000
689 ffff88001f2c5bd8      vm_start = 0x7ff6aa298000
690 ffff880036e54af8      vm_start = 0x7ff6aa497000
691 ffff880036e54f30      vm_start = 0x7ff6aa498000
692 ffff88000e06aa20      vm_start = 0x7ff6aa499000
693 ffff88000e06b368      vm_start = 0x7ff6ab95f000
694 ...
695 ffff88001f2c5e60      vm_start = 0x7ff6bc1af000
696 ffff88001f2c4ca8      vm_start = 0x7ff6bc1b6000
697 ffff88001f2c5008      vm_start = 0x7ff6bc200000
698 ffff88001f2c5d88      vm_start = 0x7ff6bc205000
699 ffff880074bfd6c8      vm_start = 0x7ff6bc206000
700 ffff88001f2c4288      vm_start = 0x7ff6bc207000
701 ffff88001f2c4510      vm_start = 0x7ffc7a5fc000
702 ffff88001f2c5b00      vm_start = 0x7ffc7a6d1000
703
```

704 Compared to the top/down root/leaves order:

```
705
706 crash> tree -s vm_area_struct.vm_start -o vm_area_struct.vm_rb \
707 -r mm_struct.mm_rb 0xffff880074b5be80 | paste - -
708 ffff88001f2c5a28      vm_start = 0x7ff6bbbb9000
709 ffff88001f2c55f0      vm_start = 0x7ff6bb252000
710 ffff88000e06a360      vm_start = 0x7ff6ac6c3000
711 ffff88001f2c4bd0      vm_start = 0xd10000
712 ffff88001f2c5290      vm_start = 0xceb000
713 ffff88001f2c50e0      vm_start = 0x400000
714 ffff880074bfc6c0      vm_start = 0xcec000
715 ffff88000e06b368      vm_start = 0x7ff6ab95f000
716 ffff88001f2c5bd8      vm_start = 0x7ff6aa298000
717 ffff880074bfc948      vm_start = 0x1fe9000
718 ffff880036e54510      vm_start = 0x7ff6aa296000
719 ffff880036e54f30      vm_start = 0x7ff6aa498000
720 ffff880036e54af8      vm_start = 0x7ff6aa497000
721 ffff88000e06aa20      vm_start = 0x7ff6aa499000
722 ffff88000e06ae58      vm_start = 0x7ff6ac1df000
723 ffff88000e06ba28      vm_start = 0x7ff6abefc000
724 ffff88000e06a6c0      vm_start = 0x7ff6ac41b000
```

```

725     ffff88001f2c4000      vm_start = 0x7ff6bac75000
726     ffff88000e06bd88      vm_start = 0x7ff6b2d00000
727     ffff88000e06b440      vm_start = 0x7ff6b28de000
728     ...
729     ffff880074bfd6c8      vm_start = 0x7ff6bc206000
730     ffff88001f2c4510      vm_start = 0x7ffc7a5fc000
731     ffff88001f2c5b00      vm_start = 0x7ffc7a6d1000
732

```

Display a list of the page structs in the radix tree of an address_space structure located at ffff88012d364de0:

```

735
736     crash> tree -t radix -r address_space.page_tree ffff88012d364de0
737     fffffea00040d12c0
738     fffffea00040d9a60
739     fffffea00040d9b08
740     fffffea000407eda8
741     fffffea0004084288
742     ...
743     fffffea000407bc70
744     fffffea00040baf48
745     fffffea0004043f48
746     fffffea000407de58
747

```

Add the -p option to the command above to show position information:

```

748
749
750     crash> tree -t radix -r address_space.page_tree ffff88012d364de0 -p
751     fffffea00040d12c0
752         index: 0   position: root/0/0
753     fffffea00040d9a60
754         index: 1   position: root/0/1
755     fffffea00040d9b08
756         index: 2   position: root/0/2
757     fffffea000407eda8
758         index: 3   position: root/0/3
759     fffffea0004084288
760         index: 4   position: root/0/4
761     ...
762     fffffea000407bc70
763         index: 217  position: root/3/25
764     fffffea00040baf48
765         index: 218  position: root/3/26
766     fffffea0004043f48
767         index: 219  position: root/3/27
768     fffffea000407de58
769         index: 220  position: root/3/28
770

```

Alternatively, take the address of the radix_tree_node from the radix_tree_root structure in the address_space structure above, and display the tree with the -N option:

```

774
775     crash> struct address_space.page_tree ffff88012d364de0
776         page_tree = {
777             height = 0x2,
778             gfp_mask = 0x20,
779             rnode = 0xfffff8801238add71
780         }
781     crash> tree -t radix -N 0xfffff8801238add71
782     fffffea00040d12c0
783     fffffea00040d9a60
784     fffffea00040d9b08
785     fffffea000407eda8
786     fffffea0004084288
787     fffffea00040843a0
788     ...
789

```

Using the same radix tree as above, display the flags and _count members of each page struct in the list, and force the output format to be hexadecimal:

```

793
794     crash> tree -t radix -N 0xfffff8801238add71 -s page.flags,_count -x
795     fffffea00040d12c0
796         flags = 0x40000000002006c
797         _count = {

```

```

798         counter = 0x7
799     }
800     fffffea00040d9a60
801     flags = 0x4000000002006c
802     _count = {
803         counter = 0x7
804     }
805     fffffea00040d9b08
806     flags = 0x4000000002006c
807     _count = {
808         counter = 0x7
809     }
810     fffffea000407eda8
811     flags = 0x4000000002006c
812     _count = {
813         counter = 0x7
814     }
815     ...
816

```

In more recent kernels, the XArray facility has replaced radix trees. Display a list of the page structs in the XArray of an address_space structure located at 0xfffff94c235e76828, where the i_pages field is an embedded xarray structure:

```

821
822     crash> tree -t xarray -r address_space.i_pages 0xfffff94c235e76828
823     ffffffcc005aa8380
824     ffffffcc005cafa80
825     ffffffcc005a79c80
826     ffffffcc005ccad80
827     ffffffcc005a72ec0
828     ffffffcc005e27c00
829     ffffffcc005ce3100
830     ffffffcc005ff8dc0
831     ffffffcc005c9a100
832     ffffffcc005a49e40
833     ffffffcc005c95a80
834

```

Add the -p option to the command above to show position information:

```

835
836
837     crash> tree -t xarray -r address_space.i_pages 0xfffff94c235e76828 -p
838     ffffffcc005aa8380
839         index: 90  position: root/1/26
840     ffffffcc005cafa80
841         index: 91  position: root/1/27
842     ffffffcc005a79c80
843         index: 92  position: root/1/28
844     ffffffcc005ccad80
845         index: 93  position: root/1/29
846     ffffffcc005a72ec0
847         index: 94  position: root/1/30
848     ffffffcc005e27c00
849         index: 95  position: root/1/31
850     ffffffcc005ce3100
851         index: 96  position: root/1/32
852     ffffffcc005ff8dc0
853         index: 97  position: root/1/33
854     ffffffcc005c9a100
855         index: 98  position: root/1/34
856     ffffffcc005a49e40
857         index: 99  position: root/1/35
858     ffffffcc005c95a80
859         index: 100 position: root/1/36
860

```

Alternatively, take the value found in the xa_head field from the xarray structure, and display the tree with the -N option:

```

861
862
863
864     crash> address_space.i_pages 0xfffff94c235e76828
865     i_pages = {
866         ... [ xa_lock field not shown ] ...
867         xa_flags = 1,
868         xa_head = 0xfffff94c23c1566ca
869     }
870     crash> tree -t x -N 0xfffff94c23c1566ca

```

```
871 ffffffcc005aa8380
872 ffffffcc005cafa80
873 ffffffcc005a79c80
874 ffffffcc005ccad80
875 ffffffcc005a72ec0
876 ffffffcc005e27c00
877 ffffffcc005ce3100
878 ffffffcc005ff8dc0
879 ffffffcc005c9a100
880 ffffffcc005a49e40
881 ffffffcc005c95a80
882
```

883 Using the same xarray command as above, display the flags and _refcount
884 members of each page struct in the list, and force the output format
885 to be hexadecimal:

```
886
887 crash> tree -t x -N 0xffff94c23c1566ca -s page.flags,_refcount -x
888 ffffffcc005aa8380
889   flags = 0x57ffffc0000014
890   _refcount = {
891     counter = 0x1
892   }
893 ffffffcc005cafa80
894   flags = 0x57ffffc0000014
895   _refcount = {
896     counter = 0x1
897   }
898 ffffffcc005a79c80
899   flags = 0x57ffffc0000014
900   _refcount = {
901     counter = 0x1
902   }
903 ffffffcc005ccad80
904   flags = 0x57ffffc0000014
905   _refcount = {
906     counter = 0x1
907   }
908 ffffffcc005a72ec0
909   flags = 0x57ffffc0000014
910   _refcount = {
911     counter = 0x1
912   }
913 ffffffcc005e27c00
914   flags = 0x57ffffc0000014
915   _refcount = {
916     counter = 0x1
917   }
918 ffffffcc005ce3100
919   flags = 0x57ffffc0000014
920   _refcount = {
921     counter = 0x1
922   }
923 ffffffcc005ff8dc0
924   flags = 0x57ffffc0000014
925   _refcount = {
926     counter = 0x1
927   }
928 ffffffcc005c9a100
929   flags = 0x57ffffc0000014
930   _refcount = {
931     counter = 0x1
932   }
933 ffffffcc005a49e40
934   flags = 0x57ffffc0000014
935   _refcount = {
936     counter = 0x1
937   }
938 ffffffcc005c95a80
939   flags = 0x57ffffc0000014
940   _refcount = {
941     counter = 0x1
942   }
943
```

```

944
945 NAME
946     alias - command aliases
947
948 SYNOPSIS
949     alias [alias] [command string]
950
951 DESCRIPTION
952     This command creates an alias for a given command string.  If no arguments
953     are entered, the current list of aliases are displayed.  If one argument is
954     entered, the command string for that alias, if any, is displayed.
955
956         alias    the single word to be used as an alias
957     command string  the word(s) that will be substituted for the alias
958
959     Aliases may be created in four manners:
960
961         1. entering the alias in $HOME/.crashrc.
962         2. entering the alias in .crashrc in the current directory.
963         3. executing an input file containing the alias command.
964         4. during runtime with this command.
965
966     During initialization, $HOME/.crashrc is read first, followed by the
967     .crashrc file in the current directory.  Aliases in the .crashrc file
968     in the current directory override those in $HOME/.crashrc.  Aliases
969     entered with this command or by runtime input file override those
970     defined in either .crashrc file.  Aliases may be deleted by entering an
971     empty string for the second argument.  If redirection characters are to
972     be part of the command string, the command string must be enclosed by
973     quotation marks.
974
975     Note that there are a number of helpful built-in aliases -- see the
976     first example below.
977
978 EXAMPLES
979     Display the currently-defined aliases, which in this example, only
980     consist of the built-in aliases:
981
982     crash> alias
983     ORIGIN  ALIAS      COMMAND
984     builtin man        help
985     builtin ?          help
986     builtin quit       q
987     builtin sf         set scroll off
988     builtin sn         set scroll on
989     builtin hex        set radix 16
990     builtin dec        set radix 10
991     builtin g          gdb
992     builtin px         p -x
993     builtin pd         p -d
994     builtin for        foreach
995     builtin size       *
996     builtin dmesg      log
997     builtin lsmod      mod
998     builtin last       ps -l
999
1000     Create a new alias to be added to the list:
1001
1002     crash> alias kp kmem -p
1003     ORIGIN  ALIAS      COMMAND
1004     runtime kp        kmem -p
1005
1006     Create an alias with redirection characters:
1007
1008     crash> alias ksd "kmem -p | grep slab | grep DMA"
1009     ORIGIN  ALIAS      COMMAND
1010     runtime ksd      kmem -p | grep slab | grep DMA
1011
1012     Remove an alias:
1013
1014     crash> alias kp ""
1015     alias deleted: kp
1016

```

```

1017
1018 NAME
1019     foreach - display command data for multiple tasks in the system
1020
1021 SYNOPSIS
1022     foreach [[pid | taskp | name | state | [kernel | user | gleader]] ...]
1023         command [flag] [argument]
1024
1025 DESCRIPTION
1026     This command allows for an examination of various kernel data associated
1027     with any, or all, tasks in the system, without having to set the context
1028     to each targeted task.
1029
1030     pid    perform the command(s) on this PID.
1031     taskp  perform the command(s) on task referenced by this hexadecimal
1032            task_struct pointer.
1033     name   perform the command(s) on all tasks with this name.  If the
1034            task name can be confused with a foreach command name, then
1035            precede the name string with a "\".  If the name string is
1036            enclosed within "'" characters, then the encompassed string
1037            must be a POSIX extended regular expression that will be used
1038            to match task names.
1039     user   perform the command(s) on all user (non-kernel) threads.
1040     gleader perform the command(s) on all user (non-kernel) thread group leaders.
1041     kernel perform the command(s) on all kernel threads.
1042     active perform the command(s) on the active thread on each CPU.
1043     state  perform the command(s) on all tasks in the specified state, which
1044            may be one of: RU, IN, UN, ST, ZO, TR, SW, DE, WA, PA, ID or NE.
1045
1046     If none of the task-identifying arguments above are entered, the command
1047     will be performed on all tasks.
1048
1049     command select one or more of the following commands to be run on the tasks
1050            selected, or on all tasks:
1051
1052            bt    run the "bt" command (optional flags: -r -t -l -e -R -f -F
1053                   -o -s -x -d)
1054            vm    run the "vm" command (optional flags: -p -v -m -R -d -x)
1055            task  run the "task" command (optional flags: -R -d -x)
1056            files run the "files" command (optional flag: -c -R)
1057            net   run the "net" command (optional flags: -s -S -R -d -x)
1058            set   run the "set" command
1059            ps    run the "ps" command (optional flags: -G -s -p -c -t -l -a
1060                   -g -r -y)
1061            sig   run the "sig" command (optional flag: -g)
1062            vtop  run the "vtop" command (optional flags: -c -u -k)
1063
1064     flag    Pass this optional flag to the command selected.
1065     argument Pass this argument to the command selected.
1066
1067     A header containing the PID, task address, cpu and command name will be
1068     pre-pended before the command output for each selected task.  Consult the
1069     help page of each of the command types above for details.
1070
1071 EXAMPLES
1072     Display the stack traces for all tasks:
1073
1074     crash> foreach bt
1075     PID: 4752    TASK: c7680000 CPU: 1    COMMAND: "xterm"
1076     #0 [c7681edc] schedule at c01135f6
1077         (void)
1078     #1 [c7681f34] schedule_timeout at c01131ff
1079         (24)
1080     #2 [c7681f64] do_select at c0132838
1081         (5, c7681fa4, c7681fa0)
1082     #3 [c7681fbc] sys_select at c0132dad
1083         (5, 8070300, 8070380, 0, 0)
1084     #4 [bffffb0c] system_call at c0109944
1085         EAX: 0000008e EBX: 00000005 ECX: 08070300 EDX: 08070380
1086         DS: 002b     ESI: 00000000 ES: 002b     EDI: 00000000
1087         SS: 002b     ESP: bffffad4 EBP: bffffb0c
1088         CS: 0023     EIP: 402259ee ERR: 0000008e EFLAGS: 00000246
1089

```

```

1090 PID: 557 TASK: c5600000 CPU: 0 COMMAND: "nfsd"
1091 #0 [c5601f38] schedule at c01135f6
1092 (void)
1093 #1 [c5601f90] schedule_timeout at c01131ff
1094 (c5600000)
1095 #2 [c5601fb8] svc_recv at c805363a
1096 (c0096f40, c5602800, 7fffffff, 100, c65c9f1c)
1097 #3 [c5601fec] (nfsd module) at c806e303
1098 (c5602800, c5602800, c0096f40, 6c6e0002, 50)
1099 #4 [c65c9f24] kernel_thread at c010834f
1100 (0, 0, ext2_file_inode_operations)
1101
1102 PID: 824 TASK: c7c84000 CPU: 0 COMMAND: "mingetty"
1103 ...
1104
1105 Display the task_struct structure for each "bash" command:
1106
1107 crash> foreach bash task
1108 ...
1109
1110 Display the open files for all tasks:
1111
1112 crash> foreach files
1113 ...
1114
1115 Display the state of tasks whose name contains a match to "event.*":
1116
1117 crash> foreach 'event.*' task -R state
1118 PID: 99 TASK: ffff8804750d5500 CPU: 0 COMMAND: "events/0"
1119 state = 1,
1120
1121 PID: 100 TASK: ffff8804750d4ac0 CPU: 1 COMMAND: "events/1"
1122 state = 1,
1123
1124 PID: 101 TASK: ffff8804750d4080 CPU: 2 COMMAND: "events/2"
1125 state = 1,
1126 ...
1127
1128 Display the stack traces for all blocked (TASK_UNINTERRUPTIBLE) tasks:
1129
1130 crash> foreach UN bt
1131 PID: 428 TASK: ffff880036b6c560 CPU: 1 COMMAND: "jbd2/dm-1-8"
1132 #0 [ffff880035779a70] __schedule at ffffffff815df272
1133 #1 [ffff880035779b08] schedule at ffffffff815dfacf
1134 #2 [ffff880035779b18] io_schedule at ffffffff815dfb7f
1135 #3 [ffff880035779b38] sleep_on_page at ffffffff81119a4e
1136 #4 [ffff880035779b48] __wait_on_bit at ffffffff815e039f
1137 #5 [ffff880035779b98] wait_on_page_bit at ffffffff81119bb8
1138 #6 [ffff880035779be8] filemap_fdatawait_range at ffffffff81119ccc
1139 #7 [ffff880035779cd8] filemap_fdatawait at ffffffff81119d8b
1140 #8 [ffff880035779ce8] jbd2_journal_commit_transaction at ffffffff8123a99c
1141 #9 [ffff880035779e58] kjournald2 at ffffffff8123ee7b
1142 #10 [ffff880035779ee8] kthread at ffffffff8108fb9c
1143 #11 [ffff880035779f48] kernel_thread_helper at ffffffff815ebaf4
1144 ...
1145
1146
1147
1148 NAME
1149 mount - mounted filesystem data
1150
1151 SYNOPSIS
1152 mount [-f][-i] [-n pid|task] [mount|vfsmount|superblock|dev|dir|dentry|inode]
1153
1154 DESCRIPTION
1155 This command displays basic information about the currently-mounted
1156 filesystems. The per-filesystem dirty inode list or list of open
1157 files for the filesystem may also be displayed.
1158
1159 -f dump dentries and inodes for open files in each filesystem; only
1160 supported on kernels prior to Linux 3.13.
1161 -i dump all dirty inodes associated with each filesystem; only
1162 supported on kernels prior to Linux 2.6.32.

```


For kernels supporting namespaces, the -n option may be used to display the mounted filesystems with respect to the namespace of a specified task:

-n pid a process PID.
-n task a hexadecimal task_struct pointer.

Specific filesystems may be selected using the following forms:

vfsmount hexadecimal address of a filesystem vfsmount structure.
mount hexadecimal address of a filesystem mount structure (Linux 3.3 and later).
superblock hexadecimal address of a filesystem super_block structure.
dev device name of a filesystem.
dir directory where a filesystem is mounted.
dentry hexadecimal address of an open dentry of a filesystem.
inode hexadecimal address of an open inode of a filesystem.

The first column of the command output displays the filesystem's vfsmount structure address for kernels prior to Linux 3.3. For Linux 3.3 and later kernels, the first column displays the filesystem's mount structure address, which contains an embedded vfsmount structure.

EXAMPLES

Display mounted filesystem data:

```
crash> mount
VFSMOUNT SUPERBLK TYPE DEVNAME DIRNAME
c0089ea0 c0088a00 ext2 /dev/root /
c0089cf0 c0088c00 proc /proc /proc
c0089e10 c0088800 ext2 /dev/sda5 /boot
c0089d80 c0088600 ext2 /dev/sda6 /usr
c0089f30 c0088400 devpts none /dev/pts
c3f4b010 c0088200 ext2 /dev/sda1 /home
```

On Linux 3.3 and later kernels, the filesystem's mount structure address is shown:

```
crash> mount
MOUNT SUPERBLK TYPE DEVNAME DIRNAME
ffff880212fb8200 ffff880212fc0800 rootfs rootfs /
ffff88020ffbea00 ffff880212fc2000 proc proc proc
ffff880211db7f00 ffff88020e01a800 sysfs sysfs /sys
ffff88020ffe1300 ffff880212a40000 devtmpfs devtmpfs /dev
ffff88020ff15000 ffff880212bbc800 devpts devpts /dev/pts
ffff88020e542800 ffff88020e62b800 tmpfs tmpfs /dev/shm
...
```

Display the open files associated with each mounted filesystem:

```
crash> mount -f
VFSMOUNT SUPERBLK TYPE DEVNAME DIRNAME
c7fb2b80 c7fb3200 ext2 /dev/root /
OPEN FILES:
DENTRY INODE TYPE PATH
c6d02200 c6d0f7a0 REG usr/X11R6/lib/libX11.so.6.1
c6d02100 c6d0f9e0 REG usr/X11R6/lib/libXext.so.6.3
c6d02000 c6d0fc20 REG usr/X11R6/lib/libICE.so.6.3
c6d02680 c6d0f320 REG usr/X11R6/bin/xf86
c7106580 c70c5440 CHR dev/psaux
...
```

Display the dirty inodes associated with each mounted filesystem:

```
crash> mount -i
VFSMOUNT SUPERBLK TYPE DEVNAME DIRNAME
c0089ea0 c0088a00 ext2 /dev/root /
DIRTY INODES
c7ad4008
c2233438
c72c4008
c7d6b548
```

```

1236     c3af1a98
1237     c7d6b768
1238     c3c4e228
1239     ...
1240
1241 Display the mounted filesystem containing inode c5000aa8:
1242
1243 crash> mount c5000aa8
1244 VFSMOUNT SUPERBLK TYPE      DEVNAME  DIRNAME
1245 c0089f30 c0088600 ext2      /dev/sda6 /usr
1246
1247 Display the mounted filesystem containing inode ffff8801f4245e40:
1248
1249 crash> mount ffff8801f4245e40
1250 MOUNT      SUPERBLK      TYPE      DEVNAME  DIRNAME
1251 ffff88020ffbea00 ffff880212fc2000 proc      proc      /proc
1252
1253
1254
1255 NAME
1256     search - search memory
1257
1258 SYNOPSIS
1259     search [-s start] [ -[kKV] | -u | -p | -t | -T ] [-e end | -l length] [-m mask]
1260             [-x count] [-cwh] [value | (expression) | symbol | string] ...
1261
1262 DESCRIPTION
1263     This command searches for a given value within a range of user virtual, kernel
1264     virtual, or physical memory space. If no end nor length value is entered,
1265     then the search stops at the end of user virtual, kernel virtual, or physical
1266     address space, whichever is appropriate.
1267
1268     An optional mask value may be entered to mask off "don't care" bits.
1269
1270     -s start    Start the search at this hexadecimal user or kernel virtual
1271                 address, physical address, or kernel symbol. The start address
1272                 must be appropriate for the memory type specified; if no memory
1273                 type is specified, the default is kernel virtual address space.
1274     -k          If no start address is specified, start the search at the base
1275                 of kernel virtual address space. This option is the default.
1276     -K          Same as -k, except that mapped kernel virtual memory that was
1277                 allocated by vmalloc(), module memory, or virtual mem_map regions
1278                 will not be searched.
1279     -V          Same as -k, except that unity-mapped kernel virtual memory and
1280                 mapped kernel-text/static-data (x86_64 and ia64) will not be
1281                 searched.
1282     -u          If no start address is specified, start the search at the base
1283                 of the current context's user virtual address space. If a start
1284                 address is specified, then this option specifies that the start
1285                 address is a user virtual address.
1286     -p          If no start address is specified, start the search at the base
1287                 of physical address space. If a start address is specified,
1288                 then this option specifies that the start address is a physical
1289                 address.
1290     -t          Search only the kernel stack pages of every task. If one or more
1291                 matches are found in a task's kernel stack, precede the output
1292                 with a task-identifying header.
1293     -T          Same as -t, except only the active task(s) are considered.
1294     -e end      Stop the search at this hexadecimal user or kernel virtual
1295                 address, kernel symbol, or physical address. The end address
1296                 must be appropriate for the memory type specified.
1297     -l length   Length in bytes of address range to search.
1298     -m mask     Ignore the bits that are set in the hexadecimal mask value.
1299     -c          Search for character string values instead of unsigned longs. If
1300                 the string contains any space(s), it must be encompassed by double
1301                 quotes.
1302     -w          Search for unsigned hexadecimal ints instead of unsigned longs.
1303                 This is only meaningful on 64-bit systems in order to search both
1304                 the upper and lower 32-bits of each 64-bit long for the value.
1305     -h          Search for unsigned hexadecimal shorts instead of unsigned longs.
1306     -x count    Display the memory contents before and after any found value. The
1307                 before and after memory context will consist of "count" memory
1308                 items of the same size as the "value" argument. This option is

```

1309 not applicable with the -c option.
 1310 value Search for this hexadecimal long, unless modified by the -c, -w,
 1311 or -h options.
 1312 (expression) Search for the value of this expression; the expression value must
 1313 not overflow the designated size when -h or -w are used; not
 1314 applicable when used with the -c option.
 1315 symbol Search for this symbol value; the symbol value must not overflow
 1316 the designated size when -h or -w are used; not applicable when
 1317 used with the -c option.
 1318 string Search for character string values; if the string contains any
 1319 space(s), it must be encompassed by double quotes; only applicable
 1320 with the -c option.
 1321

1322 If -k, -K, -V, -u, -p or -t are not used, then the search defaults to kernel
 1323 virtual address space. The starting address must be long-word aligned.
 1324 Address ranges that start in user space and end in kernel space are not
 1325 accepted.
 1326

1327 EXAMPLES

1328 Search the current context's address space for all instances of 0xdeadbeef:

```
1329
1330 crash> search -u deadbeef
1331 81aba5c: deadbeef
1332 81abaa8: deadbeef
1333 bffffc698: deadbeef
1334 bffff390: deadbeef
1335
```

1336 Search all kernel memory above the kernel text space for all instances
 1337 of 0xabcd occurring in the lower 16-bits of each 32-bit word:

```
1338
1339 crash> search -s _etext -m ffff0000 abcd
1340 c071481c: abcd
1341 c0c2b0fc: 804abcd
1342 c0cf5e74: 7489abcd
1343 c17c0b44: c012abcd
1344 c1dac730: 3dbeabcd
1345 c226d0e8: ffffabcd
1346 c23ed5dc: abcd
1347 c3022544: 3dbeabcd
1348 c3069b58: 3dbeabcd
1349 c3e86e84: aabcd
1350 c3e88ed0: aabcd
1351 c3e8ee5c: aabcd
1352 c3e9df50: aabcd
1353 c3e9e930: aabcd
1354 c440a778: 804abcd
1355 c486eb44: 3dbeabcd
1356 c578f0fc: 804abcd
1357 c6394f90: 8ababcd
1358 c65219f0: 3abcd
1359 c661399c: abcd
1360 c68514ac: 8abcd
1361 c7e036bc: 3dbeabcd
1362 c7e12568: 5abcd
1363 c7e1256c: 5abcd
1364
```

1365 Search the 4K page at c532c000 for all instances of 0xffffffff:

```
1366
1367 crash> search -s c532c000 -l 4096 ffffffff
1368 c532c33c: ffffffff
1369 c532c3fc: ffffffff
1370
```

1371 Search the static kernel data area for all instances of c2d400eb:

```
1372
1373 crash> search -s _etext -e _edata c2d400eb
1374 c022b550: c2d400eb
1375 c022b590: c2d400eb
1376 c022b670: c2d400eb
1377 c022b6e0: c2d400eb
1378 c022b7b0: c2d400eb
1379 c022b7e0: c2d400eb
1380 c022b8b0: c2d400eb
1381
```

1382 Search physical memory for all instances of 0xbabe occurring in the
1383 upper 16 bits of each 32-bit word:
1384

```
1385 crash> search -p babe0000 -m ffff
1386 2a1dc4: babe671e
1387 2b6928: babe3de1
1388 2f99ac: babe0d54
1389 31843c: babe70b9
1390 3ba920: babeb5d7
1391 413ce4: babe7540
1392 482747c: babe2600
1393 48579a4: babe2600
1394 4864a68: babe2600
1395 ...
```

1396
1397 Search physical memory for all instances of 0xbabe occurring in the
1398 upper 16 bits of each 32-bit word on a 64-bit system:
1399

```
1400 crash> search -p babe0000 -m ffff -w
1401 102e248: babel174
1402 11d2f90: babe813d
1403 122d3ad70: babe6b27
1404 124d8cd30: babe3dc8
1405 124d8eefc: babef981
1406 124d8f060: babe3dc8
1407 124d8f17c: babefc81
1408 ...
```

1409
1410 Search kernel memory for all instances of 32-bit value 0xbabel174
1411 on a 64-bit system:
1412

```
1413 crash> search -k -w babel174
1414 ffff88000102e248: babel174
1415 ffffffff8102e248: babel174
```

1416
1417 Search kernel memory for two strings:
1418

```
1419 crash> search -k -c "can't allocate memory" "Failure to"
1420 ffff8800013dde1: can't allocate memory for key lists..<3>%s %s: error con
1421 ffff8801258be748: Failure to install fence: %d..<3>[drm:%s] *ERROR* Failed
1422 ffff880125f07ec9: can't allocate memory..<3>ACPI: Invalid data..Too many d
1423 ffffffff813dde1: can't allocate memory for key lists..<3>%s %s: error con
```

1424
1425 Search the kernel stacks of all tasks for those that contain the inode
1426 address ffff81002c0a3050:
1427

```
1428 crash> search -t ffff81002c0a3050
1429 PID: 4876 TASK: ffff81003e9f5860 CPU: 7 COMMAND: "automount"
1430 ffff8100288fbe98: ffff81002c0a3050
1431
1432 PID: 4880 TASK: ffff81003ce967a0 CPU: 0 COMMAND: "automount"
1433 ffff81002c0fbdd8: ffff81002c0a3050
1434 ffff81002c0fbe78: ffff81002c0a3050
```

1435
1436 When a kernel symbol or an (expression) is used an argument, both the
1437 resultant value and the input string are displayed:
1438

```
1439 crash> search anon_inode_inode (__down_interruptible+191)
1440 ffff81000222a728: ffffffff80493d60 (anon_inode_inode)
1441 ffff810005a1e918: ffffffff800649d6 (__down_interruptible+191)
1442 ffff810005a1e9d0: ffffffff800649d6 (__down_interruptible+191)
1443 ffff810005a1eb48: ffffffff800649d6 (__down_interruptible+191)
1444 ffff81000b409c60: ffffffff80493d60 (anon_inode_inode)
1445 ffff81000c155b98: ffffffff80493d60 (anon_inode_inode)
1446 ffff8100194fac70: ffffffff80493d60 (anon_inode_inode)
1447 ffff81001daa1008: ffffffff80493d60 (anon_inode_inode)
1448 ffff810028b95830: ffffffff800649d6 (__down_interruptible+191)
1449 ffff81002cea0c70: ffffffff80493d60 (anon_inode_inode)
1450 ffff810031327268: ffffffff80493d60 (anon_inode_inode)
1451 ffff810031327270: ffffffff800649d6 (__down_interruptible+191)
1452 ffff810034b1ccd0: ffffffff800649d6 (__down_interruptible+191)
1453 ffff8100399565a8: ffffffff80493d60 (anon_inode_inode)
1454 ffff81003a278cd0: ffffffff800649d6 (__down_interruptible+191)
```

ffff81003cc23e08: ffffffff800649d6 (__down_interruptible+191)

NAME

union - union contents

SYNOPSIS

union union_name[.member[,member]] [-o][-l offset][-rfuxdp]
[address | symbol][:cpuspec] [count | -c count]

DESCRIPTION

This command displays either a union definition, or a formatted display of the contents of a union at a specified address. When no address is specified, the union definition is shown along with the union size. A union member may be appended to the structure name in order to limit the scope of the data displayed to that particular member; when no address is specified, the member's offset (always 0) and definition are shown.

union_name name of a C-code union used by the kernel.
.member name of a union member; to display multiple members of a union, use a comma-separated list of members. If any member contains an embedded structure, or the member is an array, the output may be restricted to just the embedded structure or an array element by expressing the argument as "member.member" or "member[index]"; embedded member specifications may extend beyond one level deep, by expressing the member argument as "member.member.member...".

-o show member offsets when displaying union definitions; the offset is always 0 unless used with an address or symbol argument, in which case each member will be preceded by its virtual address.

-l offset if the address argument is a pointer to a list_head structure that is embedded in the target union structure, the offset to the list_head member may be entered in either of the following manners:

1. in "structure.member" format.
2. a number of bytes.

-r raw dump of union data.

-f address argument is a dumpfile offset.

-x override default output format with hexadecimal format.

-d override default output format with decimal format.

-p if a union member is a pointer value, show the member's data type on the output line; and on the subsequent line(s), dereference the pointer, display the pointer target's symbol value in brackets if appropriate, and if possible, display the target data; requires an address argument.

-u address argument is a user virtual address in the current context.

address hexadecimal address of a union; if the address points to an embedded list_head structure contained within the target union structure, then the "-l" option must be used.

symbol symbolic reference to the address of a union.

:cpuspec CPU specification for a per-cpu address or symbol:

- : CPU of the currently selected task.
- :a[ll] all CPUs.
- :#[-#][,...] CPU list(s), e.g. "1,3,5", "1-3", or "1,3,5-7,10".

count count of unions to dump from an array of unions; if used, this must be the last argument entered.

-c count "-c" is only required if "count" is not the last argument entered or if a negative number is entered; if a negative value is entered, the (positive) "count" structures that lead up to and include the target structure will be displayed.

Union data, sizes, and member offsets are shown in the current output radix unless the -x or -d option is specified.

Please note that in the vast majority of cases, the "union" command name may be dropped; if the union name does not conflict with any crash or gdb command name, then the "union_name[.member]" argument will be recognized as a union name, and this command automatically executed. See the NOTE below.

```

1528 EXAMPLES
1529
1530 Display the bdflush_param union definition, and then an instance of it:
1531
1532 crash> union bdflush_param
1533 union bdflush_param {
1534     struct {
1535         int nfraction;
1536         int ndirty;
1537         int nrefill;
1538         int nref_dirt;
1539         int dummy1;
1540         int age_buffer;
1541         int age_super;
1542         int dummy2;
1543         int dummy3;
1544     } b_un;
1545     unsigned int data[9];
1546 }
1547
1548 SIZE: 36 (0x24)

```

```

1549
1550 crash> union bdflush_param bdf_prm
1551 union bdflush_param {
1552     b_un = {
1553         nfraction = 40,
1554         ndirty = 500,
1555         nrefill = 64,
1556         nref_dirt = 256,
1557         dummy1 = 15,
1558         age_buffer = 3000,
1559         age_super = 500,
1560         dummy2 = 1884,
1561         dummy3 = 2
1562     },
1563     data = {40, 500, 64, 256, 15, 3000, 500, 1884, 2}
1564 }
1565

```

NOTE

If the union name does not conflict with any crash command name, the "union" command may be dropped. Accordingly, the examples above could also have been accomplished like so:

```

1570
1571 crash> bdflush_param
1572 crash> bdflush_param bdf_prm
1573

```

Lastly, the short-cut "*" (pointer-to) command may also be used to negate the need to enter the "union" command name (enter "help *" for details).

NAME

ascii - translate a hexadecimal string to ASCII

SYNOPSIS

ascii value ...

DESCRIPTION

Translates 32-bit or 64-bit hexadecimal values to ASCII. If no argument is entered, an ASCII chart is displayed.

EXAMPLES

Translate the hexadecimal value of 0x62696c2f7273752f to ASCII:

```

1591 crash> ascii 62696c2f7273752f
1592 62696c2f7273752f: /usr/lib
1593

```

Display an ASCII chart:

```

1594 crash> ascii
1595
1596
1597
1598      0      1      2      3      4      5      6      7
1599  +-----+
1600  0 | NUL DLE  SP   0   @   P   '   p

```

```

1601      1 | SOH DC1  !    1    A    Q    a    q
1602      2 | STX DC2  "    2    B    R    b    r
1603      3 | ETX DC3  #    3    C    S    c    s
1604      4 | EOT DC4  $    4    D    T    d    t
1605      5 | ENQ NAK  %    5    E    U    e    u
1606      6 | ACK SYN  &    6    F    V    f    v
1607      7 | BEL ETB  `    7    G    W    g    w
1608      8 |  BS CAN  (    8    H    X    h    x
1609      9 |  HT EM   )    9    I    Y    i    y
1610     A |  LF SUB  *    :    J    Z    j    z
1611     B |  VT ESC  +    ;    K    [    k    {
1612     C |  FF  FS  ,    <    L    \    l    |
1613     D |  CR  GS  _    =    M    ]    m    }
1614     E |  SO  RS  .    >    N    ^    n    ~
1615     F |  SI  US  /    ?    O    -    o    DEL

```

```

1616
1617

```

NAME

```

1618     fuser - file users

```

```

1620

```

SYNOPSIS

```

1621     fuser [pathname | inode]

```

```

1623

```

DESCRIPTION

```

1625     This command displays the tasks using specified files or sockets.
1626     Tasks will be listed that reference the file as the current working
1627     directory, root directory, an open file descriptor, or that mmap the
1628     file. If the file is held open in the kernel by the lockd server on
1629     behalf of a client discretionary file lock, the client hostname is
1630     listed.

```

```

1631

```

```

1632     pathname  the full pathname of the file.
1633     inode     the hexadecimal inode address for the file.

```

```

1634

```

EXAMPLES

```

1635     Display the tasks using file /usr/lib/libkfm.so.2.0.0

```

```

1637

```

```

1638     crash> fuser /usr/lib/libkfm.so.2.0.0
1639     PID    TASK    COMM                USAGE
1640     779    c5e82000  "kwm"              mmap
1641     808    c5a8e000  "krootwm"          mmap
1642     806    c5b42000  "kfm"              mmap
1643     809    c5dde000  "kpanel"           mmap

```

```

1644

```

```

1645

```

NAME

```

1646     net - network command

```

```

1648

```

SYNOPSIS

```

1650     net [[-s | -S] [-xd] [-R ref] [pid | task]] [-a] [ -n [pid | task]] [-N addr]

```

```

1651

```

DESCRIPTION

```

1652     Displays various network related data.

```

```

1654

```

```

1655     If no arguments are entered, the list of network devices, names and IP
1656     addresses are displayed. For kernels supporting namespaces, the -n option
1657     may be used to display the list of network devices with respect to the
1658     network namespace of a current context or a task specified by pid or task:

```

```

1659

```

```

1660     -n      the namespace of the current context.
1661     -n pid  a process PID.
1662     -n task a hexadecimal task_struct pointer.

```

```

1663

```

```

1664     The -s and -S options display data with respect to the current context, but
1665     may be appended with an argument to show the socket data with respect
1666     to a specified task:

```

```

1667

```

```

1668     -s      display open network socket/sock addresses, their family and type,
1669             and for INET and INET6 families, their source and destination
1670             addresses and ports.
1671     -s pid  same as above, for task with process PID pid.
1672     -s task same as above, for task with hexadecimal task_struct pointer task.

```

```

1673

```

```

1674         -S displays open network socket/sock addresses followed by a dump
1675         of both data structures.
1676         -S pid same as above, with respect to process PID.
1677         -S task same as above, with respect to hexadecimal task_struct pointer.
1678
1679 The -R option, typically invoked from "foreach net", and in conjunction
1680 with the -s or -S options, searches for references to a socket address,
1681 sock address, or a file descriptor; if found, only the referenced fd, socket
1682 or sock data will be displayed:

```

```

1683
1684     -R ref  socket or sock address, or file descriptor.
1685
1686 Other options:

```

```

1687
1688     -a display the ARP cache.
1689     -N addr translates an IPv4 address expressed as a decimal or hexadecimal
1690         value into a standard numbers-and-dots notation.
1691     -x override default output format with hexadecimal format.
1692     -d override default output format with decimal format.
1693

```

EXAMPLES

```

1694 Display the system's network device list:

```

```

1695
1696 crash> net
1697
1698     NET_DEVICE      NAME      IP ADDRESS(ES)
1699 ffff8803741c0000   lo       127.0.0.1
1700 fff88037059c0000  eth0     10.226.229.141
1701 ffff8803705c0000  eth1     10.226.228.250
1702 ffff880374ad6000  usb0     169.254.95.120
1703

```

```

1704 Display the network device list with respect to the network namespace
1705 of PID 2618:

```

```

1706
1707 crash> net -n 2618
1708
1709     NET_DEVICE      NAME      IP ADDRESS(ES)
1710 ffff880456ee7020   lo       127.0.0.1
1711 ffff8804516a1020  eth0     10.1.9.223
1712

```

```

1713 Dump the ARP cache:

```

```

1714 crash> net -a
1715
1716 NEIGHBOUR      IP ADDRESS      HW TYPE      HW ADDRESS      DEVICE      STATE
1717 f38d1b00       10.16.64.14     ETHER        00:16:3e:4b:a5:4a eth1        STALE
1718 f38d1080       0.0.0.0         UNKNOWN      00 00 00 00 00 00 lo          NOARP
1719 f38d1bc0       10.16.71.254    ETHER        00:00:0c:07:ac:00 eth1        REACHABLE
1720 f38d1200       10.16.64.21     ETHER        00:16:3e:51:d8:09 eth1        REACHABLE
1721

```

```

1722 Display the sockets for PID 2517, using both -s and -S output formats:

```

```

1723 crash> net -s 2517
1724 PID: 2517  TASK: c1598000 CPU: 1  COMMAND: "rlogin"
1725 FD  SOCKET      SOCK      FAMILY:TYPE      SOURCE-PORT      DESTINATION-PORT
1726 3   c57375dc  c1fff1850  INET:STREAM      10.1.8.20-1023    10.1.16.62-513
1727

```

```

1728 crash> net -S 2517
1729 PID: 2517  TASK: c1598000 CPU: 1  COMMAND: "rlogin"
1730 FD  SOCKET      SOCK
1731 3   c57375dc  c1fff1850
1732

```

```

1733 struct socket {
1734     state = SS_CONNECTED,
1735     flags = 131072,
1736     ops = 0xc023f820,
1737     inode = 0xc5737540,
1738     fasync_list = 0x0,
1739     file = 0xc58892b0,
1740     sk = 0xc1fff1850,
1741     wait = 0xc14d9ed4,
1742     type = 1,
1743     passcred = 0 '\000',
1744     tli = 0 '\000'
1745 }

```

```

1746 struct sock {

```



```

1747     sklist_next = 0xc1fff12f0,
1748     sklist_prev = 0xc216bc00,
1749     bind_next = 0x0,
1750     bind_pprev = 0xc0918448,
1751     daddr = 1041236234,
1752     rcv_saddr = 336068874,
1753     dport = 258,
1754     num = 1023,
1755     bound_dev_if = 0,
1756     next = 0x0,
1757     pprev = 0xc0286dd4,
1758     state = 1 '\001',
1759     zapped = 0 '\000',
1760     sport = 65283,
1761     family = 2,
1762     reuse = 0 '\000',
1763     ...

```

Translate the rcv_saddr from above into dotted-decimal notation:

```

1766 crash> net -N 1041236234
1767 10.1.16.62

```

From "foreach", find all tasks with references to socket c08ea3cc:

```

1770 crash> foreach net -s -R c08ea3cc
1771 PID: 2184  TASK: c7026000  CPU: 1  COMMAND: "klines.kss"
1772 FD  SOCKET      SOCK      FAMILY:TYPE      SOURCE-PORT      DESTINATION-PORT
1773 5  c08ea3cc  c50d3c80  INET:STREAM      0.0.0.0-1026      0.0.0.0-0
1774
1775
1776 PID: 2200  TASK: c670a000  CPU: 1  COMMAND: "kpanel"
1777 FD  SOCKET      SOCK      FAMILY:TYPE      SOURCE-PORT      DESTINATION-PORT
1778 5  c08ea3cc  c50d3c80  INET:STREAM      0.0.0.0-1026      0.0.0.0-0
1779
1780 PID: 2201  TASK: c648a000  CPU: 1  COMMAND: "kbgndwm"
1781 FD  SOCKET      SOCK      FAMILY:TYPE      SOURCE-PORT      DESTINATION-PORT
1782 5  c08ea3cc  c50d3c80  INET:STREAM      0.0.0.0-1026      0.0.0.0-0
1783
1784 PID: 19294 TASK: c250a000  CPU: 0  COMMAND: "prefdm"
1785 FD  SOCKET      SOCK      FAMILY:TYPE      SOURCE-PORT      DESTINATION-PORT
1786 5  c08ea3cc  c50d3c80  INET:STREAM      0.0.0.0-1026      0.0.0.0-0
1787
1788 PID: 2194  TASK: c62dc000  CPU: 1  COMMAND: "kaudioserver"
1789 FD  SOCKET      SOCK      FAMILY:TYPE      SOURCE-PORT      DESTINATION-PORT
1790 5  c08ea3cc  c50d3c80  INET:STREAM      0.0.0.0-1026      0.0.0.0-0
1791
1792 PID: 2195  TASK: c6684000  CPU: 1  COMMAND: "maudio"
1793 FD  SOCKET      SOCK      FAMILY:TYPE      SOURCE-PORT      DESTINATION-PORT
1794 5  c08ea3cc  c50d3c80  INET:STREAM      0.0.0.0-1026      0.0.0.0-0
1795
1796 PID: 2196  TASK: c6b58000  CPU: 1  COMMAND: "kwmsound"
1797 FD  SOCKET      SOCK      FAMILY:TYPE      SOURCE-PORT      DESTINATION-PORT
1798 5  c08ea3cc  c50d3c80  INET:STREAM      0.0.0.0-1026      0.0.0.0-0
1799
1800 PID: 2197  TASK: c6696000  CPU: 0  COMMAND: "kfm"
1801 FD  SOCKET      SOCK      FAMILY:TYPE      SOURCE-PORT      DESTINATION-PORT
1802 5  c08ea3cc  c50d3c80  INET:STREAM      0.0.0.0-1026      0.0.0.0-0
1803
1804 PID: 2199  TASK: c65ec000  CPU: 0  COMMAND: "krootwm"
1805 FD  SOCKET      SOCK      FAMILY:TYPE      SOURCE-PORT      DESTINATION-PORT
1806 5  c08ea3cc  c50d3c80  INET:STREAM      0.0.0.0-1026      0.0.0.0-0
1807
1808 PID: 694   TASK: c1942000  CPU: 0  COMMAND: "prefdm"
1809 FD  SOCKET      SOCK      FAMILY:TYPE      SOURCE-PORT      DESTINATION-PORT
1810 5  c08ea3cc  c50d3c80  INET:STREAM      0.0.0.0-1026      0.0.0.0-0
1811
1812 PID: 698   TASK: c6a2c000  CPU: 1  COMMAND: "x"
1813 FD  SOCKET      SOCK      FAMILY:TYPE      SOURCE-PORT      DESTINATION-PORT
1814 5  c08ea3cc  c50d3c80  INET:STREAM      0.0.0.0-1026      0.0.0.0-0
1815
1816 PID: 2159  TASK: c4a5a000  CPU: 1  COMMAND: "kwm"
1817 FD  SOCKET      SOCK      FAMILY:TYPE      SOURCE-PORT      DESTINATION-PORT
1818 5  c08ea3cc  c50d3c80  INET:STREAM      0.0.0.0-1026      0.0.0.0-0
1819

```

1820
1821
1822
1823
1824
1825
1826
1827
1828
1829
1830
1831
1832
1833
1834
1835
1836
1837
1838
1839
1840
1841
1842
1843
1844
1845
1846
1847
1848
1849
1850
1851
1852
1853
1854
1855
1856
1857
1858
1859
1860
1861
1862
1863
1864
1865
1866
1867
1868
1869
1870
1871
1872
1873
1874
1875
1876
1877
1878
1879
1880
1881
1882
1883
1884
1885
1886
1887
1888
1889
1890
1891
1892

NAME

set - set a process context or internal crash variable

SYNOPSIS

set [[-a] [pid | taskp] | [-c cpu] | -p] | [crash_variable [setting]] | -v

DESCRIPTION

This command either sets a new context, or gets the current context for display. The context can be set by the use of:

- pid a process PID.
- taskp a hexadecimal task_struct pointer.
- a sets the pid or task as the active task on its cpu (dumpfiles only).
- c cpu sets the context to the active task on a cpu (dumpfiles only).
- p sets the context to the panic task, or back to the crash task on a live system.
- v display the current state of internal crash variables.

If no argument is entered, the current context is displayed. The context consists of the PID, the task pointer, the CPU, and task state. The task state shows the bits found in both the task_struct state and exit_state fields.

This command may also be used to set internal crash variables. If no value argument is entered, the current value of the crash variable is shown. These are the crash variables, acceptable arguments, and purpose:

scroll	on off	controls output scrolling.
scroll	less	/usr/bin/less as the output scrolling program.
scroll	more	/bin/more as the output scrolling program.
scroll	CRASHPAGER	use CRASHPAGER environment variable as the output scrolling program.
radix	10 16	sets output radix to 10 or 16.
refresh	on off	controls internal task list refresh.
print_max	number	set maximum number of array elements to print.
print_array	on off	if on, set gdb's printing of arrays to "pretty" format, with one line per element.
console	device-name	sets debug console device.
debug	number	sets crash debug level.
core	on off	if on, drops core when the next error message is displayed.
hash	on off	controls internal list verification.
silent	on off	turns off initialization messages; turns off crash prompt during input file execution. (scrolling is turned off if silent is on)
edit	vi emacs	set line editing mode (from .crashrc file only).
namelist	filename	name of kernel (from .crashrc file only).
zero_excluded	on off	controls whether excluded pages, or pages that are missing from an incomplete dumpfile, should return zero-filled memory when read.
null-stop	on off	if on, gdb's printing of character arrays will stop at the first NULL encountered.
gdb	on off	if on, the crash session will be run in a mode where all commands will be passed directly to gdb, and the command prompt will change to "gdb>"; when running in this mode, native crash commands may be executed by preceding them with the "crash" directive.
scope	text-addr	sets the text scope for viewing the definition of data structures; the "text-addr" argument must be a kernel or module text address, which may be expressed symbolically or as a hexadecimal value; set scope 0 to un-set.
offline	show hide	show or hide command output that is associated with offline cpus.
redzone	on off	if on, CONFIG_SLUB object addresses displayed by the kmem command will point to the SLAB_RED_ZONE padding inserted at the beginning of the object.
error	default redirect filename	set the destination of error messages. "default": error messages are always displayed on the console; if the output of a command is

1893 piped to an external command or redirected
1894 to a file, the error messages are also sent
1895 to the pipe or file.
1896 "redirect": if the output of a command is piped
1897 to an external command or redirected to a file,
1898 error messages are only sent to the pipe or
1899 file; otherwise they are displayed on the
1900 console.
1901 "filename": error messages are only sent to the
1902 specified filename; they are not displayed on
1903 the console and are not sent to a pipe or file.
1904

1905 Internal variables may be set in four manners:
1906

- 1907 1. entering the set command in \$HOME/.crashrc.
 - 1908 2. entering the set command in .crashrc in the current directory.
 - 1909 3. executing an input file containing the set command.
 - 1910 4. during runtime with this command.
- 1911

1912 During initialization, \$HOME/.crashrc is read first, followed by the
1913 .crashrc file in the current directory. Set commands in the .crashrc file
1914 in the current directory override those in \$HOME/.crashrc. Set commands
1915 entered with this command or by runtime input file override those
1916 defined in either .crashrc file. Multiple set command arguments or argument
1917 pairs may be entered in one command line.
1918

1919 EXAMPLES

1920 Set the current context to task c2fe8000:
1921

```
1922 crash> set c2fe8000
1923     PID: 15917
1924  COMMAND: "bash"
1925     TASK: c2fe8000
1926     CPU: 0
1927  STATE: TASK_INTERRUPTIBLE
1928
```

1929 Set the context back to the panicking task:
1930

```
1931 crash> set -p
1932     PID: 698
1933  COMMAND: "gen12"
1934     TASK: f9d78000
1935     CPU: 2
1936  STATE: TASK_RUNNING (PANIC)
1937
```

1938 Turn off output scrolling:
1939

```
1940 crash> set scroll off
1941 scroll: off (/usr/bin/less)
1942
```

1943 Show the current state of crash internal variables:
1944

```
1945 crash> set -v
1946     scroll: on (/usr/bin/less)
1947     radix: 10 (decimal)
1948     refresh: on
1949   print_max: 256
1950  print_array: off
1951     console: /dev/pts/2
1952     debug: 0
1953     core: off
1954     hash: on
1955     silent: off
1956     edit: vi
1957   namelist: vmlinux
1958  zero_excluded: off
1959    null-stop: on
1960     gdb: off
1961     scope: (not set)
1962  offline: show
1963   redzone: on
1964     error: default
1965
```

```

1966 Show the current context:
1967
1968 crash> set
1969     PID: 1525
1970     COMMAND: "bash"
1971     TASK: c1ede000
1972     CPU: 0
1973     STATE: TASK_INTERRUPTIBLE
1974
1975
1976
1977 NAME
1978     vm - virtual memory
1979
1980 SYNOPSIS
1981     vm [-p | -P vma | -M mm | -v | -m | -x | -d | [-R reference] [pid | task]]
1982     [-f vm_flags]
1983
1984 DESCRIPTION
1985     This command displays basic virtual memory information of a context,
1986     consisting of a pointer to its mm_struct and page directory, its RSS and
1987     total virtual memory size; and a list of pointers to each vm_area_struct,
1988     its starting and ending address, vm_flags value, and file pathname. If no
1989     arguments are entered, the current context is used. Additionally, the -p
1990     option translates each virtual page of each VM area to its physical address.
1991     The -R option, typically invoked from "foreach vm", searches for references
1992     to a supplied number, address, or filename argument, and prints only the
1993     essential information leading up to and including the reference.
1994     Alternatively, the -m or -v options may be used to dump the task's mm_struct
1995     or all of its vm_area_structs respectively. The -p, -v, -m, -R and -f
1996     options are all mutually exclusive.
1997
1998     -p translate each virtual page to its physical address, or if
1999     the page is not mapped, its swap device and offset, or
2000     filename and offset.
2001     -P vma similar to -p, but only translate the pages belonging to the
2002     specified VM area of a context.
2003     -M mm if the mm_struct address has been removed from the task_struct
2004     of an exiting task, the virtual memory data cannot be displayed.
2005     However, if the address can be determined from the kernel stack,
2006     it can be entered manually in order to try to resurrect the
2007     virtual memory data of the task.
2008     -R reference search for references to this number or filename.
2009     -m dump the mm_struct associated with the task.
2010     -v dump all of the vm_area_structs associated with the task.
2011     -x override the default output format for the -m or -v options
2012     with hexadecimal format.
2013     -d override the default output format for the -m or -v options
2014     with decimal format.
2015     -f vm_flags translate the bits of a FLAGS (vm_flags) value.
2016     pid a process PID.
2017     task a hexadecimal task_struct pointer.
2018
2019 EXAMPLES
2020     Display the virtual memory data of the current context:
2021
2022 crash> vm
2023     PID: 30986 TASK: c0440000 CPU: 0 COMMAND: "bash"
2024     MM PGD RSS TOTAL_VM
2025     c303fe20 c4789000 88k 1728k
2026     VMA START END FLAGS FILE
2027     c0d1f540 8048000 80ad000 1875 /bin/bash
2028     c0d1f400 80ad000 80b3000 1873 /bin/bash
2029     c0d1f880 80b3000 80ec000 77
2030     c0d1f0c0 40000000 40012000 875 /lib/ld-2.1.1.so
2031     c0d1f700 40012000 40013000 873 /lib/ld-2.1.1.so
2032     c0d1fe00 40013000 40014000 77
2033     c0d1f580 40014000 40016000 73
2034     c0d1f280 4001a000 4004b000 75 /usr/lib/libncurses.so.4.2
2035     c0d1f100 4004b000 40054000 73 /usr/lib/libncurses.so.4.2
2036     c0d1f600 40054000 40057000 73
2037     c0d1f9c0 40057000 40059000 75 /lib/libdl-2.1.1.so
2038     c0d1f800 40059000 4005a000 73 /lib/libdl-2.1.1.so

```

2039	c0d1fd00	4005a000	40140000	75	/lib/libc-2.1.1.so
2040	c0d1fe40	40140000	40145000	73	/lib/libc-2.1.1.so
2041	c0d1f780	40145000	40148000	73	
2042	c0d1f140	40148000	40150000	75	/lib/libnss_files-2.1.1.so
2043	c0d1fa80	40150000	40151000	73	/lib/libnss_files-2.1.1.so
2044	c0d1fb00	40151000	4015a000	75	/lib/libnss_nisplus-2.1.1.so
2045	c5f754e0	4015a000	4015b000	73	/lib/libnss_nisplus-2.1.1.so
2046	c0d1fec0	4015b000	4016d000	75	/lib/libnsl-2.1.1.so
2047	c5f75460	4016d000	4016e000	73	/lib/libnsl-2.1.1.so
2048	c5f75420	4016e000	40170000	73	
2049	c5f753e0	40170000	40178000	75	/lib/libnss_nis-2.1.1.so
2050	c5f753a0	40178000	40179000	73	/lib/libnss_nis-2.1.1.so
2051	c0d1f240	bffffc00	c0000000	177	

2052

2053 Display the virtual memory data along with page translations for PID 386:

```

2054
2055 crash> vm -p 386
2056 PID: 386      TASK: c11cc000  CPU: 0    COMMAND: "atd"
2057   MM          PGD          RSS      TOTAL_VM
2058 c7e30560     c10e5000      104k      1112k
2059   VMA          START        END      FLAGS  FILE
2060 c0fbe6a0      8048000      804b000   1875   /usr/sbin/atd
2061 VIRTUAL      PHYSICAL
2062 8048000      20e1000
2063 8049000      17c6000
2064 804a000      1f6f000
2065   VMA          START        END      FLAGS  FILE
2066 c61e0ba0      804b000      804d000   1873   /usr/sbin/atd
2067 VIRTUAL      PHYSICAL
2068 804b000      254d000
2069 804c000      6a9c000
2070   VMA          START        END      FLAGS  FILE
2071 c61e04e0      804d000      8050000    77
2072 VIRTUAL      PHYSICAL
2073 804d000      219d000
2074 804e000      2617000
2075 804f000      SWAP: /dev/sda8  OFFSET: 24225
2076   VMA          START        END      FLAGS  FILE
2077 c61e0720     40000000     40012000    875   /lib/ld-2.1.1.so
2078 VIRTUAL      PHYSICAL
2079 40000000     FILE: /lib/ld-2.1.1.so  OFFSET: 0
2080 40001000     FILE: /lib/ld-2.1.1.so  OFFSET: 1000
2081 40002000     FILE: /lib/ld-2.1.1.so  OFFSET: 2000
2082 40003000     FILE: /lib/ld-2.1.1.so  OFFSET: 3000
2083 40004000     FILE: /lib/ld-2.1.1.so  OFFSET: 4000
2084 40005000     FILE: /lib/ld-2.1.1.so  OFFSET: 5000
2085 ...
2086

```

2087 Although the -R option is typically invoked from "foreach vm", it can be

2088 executed directly. This example displays all VM areas with vm_flags of 75:

```

2089
2090 crash> vm -R 75
2091 PID: 694      TASK: c0c76000  CPU: 1    COMMAND: "crash"
2092   MM          PGD          RSS      TOTAL_VM
2093 c6c43110     c0fe9000      8932k      10720k
2094   VMA          START        END      FLAGS  FILE
2095 c322c0d0     40019000     4004a000    75   /usr/lib/libncurses.so.4.2
2096 c67537c0     40056000     40071000    75   /lib/libm-2.1.1.so
2097 c6753d00     40072000     40074000    75   /lib/libdl-2.1.1.so
2098 c6753540     40075000     40081000    75   /usr/lib/libz.so.1.1.3
2099 c6753740     40085000     4016b000    75   /lib/libc-2.1.1.so
2100

```

2101 One reason to use -R directly is to pare down the output associated with

2102 the -p option on a task with a huge address space. This example displays

2103 the page data associated with virtual address 40121000:

```

2104
2105 crash> vm -R 40121000
2106 PID: 694      TASK: c0c76000  CPU: 0    COMMAND: "crash"
2107   MM          PGD          RSS      TOTAL_VM
2108 c6c43110     c0fe9000      8928k      10720k
2109   VMA          START        END      FLAGS  FILE
2110 c6753740     40085000     4016b000    75   /lib/libc-2.1.1.so
2111 VIRTUAL      PHYSICAL

```

2112 40121000 FILE: /lib/libc-2.1.1.so OFFSET: 9c000

2113

2114 Display the mm_struct for PID 4777:

2115

2116 crash> vm -m 4777

2117 PID: 4777 TASK: c0896000 CPU: 0 COMMAND: "bash"

2118 struct mm_struct {

2119 mmap = 0xc6caalc0,

2120 mmap_avl = 0x0,

2121 mmap_cache = 0xc6caabc0,

2122 pgd = 0xc100a000,

2123 count = {

2124 counter = 0x1

2125 },

2126 map_count = 0x14,

2127 mmap_sem = {

2128 count = {

2129 counter = 0x1

2130 },

2131 waking = 0x0,

2132 wait = 0x0

2133 },

2134 context = 0x0,

2135 start_code = 0x8048000,

2136 end_code = 0x809c6f7,

2137 start_data = 0x0,

2138 end_data = 0x80a2090,

2139 start_brk = 0x80a5420,

2140 brk = 0x80b9000,

2141 start_stack = 0xbffff9d0,

2142 arg_start = 0xbffffad1,

2143 arg_end = 0xbffffad7,

2144 env_start = 0xbffffad7,

2145 env_end = 0xbffffbf2,

2146 rss = 0xf6,

2147 total_vm = 0x1a3,

2148 locked_vm = 0x0,

2149 def_flags = 0x0,

2150 cpu_vm_mask = 0x0,

2151 swap_cnt = 0x23d,

2152 swap_address = 0x0,

2153 segments = 0x0

2154 }

2155

2156 Display all of the vm_area_structs for task c47d4000:

2157

2158 crash> vm -v c47d4000

2159 PID: 4971 TASK: c47d4000 CPU: 1 COMMAND: "login"

2160 struct vm_area_struct {

2161 vm_mm = 0xc4b0d200,

2162 vm_start = 0x8048000,

2163 vm_end = 0x804d000,

2164 vm_next = 0xc3e3abd0,

2165 vm_page_prot = {

2166 pgprot = 0x25

2167 },

2168 vm_flags = 0x1875,

2169 vm_avl_height = 0x1,

2170 vm_avl_left = 0x0,

2171 vm_avl_right = 0x0,

2172 vm_next_share = 0x0,

2173 vm_pprev_share = 0xc3e3abf0,

2174 vm_ops = 0xc02392a0,

2175 vm_offset = 0x0,

2176 vm_file = 0xc1e23660,

2177 vm_pte = 0x0

2178 }

2179 struct vm_area_struct {

2180 vm_mm = 0xc4b0d200,

2181 vm_start = 0x804d000,

2182 vm_end = 0x804e000,

2183 vm_next = 0xc3e3a010,

2184 vm_page_prot = {

```

2185         pgprot = 0x25
2186     },
2187     vm_flags = 0x1873,
2188     vm_avl_height = 0x2,
2189     vm_avl_left = 0xc3e3a810,
2190     vm_avl_right = 0xc3e3a010,
2191     vm_next_share = 0xc3e3a810,
2192     vm_pprev_share = 0xc3699c14
2193     ...
2194

```

Translate a FLAGS value:

```

2196
2197     crash> vm -f 3875
2198     3875: (READ|EXEC|MAYREAD|MAYWRITE|MAYEXEC|DENYWRITE|EXECUTABLE|LOCKED)
2199

```

Display the page translations of the VM area at address f5604f2c:

```

2200
2201
2202     crash> vm -P f5604f2c
2203     PID: 5508      TASK: f56a9570  CPU: 0      COMMAND: "crond"
2204     VMA          START      END      FLAGS    FILE
2205     f5604f2c      f5b000      f67000  8000075  /lib/libnss_files-2.12.so
2206     VIRTUAL      PHYSICAL
2207     f5b000        3fec1000
2208     f5c000        3d3a4000
2209     f5d000        FILE: /lib/libnss_files-2.12.so  OFFSET: 2000
2210     f5e000        FILE: /lib/libnss_files-2.12.so  OFFSET: 3000
2211     f5f000        FILE: /lib/libnss_files-2.12.so  OFFSET: 4000
2212     f60000        3fd31000
2213     f61000        3fd32000
2214     f62000        FILE: /lib/libnss_files-2.12.so  OFFSET: 7000
2215     f63000        FILE: /lib/libnss_files-2.12.so  OFFSET: 8000
2216     f64000        3ff35000
2217     f65000        FILE: /lib/libnss_files-2.12.so  OFFSET: a000
2218     f66000        FILE: /lib/libnss_files-2.12.so  OFFSET: b000
2219

```

NAME

bpf - extended Berkeley Packet Filter (eBPF)

SYNOPSIS

bpf [[-p ID | -P] [-tTj]] [[-m ID] | -M] [-s] [-xd]

DESCRIPTION

This command provides information on currently-loaded eBPF programs and maps. With no arguments, basic information about each loaded eBPF program and map is displayed. For each eBPF program, its ID number, the addresses of its bpf_prog and bpf_prog_aux data structures, its type, tag, and the IDs of the eBPF maps that it uses are displayed. For each eBPF map, its ID number, the address of its bpf_map data structure, its type, and the hexadecimal value of its map_flags are displayed.

```

2237     -p ID    displays the basic information specific to the program ID, plus the
2238              size in bytes of its translated bytecode, the size in bytes of its
2239              jited code, the number of bytes locked into memory, the time that
2240              the program was loaded, whether it is GPL compatible, its name
2241              string, and its UID.
2242     -P      same as -p, but displays the basic and extra data for all programs.
2243     -m ID    displays the basic information specific to the map ID, plus the
2244              size in bytes of its key and value, the maximum number of key-value
2245              pairs that can be stored within the map, the number of bytes locked
2246              into memory, its name string, and its UID.
2247     -M      same as -m, but displays the basic and extra data for all maps.
2248     -t      translate the bytecode of the specified program ID.
2249     -T      same as -t, but also dump the bytecode of each instruction.
2250     -j      disassemble the jited code of the specified program ID.
2251     -s      with -p or -P, dump the bpf_prog and bpf_prog_aux datastructures.
2252              with -m or -M, dump the bpf_map structure.
2253     -x      with -s, override default output format with hexadecimal format.
2254     -d      with -s, override default output format with decimal format.
2255

```

EXAMPLES

Display all loaded eBPF programs and maps:

```

2258
2259 crash> bpf
2260 ID      BPF_PROG      BPF_PROG_AUX      BPF_PROG_TYPE      TAG      USED_MAPS
2261 13 fffffbc00c06d1000 ffff9ff260f0c400 CGROUP_SKB 7be49e3934a125ba 13,14
2262 14 fffffbc00c0761000 ffff9ff260f0f600 CGROUP_SKB 2a142ef67aaad174 13,14
2263 15 fffffbc00c001d000 ffff9ff2618f9e00 CGROUP_SKB 7be49e3934a125ba 15,16
2264 16 fffffbc00c06c9000 ffff9ff2618f9400 CGROUP_SKB 2a142ef67aaad174 15,16
2265 19 fffffbc00c0d39000 ffff9ff2610fa000 CGROUP_SKB 7be49e3934a125ba 19,20
2266 20 fffffbc00c0d41000 ffff9ff2610f8e00 CGROUP_SKB 2a142ef67aaad174 19,20
2267 30 fffffbc00c065f000 ffff9ff1b64de200 KPROBE 69fed6de18629d7a 32
2268 31 fffffbc00c065b000 ffff9ff1b64df200 KPROBE 69fed6de18629d7a 37
2269 32 fffffbc00c0733000 ffff9ff1b64dc600 KPROBE 69fed6de18629d7a 38
2270 33 fffffbc00c0735000 ffff9ff1b64dca00 KPROBE 69fed6de18629d7a 39
2271 34 fffffbc00c0737000 ffff9ff1b64dfc00 KPROBE 4abbdade72a6ee17 33,36,34
2272 36 fffffbc00c0839000 ffff9ff1b64dd000 KPROBE da4fc6a3f41761a2 32
2273 41 fffffbc00c07ec000 ffff9ff207b70400 TRACEPOINT e2094f9f46284bf6 55,54
2274 44 fffffbc00c07ee000 ffff9ff1b64dc800 PERF_EVENT 19578a12836c4115 62
2275 46 fffffbc00c07f0000 ffff9ff207b70400 SOCKET_FILTER 1fcfc04afd689133 64
2276
2277 ID      BPF_MAP      BPF_MAP_TYPE      MAP_FLAGS
2278 13 ffff9ff260f0ec00 LPM_TRIE 00000001
2279 14 ffff9ff260f0de00 LPM_TRIE 00000001
2280 15 ffff9ff2618fbe00 LPM_TRIE 00000001
2281 16 ffff9ff2618fb800 LPM_TRIE 00000001
2282 19 ffff9ff2610faa00 LPM_TRIE 00000001
2283 20 ffff9ff2610fb800 LPM_TRIE 00000001
2284 32 ffff9ff260d74000 HASH 00000000
2285 33 ffff9ff260d76400 LRU_HASH 00000000
2286 34 ffff9ff260d70000 LRU_HASH 00000002
2287 35 ffff9ff260d73800 LRU_HASH 00000004
2288 36 ffff9ff1b4f44000 ARRAY_OF_MAPS 00000000
2289 37 ffff9ff260d77c00 PERCPU_HASH 00000000
2290 38 ffff9ff260d70800 HASH 00000001
2291 39 ffff9ff260d76c00 PERCPU_HASH 00000001
2292 54 ffff9ff260dd2c00 HASH 00000000
2293 55 ffff9ff260dd1400 HASH 00000000
2294 62 ffff9ff1ae784000 HASH 00000000
2295 64 ffff9ff1aeal5000 ARRAY 00000000
2296
2297 Display additional data about program ID 20:
2298
2299 crash> bpf -p 20
2300 ID      BPF_PROG      BPF_PROG_AUX      BPF_PROG_TYPE      TAG      USED_MAPS
2301 20 fffffbc00c0d41000 ffff9ff2610f8e00 CGROUP_SKB 2a142ef67aaad174 19,20
2302 XLATED: 296 JITED: 229 MEMLOCK: 4096
2303 LOAD_TIME: Fri Apr 20 19:39:21 2018
2304 GPL_COMPATIBLE: yes UID: 0
2305
2306 Display additional data about map ID 34:
2307
2308 crash> bpf -m 34
2309 ID      BPF_MAP      BPF_MAP_TYPE      MAP_FLAGS
2310 34 ffff9ff260d70000 LRU_HASH 00000000
2311 KEY_SIZE: 4 VALUE_SIZE: 8 MAX_ENTRIES: 10000 MEMLOCK: 1953792
2312 NAME: "lru_hash_map" UID: 0
2313
2314 Disassemble the jited program of program ID 20:
2315
2316 crash> bpf -p 20 -j
2317 ID      BPF_PROG      BPF_PROG_AUX      BPF_PROG_TYPE      TAG      USED_MAPS
2318 20 fffffbc00c0d41000 ffff9ff2610f8e00 CGROUP_SKB 2a142ef67aaad174 19,20
2319 XLATED: 296 JITED: 229 MEMLOCK: 4096
2320 LOAD_TIME: Fri Apr 20 19:39:21 2018
2321 GPL_COMPATIBLE: yes UID: 0
2322
2323 0xfffffffffc06887a2: push %rbp
2324 0xfffffffffc06887a3: mov %rsp,%rbp
2325 0xfffffffffc06887a6: sub $0x40,%rsp
2326 0xfffffffffc06887ad: sub $0x28,%rbp
2327 0xfffffffffc06887b1: mov %rbx,0x0(%rbp)
2328 0xfffffffffc06887b5: mov %r13,0x8(%rbp)
2329 0xfffffffffc06887b9: mov %r14,0x10(%rbp)
2330 0xfffffffffc06887bd: mov %r15,0x18(%rbp)

```



```

2331 0xfffffffffc06887c1: xor    %eax,%eax
2332 0xfffffffffc06887c3: mov    %rax,0x20(%rbp)
2333 0xfffffffffc06887c7: mov    %rdi,%rbx
2334 0xfffffffffc06887ca: movzwb 0xc0(%rbx),%r13
2335 0xfffffffffc06887d2: xor    %r14d,%r14d
2336 0xfffffffffc06887d5: cmp    $0x8,%r13
2337 0xfffffffffc06887d9: jne    0xfffffffffc068881b
2338 0xfffffffffc06887db: mov    %rbx,%rdi
2339 0xfffffffffc06887de: mov    $0xc,%esi
2340 0xfffffffffc06887e3: mov    %rbp,%rdx
2341 0xfffffffffc06887e6: add    $0xfffffffffffffffffc,%rdx
2342 0xfffffffffc06887ea: mov    $0x4,%ecx
2343 0xfffffffffc06887ef: callq  0xfffffffffb0865340 <bpf_skb_load_bytes>
2344 0xfffffffffc06887f4: movabs $0xffff9ff2610faa00,%rdi
2345 0xfffffffffc06887fe: mov    %rbp,%rsi
2346 0xfffffffffc0688801: add    $0xfffffffffffffffff8,%rsi
2347 0xfffffffffc0688805: movl   $0x20,0x0(%rsi)
2348 0xfffffffffc068880c: callq  0xfffffffffb01fcba0 <bpf_map_lookup_elem>
2349 0xfffffffffc0688811: cmp    $0x0,%rax
2350 0xfffffffffc0688815: je     0xfffffffffc068881b
2351 0xfffffffffc0688817: or     $0x2,%r14d
2352 0xfffffffffc068881b: cmp    $0xdd86,%r13
2353 0xfffffffffc0688822: jne    0xfffffffffc0688864
2354 0xfffffffffc0688824: mov    %rbx,%rdi
2355 0xfffffffffc0688827: mov    $0x8,%esi
2356 0xfffffffffc068882c: mov    %rbp,%rdx
2357 0xfffffffffc068882f: add    $0xfffffffffffffffff0,%rdx
2358 0xfffffffffc0688833: mov    $0x10,%ecx
2359 0xfffffffffc0688838: callq  0xfffffffffb0865340 <bpf_skb_load_bytes>
2360 0xfffffffffc068883d: movabs $0xffff9ff2610fb800,%rdi
2361 0xfffffffffc0688847: mov    %rbp,%rsi
2362 0xfffffffffc068884a: add    $0xfffffffffffffec,%rsi
2363 0xfffffffffc068884e: movl   $0x80,0x0(%rsi)
2364 0xfffffffffc0688855: callq  0xfffffffffb01fcba0 <bpf_map_lookup_elem>
2365 0xfffffffffc068885a: cmp    $0x0,%rax
2366 0xfffffffffc068885e: je     0xfffffffffc0688864
2367 0xfffffffffc0688860: or     $0x2,%r14d
2368 0xfffffffffc0688864: mov    $0x1,%eax
2369 0xfffffffffc0688869: cmp    $0x2,%r14
2370 0xfffffffffc068886d: jne    0xfffffffffc0688871
2371 0xfffffffffc068886f: xor    %eax,%eax
2372 0xfffffffffc0688871: mov    0x0(%rbp),%rbx
2373 0xfffffffffc0688875: mov    0x8(%rbp),%r13
2374 0xfffffffffc0688879: mov    0x10(%rbp),%r14
2375 0xfffffffffc068887d: mov    0x18(%rbp),%r15
2376 0xfffffffffc0688881: add    $0x28,%rbp
2377 0xfffffffffc0688885: leaveq
2378 0xfffffffffc0688886: retq

```

Translate each bytecode instruction of program ID 13:

```

2380
2381
2382 crash> bpf -p 13 -t
2383 ID      BPF_PROG      BPF_PROG_AUX      BPF_PROG_TYPE      TAG      USED_MAPS
2384 13 ffffb00c06d1000 ffff9ff260f0c400 CGROUP_SKB 7be49e3934a125ba 13,14
2385 XLATED: 296 JITED: 229 MEMLOCK: 4096
2386 LOAD_TIME: Fri Apr 20 19:39:11 2018
2387 GPL_COMPATIBLE: yes UID: 0
2388
2389 0: (bf) r6 = r1
2390 1: (69) r7 = *(u16 *) (r6 +192)
2391 2: (b4) (u32) r8 = (u32) 0
2392 3: (55) if r7 != 0x8 goto pc+14
2393 4: (bf) r1 = r6
2394 5: (b4) (u32) r2 = (u32) 16
2395 6: (bf) r3 = r10
2396 7: (07) r3 += -4
2397 8: (b4) (u32) r4 = (u32) 4
2398 9: (85) call bpf_skb_load_bytes#6793152
2399 10: (18) r1 = map[id:13]
2400 12: (bf) r2 = r10
2401 13: (07) r2 += -8
2402 14: (62) *(u32 *) (r2 +0) = 32
2403 15: (85) call bpf_map_lookup_elem#73760

```

```

2404 16: (15) if r0 == 0x0 goto pc+1
2405 17: (44) (u32) r8 |= (u32) 2
2406 18: (55) if r7 != 0xdd86 goto pc+14
2407 19: (bf) r1 = r6
2408 20: (b4) (u32) r2 = (u32) 24
2409 21: (bf) r3 = r10
2410 22: (07) r3 += -16
2411 23: (b4) (u32) r4 = (u32) 16
2412 24: (85) call bpf_skb_load_bytes#6793152
2413 25: (18) r1 = map[id:14]
2414 27: (bf) r2 = r10
2415 28: (07) r2 += -20
2416 29: (62) *(u32 *) (r2 +0) = 128
2417 30: (85) call bpf_map_lookup_elem#73760
2418 31: (15) if r0 == 0x0 goto pc+1
2419 32: (44) (u32) r8 |= (u32) 2
2420 33: (b7) r0 = 1
2421 34: (55) if r8 != 0x2 goto pc+1
2422 35: (b7) r0 = 0
2423 36: (95) exit
2424

```

Translate, and then dump each bytecode instruction of program ID 13:

```

2425 crash> bpf -p 13 -T
2426
2427 ID          BPF_PROG          BPF_PROG_AUX          BPF_PROG_TYPE          TAG          USED_MAPS
2428 13 fffffbc00c06d1000 fffff9ff260f0c400 CGROUP_SKB 7be49e3934a125ba 13,14
2429 XLATED: 296 JITED: 229 MEMLOCK: 4096
2430 LOAD_TIME: Fri Apr 20 19:39:11 2018
2431 GPL_COMPATIBLE: yes UID: 0
2432
2433 0: (bf) r6 = r1
2434      bf 16 00 00 00 00 00 00
2435 1: (69) r7 = *(u16 *) (r6 +192)
2436      69 67 c0 00 00 00 00 00
2437 2: (b4) (u32) r8 = (u32) 0
2438      b4 08 00 00 00 00 00 00
2439 3: (55) if r7 != 0x8 goto pc+14
2440      55 07 0e 00 08 00 00 00
2441 4: (bf) r1 = r6
2442      bf 61 00 00 00 00 00 00
2443 5: (b4) (u32) r2 = (u32) 16
2444      b4 02 00 00 10 00 00 00
2445 6: (bf) r3 = r10
2446      bf a3 00 00 00 00 00 00
2447 7: (07) r3 += -4
2448      07 03 00 00 fc ff ff ff
2449 8: (b4) (u32) r4 = (u32) 4
2450      b4 04 00 00 04 00 00 00
2451 9: (85) call bpf_skb_load_bytes#6793152
2452      85 00 00 00 c0 a7 67 00
2453 10: (18) r1 = map[id:13]
2454      18 01 00 00 00 7a 96 61 00 00 00 00 b2 9d ff ff
2455 12: (bf) r2 = r10
2456      bf a2 00 00 00 00 00 00
2457 13: (07) r2 += -8
2458      07 02 00 00 f8 ff ff ff
2459 14: (62) *(u32 *) (r2 +0) = 32
2460      62 02 00 00 20 00 00 00
2461 15: (85) call bpf_map_lookup_elem#73760
2462      85 00 00 00 20 20 01 00
2463 16: (15) if r0 == 0x0 goto pc+1
2464      15 00 01 00 00 00 00 00
2465 17: (44) (u32) r8 |= (u32) 2
2466      44 08 00 00 02 00 00 00
2467 18: (55) if r7 != 0xdd86 goto pc+14
2468      55 07 0e 00 86 dd 00 00
2469 19: (bf) r1 = r6
2470      bf 61 00 00 00 00 00 00
2471 20: (b4) (u32) r2 = (u32) 24
2472      b4 02 00 00 18 00 00 00
2473 21: (bf) r3 = r10
2474      bf a3 00 00 00 00 00 00
2475 22: (07) r3 += -16
2476

```

```

2477         07 03 00 00 f0 ff ff ff
2478 23: (b4) (u32) r4 = (u32) 16
2479         b4 04 00 00 10 00 00 00
2480 24: (85) call bpf_skb_load_bytes#6793152
2481         85 00 00 00 c0 a7 67 00
2482 25: (18) r1 = map[id:14]
2483         18 01 00 00 00 68 96 61 00 00 00 00 b2 9d ff ff
2484 27: (bf) r2 = r10
2485         bf a2 00 00 00 00 00 00
2486 28: (07) r2 += -20
2487         07 02 00 00 ec ff ff ff
2488 29: (62) *(u32 *) (r2 +0) = 128
2489         62 02 00 00 80 00 00 00
2490 30: (85) call bpf_map_lookup_elem#73760
2491         85 00 00 00 20 20 01 00
2492 31: (15) if r0 == 0x0 goto pc+1
2493         15 00 01 00 00 00 00 00
2494 32: (44) (u32) r8 |= (u32) 2
2495         44 08 00 00 02 00 00 00
2496 33: (b7) r0 = 1
2497         b7 00 00 00 01 00 00 00
2498 34: (55) if r8 != 0x2 goto pc+1
2499         55 08 01 00 02 00 00 00
2500 35: (b7) r0 = 0
2501         b7 00 00 00 00 00 00 00
2502 36: (95) exit
2503         95 00 00 00 00 00 00 00
2504

```

Display the bpf_map data structure for map ID 13:

```

2505
2506
2507 crash> bpf -m 13 -s
2508 ID          BPF_MAP          BPF_MAP_TYPE      MAP_FLAGS
2509 13  ffff9ff260f0ec00      LPM_TRIE          00000001
2510      KEY_SIZE: 8  VALUE_SIZE: 8  MAX_ENTRIES: 1  MEMLOCK: 4096
2511      NAME: (unused)  UID: 0
2512

```

```

2513 struct bpf_map {
2514     ops = 0xfffffffffb0e36720,
2515     inner_map_meta = 0x0,
2516     security = 0xffff9ff26873a158,
2517     map_type = BPF_MAP_TYPE_LPM_TRIE,
2518     key_size = 8,
2519     value_size = 8,
2520     max_entries = 1,
2521     map_flags = 1,
2522     pages = 1,
2523     id = 13,
2524     numa_node = -1,
2525     unpriv_array = false,
2526     user = 0xfffffffffb14578a0,
2527     refcnt = {
2528         counter = 3
2529     },
2530     usercnt = {
2531         counter = 1
2532     },
2533     work = {
2534         data = {
2535             counter = 0
2536         },
2537         entry = {
2538             next = 0x0,
2539             prev = 0x0
2540         },
2541         func = 0x0,
2542         lockdep_map = {
2543             key = 0x0,
2544             class_cache = {0x0, 0x0},
2545             name = 0x0,
2546             cpu = 0,
2547             ip = 0
2548         }
2549     },

```

```

2550     name = "
2551 }
2552
2553 Display the bpf_prog and bpf_prog_aux structures for program ID 13:
2554
2555 crash> bpf -p 13 -s
2556 ID      BPF_PROG      BPF_PROG_AUX      BPF_PROG_TYPE      TAG      USED_MAPS
2557 13      fffffbc00c06d1000 fffff9ff260f0c400  CGROUP_SKB      7be49e3934a125ba      13,14
2558      XLATED: 296 JITED: 229 MEMLOCK: 4096
2559      LOAD_TIME: Fri Apr 20 19:39:10 2018
2560      GPL_COMPATIBLE: yes UID: 0
2561
2562 struct bpf_prog {
2563     pages = 1,
2564     jited = 1,
2565     jit_requested = 1,
2566     locked = 1,
2567     gpl_compatible = 1,
2568     cb_access = 0,
2569     dst_needed = 0,
2570     blinded = 0,
2571     is_func = 0,
2572     kprobe_override = 0,
2573     type = BPF_PROG_TYPE_CGROUP_SKB,
2574     len = 37,
2575     jited_len = 229,
2576     tag = "{\344\236\071\064\241%\272",
2577     aux = fffff9ff260f0c400,
2578     orig_prog = 0x0,
2579     bpf_func = 0xffffffffc0218a59,
2580     {
2581         insns = 0xfffffb0cf406d1030,
2582         insnsi = 0xfffffb0cf406d1030
2583     }
2584 }
2585
2586 struct bpf_prog_aux {
2587     refcnt = {
2588         counter = 2
2589     },
2590     used_map_cnt = 2,
2591     max_ctx_offset = 20,
2592     stack_depth = 20,
2593     id = 13,
2594     func_cnt = 0,
2595     offload_requested = false,
2596     func = 0x0,
2597     jit_data = 0x0,
2598     ksym_tnode = {
2599         node = {{
2600             _rb_parent_color = 18446635988194065457,
2601             rb_right = 0x0,
2602             rb_left = 0x0
2603         }, {
2604             _rb_parent_color = 18446635988194065481,
2605             rb_right = 0x0,
2606             rb_left = 0x0
2607         }}
2608     },
2609     ksym_lnode = {
2610         next = 0xfffff9db261966460,
2611         prev = 0xfffffffffb85d1150
2612     },
2613     ops = 0xfffffffffb7f09060,
2614     used_maps = 0xfffff9db261e03600,
2615     prog = 0xfffffb0cf406d1000,
2616     user = 0xfffffffffb84578a0,
2617     load_time = 23962237943,
2618     name = "
2619     security = 0xfffff9db266f9cf50,
2620     offload = 0x0,
2621     {
2622         work = {

```

```

2623     data = {
2624         counter = 0
2625     },
2626     entry = {
2627         next = 0x0,
2628         prev = 0x0
2629     },
2630     func = 0x0,
2631     lockdep_map = {
2632         key = 0x0,
2633         class_cache = {0x0, 0x0},
2634         name = 0x0,
2635         cpu = 0,
2636         ip = 0
2637     }
2638 },
2639     rcu = {
2640         next = 0x0,
2641         func = 0x0
2642     }
2643 }
2644 }
2645

```

Display the extra data about all programs:

```
crash> bpf -P
```

ID	BPF_PROG	BPF_PROG_AUX	BPF_PROG_TYPE	TAG	USED_MAPS
13	ffffbc00c06d1000	ffff9ff260f0c400	CGROUP_SKB	7be49e3934a125ba	13,14
XLATED: 296 JITED: 229 MEMLOCK: 4096					
LOAD_TIME: Fri Apr 20 19:39:10 2018					
GPL_COMPATIBLE: yes UID: 0					

ID	BPF_PROG	BPF_PROG_AUX	BPF_PROG_TYPE	TAG	USED_MAPS
14	ffffbc00c0761000	ffff9ff260f0f600	CGROUP_SKB	2a142ef67aaad174	13,14
XLATED: 296 JITED: 229 MEMLOCK: 4096					
LOAD_TIME: Fri Apr 20 19:39:10 2018					
GPL_COMPATIBLE: yes UID: 0					

ID	BPF_PROG	BPF_PROG_AUX	BPF_PROG_TYPE	TAG	USED_MAPS
15	ffffbc00c001d000	ffff9ff2618f9e00	CGROUP_SKB	7be49e3934a125ba	15,16
XLATED: 296 JITED: 229 MEMLOCK: 4096					
LOAD_TIME: Fri Apr 20 19:39:11 2018					
GPL_COMPATIBLE: yes UID: 0					

...

ID	BPF_PROG	BPF_PROG_AUX	BPF_PROG_TYPE	TAG	USED_MAPS
75	ffffbc00c0ed1000	ffff9ff2429c6400	KPROBE	da4fc6a3f41761a2	107
XLATED: 5168 JITED: 2828 MEMLOCK: 8192					
LOAD_TIME: Fri Apr 27 14:54:40 2018					
GPL_COMPATIBLE: yes UID: 0					

Display the extra data for all maps:

```
crash> bpf -M
```

ID	BPF_MAP	BPF_MAP_TYPE	MAP_FLAGS
13	ffff9ff260f0ec00	LPM_TRIE	00000001
KEY_SIZE: 8 VALUE_SIZE: 8 MAX_ENTRIES: 1 MEMLOCK: 4096			
NAME: (unused) UID: 0			

ID	BPF_MAP	BPF_MAP_TYPE	MAP_FLAGS
14	ffff9ff260f0de00	LPM_TRIE	00000001
KEY_SIZE: 20 VALUE_SIZE: 8 MAX_ENTRIES: 1 MEMLOCK: 4096			
NAME: (unused) UID: 0			

...

ID	BPF_MAP	BPF_MAP_TYPE	MAP_FLAGS
108	ffff9ff1aeab9400	LRU_HASH	00000000
KEY_SIZE: 4 VALUE_SIZE: 8 MAX_ENTRIES: 1000 MEMLOCK: 147456			
NAME: "lru_hash_lookup" UID: 0			

To display all possible information that this command offers about

```

2696     all programs and maps, enter:
2697
2698     crash> bpf -PM -jTs
2699
2700
2701 NAME
2702     gdb - gdb command
2703
2704 SYNOPSIS
2705     gdb command ...
2706
2707 DESCRIPTION
2708     This command passes its arguments directly to gdb for processing.
2709     This is typically not necessary, but where ambiguities between crash and
2710     gdb command names exist, this will force the command to be executed by gdb.
2711
2712     Alternatively, if "set gdb on" is entered, the session will be run in a
2713     mode where all commands are passed directly to gdb. When running in that
2714     mode, native crash commands may be executed by preceding them with the
2715     "crash" directive. To restore native crash mode, enter "set gdb off".
2716
2717 EXAMPLES
2718     crash> gdb help
2719     List of classes of commands:
2720
2721     aliases -- Aliases of other commands
2722     breakpoints -- Making program stop at certain points
2723     data -- Examining data
2724     files -- Specifying and examining files
2725     internals -- Maintenance commands
2726     obscure -- Obscure features
2727     running -- Running the program
2728     stack -- Examining the stack
2729     status -- Status inquiries
2730     support -- Support facilities
2731     tracepoints -- Tracing of program execution without stopping the program
2732     user-defined -- User-defined commands
2733
2734     Type "help" followed by a class name for a list of commands in that class.
2735     Type "help" followed by command name for full documentation.
2736     Command name abbreviations are allowed if unambiguous.
2737
2738
2739 NAME
2740     p - print the value of an expression
2741
2742 SYNOPSIS
2743     p [-x|-d][-u] [expression | symbol[:cpuspec]]
2744
2745 DESCRIPTION
2746     This command passes its arguments on to gdb "print" command for evaluation.
2747
2748     expression    an expression to be evaluated.
2749     symbol        a kernel symbol.
2750     :cpuspec      CPU specification for a per-cpu symbol:
2751         :          CPU of the currently selected task.
2752         :a[ll]     all CPUs.
2753         :#[-#][,...] CPU list(s), e.g. "1,3,5", "1-3",
2754                     or "1,3,5-7,10".
2755     -x            override default output format with hexadecimal format.
2756     -d            override default output format with decimal format.
2757     -u            the expression evaluates to a user address reference.
2758
2759     The default output format is decimal, but that can be changed at any time
2760     with the two built-in aliases "hex" and "dec". Alternatively, there
2761     are two other built-in aliases, "px" and "pd", which force the command
2762     output to be displayed in hexadecimal or decimal, without changing the
2763     default mode.
2764
2765 EXAMPLES
2766     Print the contents of jiffies:
2767
2768     crash> p jiffies

```

```
2769 jiffies = $6 = 166532620
2770 crash> px jiffies
2771 jiffies = $7 = 0x9ed174b
2772 crash> pd jiffies
2773 jiffies = $8 = 166533160
2774
```

Print the contents of the vm_area_struct "init_mm":

```
2775
2776
2777 crash> p init_mm
2778 init_mm = $5 = {
2779     mmap = 0xc022d540,
2780     mmap_avl = 0x0,
2781     mmap_cache = 0x0,
2782     pgd = 0xc0101000,
2783     count = {
2784         counter = 0x6
2785     },
2786     map_count = 0x1,
2787     mmap_sem = {
2788         count = {
2789             counter = 0x1
2790         },
2791         waking = 0x0,
2792         wait = 0x0
2793     },
2794     context = 0x0,
2795     start_code = 0xc0000000,
2796     end_code = 0xc022b4c8,
2797     start_data = 0x0,
2798     end_data = 0xc0250388,
2799     start_brk = 0x0,
2800     brk = 0xc02928d8,
2801     start_stack = 0x0,
2802     arg_start = 0x0,
2803     arg_end = 0x0,
2804     env_start = 0x0,
2805     env_end = 0x0,
2806     rss = 0x0,
2807     total_vm = 0x0,
2808     locked_vm = 0x0,
2809     def_flags = 0x0,
2810     cpu_vm_mask = 0x0,
2811     swap_cnt = 0x0,
2812     swap_address = 0x0,
2813     segments = 0x0
2814 }
2815
```

If a per-cpu symbol is entered as a argument, its data type and all of its per-cpu addresses are displayed:

```
2816
2817
2818
2819 crash> p irq_stat
2820 PER-CPU DATA TYPE:
2821     irq_cpustat_t irq_stat;
2822 PER-CPU ADDRESSES:
2823     [0]: ffff88021e211540
2824     [1]: ffff88021e251540
2825     [2]: ffff88021e291540
2826     [3]: ffff88021e2d1540
2827
```

To display the contents a per-cpu symbol for CPU 1, append a cpu-specifier:

```
2830
2831 crash> p irq_stat:1
2832 per_cpu(irq_stat, 1) = $29 = {
2833     __softirq_pending = 0,
2834     __nmi_count = 209034,
2835     apic_timer_irqs = 597509876,
2836     irq_spurious_count = 0,
2837     icr_read_retry_count = 2,
2838     x86_platform_ipis = 0,
2839     apic_perf_irqs = 209034,
2840     apic_irq_work_irqs = 0,
2841     irq_resched_count = 264922233,

```

```

2842     irq_call_count = 7036692,
2843     irq_tlb_count = 4750442,
2844     irq_thermal_count = 0,
2845     irq_threshold_count = 0
2846 }

```

NAME

sig - task signal handling

SYNOPSIS

sig [[-l] | [-s sigset]] | [-g] [pid | taskp] ...

DESCRIPTION

This command displays signal-handling data of one or more tasks. Multiple task or PID numbers may be entered; if no arguments are entered, the signal handling data of the current context will be displayed. The default display shows:

1. A formatted dump of the "sig" signal_struct structure referenced by the task_struct. For each defined signal, it shows the sigaction structure address, the signal handler, the signal sigset_t mask (also expressed as a 64-bit hexadecimal value), and the flags.
2. Whether the task has an unblocked signal pending.
3. The contents of the "blocked" and "signal" sigset_t structures from the task_struct/signal_struct, both of which are represented as a 64-bit hexadecimal value.
4. For each queued signal, private and/or shared, if any, its signal number and associated siginfo structure address.

The -l option lists the signal numbers and their name(s). The -s option translates a 64-bit hexadecimal value representing the contents of a sigset_t structure into the signal names whose bits are set.

```

pid    a process PID.
taskp  a hexadecimal task_struct pointer.
-g     displays signal information for all threads in a task's
       thread group.
-l     displays the defined signal numbers and names.
-s sigset translates a 64-bit hexadecimal value representing a sigset_t
       into a list of signal names associated with the bits set.

```

EXAMPLES

Dump the signal-handling data of PID 8970:

```

2888 crash> sig 8970
2889 PID: 8970  TASK: f67d8560  CPU: 1   COMMAND: "procsig"
2890 SIGNAL_STRUCT: f6018680  COUNT: 1
2891 SIG SIGACTION  HANDLER      MASK          FLAGS
2892 [1]  f7877684  SIG_DFL 0000000000000000 0
2893 [2]  f7877698  SIG_DFL 0000000000000000 0
2894 ...
2895 [8]  f7877710  SIG_DFL 0000000000000000 0
2896 [9]  f7877724  SIG_DFL 0000000000000000 0
2897 [10] f7877738  804867a 0000000000000000 80000000 (SA_RESETHAND)
2898 [11] f787774c  SIG_DFL 0000000000000000 0
2899 [12] f7877760  804867f 0000000000000000 10000004 (SA_SIGINFO|SA_RESTART)
2900 [13] f7877774  SIG_DFL 0000000000000000 0
2901 ...
2902 [31] f78778dc  SIG_DFL 0000000000000000 0
2903 [32] f78778f0  SIG_DFL 0000000000000000 0
2904 [33] f7877904  SIG_DFL 0000000000000000 0
2905 [34] f7877918  804867f 0000000000000000 10000004 (SA_SIGINFO|SA_RESTART)
2906 [35] f787792c  SIG_DFL 0000000000000000 0
2907 [36] f7877940  SIG_DFL 0000000000000000 0
2908 ...
2909 [58] f7877af8  SIG_DFL 0000000000000000 0
2910 [59] f7877b0c  SIG_DFL 0000000000000000 0
2911 [60] f7877b20  SIG_DFL 0000000000000000 0
2912 [61] f7877b34  SIG_DFL 0000000000000000 0
2913 [62] f7877b48  SIG_DFL 0000000000000000 0
2914 [63] f7877b5c  SIG_DFL 0000000000000000 0

```



```

2915     [64] f7877b70 804867f 0000000000000000 10000004 (SA_SIGINFO|SA_RESTART)
2916 SIGPENDING: no
2917     BLOCKED: 80000002000000800
2918 PRIVATE_PENDING
2919     SIGNAL: 00000002000000800
2920 SIGQUEUE: SIG SIGINFO
2921           12 f51b9c84
2922           34 f51b9594
2923 SHARED_PENDING
2924     SIGNAL: 80000000000000800
2925 SIGQUEUE: SIG SIGINFO
2926           12 f51b9188
2927           64 f51b9d18
2928           64 f51b9500
2929

```

2930 Dump the signal-handling data for all tasks in the thread group containing
 2931 PID 2578:

```

2932
2933 crash> sig -g 2578
2934 PID: 2387 TASK: f617d020 CPU: 0 COMMAND: "slapd"
2935 SIGNAL_STRUCT: f7dede00 COUNT: 6
2936 SIG SIGACTION HANDLER MASK FLAGS
2937 [1] c1f60c04 a258a7 0000000000000000 10000000 (SA_RESTART)
2938 [2] c1f60c18 a258a7 0000000000000000 10000000 (SA_RESTART)
2939 [3] c1f60c2c SIG_DFL 0000000000000000 0
2940 [4] c1f60c40 SIG_DFL 0000000000000000 0
2941 [5] c1f60c54 a258a7 0000000000000000 10000000 (SA_RESTART)
2942 [6] c1f60c68 SIG_DFL 0000000000000000 0
2943 [7] c1f60c7c SIG_DFL 0000000000000000 0
2944 [8] c1f60c90 SIG_DFL 0000000000000000 0
2945 [9] c1f60ca4 SIG_DFL 0000000000000000 0
2946 [10] c1f60cb8 a25911 0000000000000000 10000000 (SA_RESTART)
2947 ...

```

```

2948 [64] c1f610f0 SIG_DFL 0000000000000000 0
2949 SHARED_PENDING

```

```

2950     SIGNAL: 0000000000000000

```

```

2951     SIGQUEUE: (empty)
2952

```

```

2953     PID: 2387 TASK: f617d020 CPU: 0 COMMAND: "slapd"

```

```

2954     SIGPENDING: no

```

```

2955     BLOCKED: 0000000000000000

```

```

2956     PRIVATE_PENDING

```

```

2957     SIGNAL: 0000000000000000

```

```

2958     SIGQUEUE: (empty)
2959

```

```

2960     PID: 2392 TASK: f6175aa0 CPU: 0 COMMAND: "slapd"

```

```

2961     SIGPENDING: no

```

```

2962     BLOCKED: 0000000000000000

```

```

2963     PRIVATE_PENDING

```

```

2964     SIGNAL: 0000000000000000

```

```

2965     SIGQUEUE: (empty)
2966

```

```

2967     PID: 2523 TASK: f7cd4aa0 CPU: 1 COMMAND: "slapd"

```

```

2968     SIGPENDING: no

```

```

2969     BLOCKED: 0000000000000000

```

```

2970     PRIVATE_PENDING

```

```

2971     SIGNAL: 0000000000000000

```

```

2972     SIGQUEUE: (empty)
2973

```

```

2974     ...
2975

```

2976 Translate the sigset_t mask value, cut-and-pasted from the signal handling
 2977 data from signals 1 and 10 above:

```

2978
2979 crash> sig -s 800A0000000000201

```

```

2980     SIGHUP SIGUSR1 SIGRTMAX-14 SIGRTMAX-12 SIGRTMAX
2981

```

2982 List the signal numbers and their names:

```

2983
2984 crash> sig -l

```

```

2985     [1] SIGHUP

```

```

2986     [2] SIGINT

```

```

2987     [3] SIGQUIT

```

```
2988      [4] SIGILL
2989      [5] SIGTRAP
2990      [6] SIGABRT/SIGIOT
2991      [7] SIGBUS
2992      [8] SIGFPE
2993      [9] SIGKILL
2994     [10] SIGUSR1
2995     [11] SIGSEGV
2996     [12] SIGUSR2
2997     [13] SIGPIPE
2998     [14] SIGALRM
2999     [15] SIGTERM
3000     [16] SIGSTKFLT
3001     [17] SIGCHLD/SIGCLD
3002     [18] SIGCONT
3003     [19] SIGSTOP
3004     [20] SIGTSTP
3005     [21] SIGTTIN
3006     [22] SIGTTOU
3007     [23] SIGURG
3008     [24] SIGXCPU
3009     [25] SIGXFSZ
3010     [26] SIGVTALRM
3011     [27] SIGPROF
3012     [28] SIGWINCH
3013     [29] SIGIO/SIGPOLL
3014     [30] SIGPWR
3015     [31] SIGSYS
3016     [32] SIGRTMIN
3017     [33] SIGRTMIN+1
3018     [34] SIGRTMIN+2
3019     [35] SIGRTMIN+3
3020     [36] SIGRTMIN+4
3021     [37] SIGRTMIN+5
3022     [38] SIGRTMIN+6
3023     [39] SIGRTMIN+7
3024     [40] SIGRTMIN+8
3025     [41] SIGRTMIN+9
3026     [42] SIGRTMIN+10
3027     [43] SIGRTMIN+11
3028     [44] SIGRTMIN+12
3029     [45] SIGRTMIN+13
3030     [46] SIGRTMIN+14
3031     [47] SIGRTMIN+15
3032     [48] SIGRTMIN+16
3033     [49] SIGRTMAX-15
3034     [50] SIGRTMAX-14
3035     [51] SIGRTMAX-13
3036     [52] SIGRTMAX-12
3037     [53] SIGRTMAX-11
3038     [54] SIGRTMAX-10
3039     [55] SIGRTMAX-9
3040     [56] SIGRTMAX-8
3041     [57] SIGRTMAX-7
3042     [58] SIGRTMAX-6
3043     [59] SIGRTMAX-5
3044     [60] SIGRTMAX-4
3045     [61] SIGRTMAX-3
3046     [62] SIGRTMAX-2
3047     [63] SIGRTMAX-1
3048     [64] SIGRTMAX
```

```
3049
3050
```

NAME

```
3052     vtop - virtual to physical
3053
```

SYNOPSIS

```
3055     vtop [-c [pid | taskp]] [-u|-k] address ...
3056
```

DESCRIPTION

```
3058     This command translates a user or kernel virtual address to its physical
3059     address. Also displayed is the PTE translation, the vm_area_struct data
3060     for user virtual addresses, the mem_map page data associated with the
```

physical page, and the swap location or file location if the page is not mapped. The -u and -k options specify that the address is a user or kernel virtual address; -u and -k are not necessary on processors whose virtual addresses self-define themselves as user or kernel. User addresses are translated with respect to the current context unless the -c option is used. Kernel virtual addresses are translated using the swapper_pg_dir as the base page directory unless the -c option is used.

-u The address is a user virtual address; only required on processors with overlapping user and kernel virtual address spaces.

-k The address is a kernel virtual address; only required on processors with overlapping user and kernel virtual address spaces.

-c [pid | taskp] Translate the virtual address from the page directory of the specified PID or hexadecimal task_struct pointer. However, if this command is invoked from "foreach vtop", the pid or taskp argument should NOT be entered; the address will be translated using the page directory of each task specified by "foreach".

address A hexadecimal user or kernel virtual address.

EXAMPLES

Translate user virtual address 80b4000:

```
crash> vtop 80b4000
VIRTUAL    PHYSICAL
80b4000    660f000

PAGE DIRECTORY: c37f0000
  PGD: c37f0080 => e0d067
  PMD: c37f0080 => e0d067
  PTE: c0e0d2d0 => 660f067
  PAGE: 660f000

    PTE      PHYSICAL  FLAGS
660f067    660f000  (PRESENT|RW|USER|ACCESSED|DIRTY)

    VMA      START      END      FLAGS  FILE
c773daa0    80b4000    810c000    77

    PAGE      PHYSICAL  INODE      OFFSET  CNT  FLAGS
c0393258    660f000          0      17000  1  uptodate
```

Translate kernel virtual address c806e000, first using swapper_pg_dir as the page directory base, and secondly, using the page table base of PID 1359:

```
crash> vtop c806e000
VIRTUAL    PHYSICAL
c806e000    2216000

PAGE DIRECTORY: c0101000
  PGD: c0101c80 => 94063
  PMD: c0101c80 => 94063
  PTE: c00941b8 => 2216063
  PAGE: 2216000

    PTE      PHYSICAL  FLAGS
2216063    2216000  (PRESENT|RW|ACCESSED|DIRTY)

    PAGE      PHYSICAL  INODE      OFFSET  CNT  FLAGS
c02e9370    2216000          0          0  1

crash> vtop -c 1359 c806e000
VIRTUAL    PHYSICAL
c806e000    2216000

PAGE DIRECTORY: c5caf000
  PGD: c5cafc80 => 94063
  PMD: c5cafc80 => 94063
  PTE: c00941b8 => 2216063
  PAGE: 2216000
```

```

3134
3135         PTE      PHYSICAL  FLAGS
3136     2216063    2216000  (PRESENT|RW|ACCESSED|DIRTY)
3137
3138         PAGE      PHYSICAL  INODE      OFFSET  CNT  FLAGS
3139     c02e9370    2216000          0          0   1
3140
3141 Determine swap location of user virtual address 40104000:
3142
3143 crash> vtop 40104000
3144 VIRTUAL      PHYSICAL
3145 40104000    (not mapped)
3146
3147 PAGE DIRECTORY: c40d8000
3148     PGD: c40d8400 => 6bbe067
3149     PMD: c40d8400 => 6bbe067
3150     PTE: c6bbe410 => 58bc00
3151
3152         PTE      SWAP      OFFSET
3153     58bc00    /dev/sda8    22716
3154
3155         VMA      START      END      FLAGS  FILE
3156     c7200ae0    40104000    40b08000    73
3157
3158 SWAP: /dev/sda8  OFFSET: 22716
3159
3160
3161 NAME
3162     bt - backtrace
3163
3164 SYNOPSIS
3165     bt [-a|-c cpu(s)|-g|-r|-t|-T|-l|-e|-E|-f|-F|-o|-O|-v|-p] [-R ref] [-s [-x|d]]
3166         [-I ip] [-S sp] [-n idle] [pid | task]
3167
3168 DESCRIPTION
3169     Display a kernel stack backtrace.  If no arguments are given, the stack
3170     trace of the current context will be displayed.
3171
3172     -a displays the stack traces of the active task on each CPU.
3173         (only applicable to crash dumps)
3174     -A same as -a, but also displays vector registers (S390X only).
3175     -n idle filter the stack of idle tasks (x86_64, arm64).
3176         (only applicable to crash dumps)
3177     -p display the stack trace of the panic task only.
3178         (only applicable to crash dumps)
3179     -c cpu display the stack trace of the active task on one or more CPUs,
3180         which can be specified using the format "3", "1,8,9", "1-23",
3181         or "1,8,9-14". (only applicable to crash dumps)
3182     -g displays the stack traces of all threads in the thread group of
3183         the target task; the thread group leader will be displayed first.
3184     -r display raw stack data, consisting of a memory dump of the two
3185         pages of memory containing the task_union structure.
3186     -t display all text symbols found from the last known stack location
3187         to the top of the stack. (helpful if the back trace fails)
3188     -T display all text symbols found from just above the task_struct or
3189         thread_info to the top of the stack. (helpful if the back trace
3190         fails or the -t option starts too high in the process stack).
3191     -l show file and line number of each stack trace text location.
3192     -e search the stack for possible kernel and user mode exception frames.
3193     -E search the IRQ stacks (x86, x86_64, arm64, and ppc64), and the
3194         exception stacks (x86_64) for possible exception frames; all other
3195         arguments except for -c will be ignored since this is not a context-
3196         sensitive operation.
3197     -f display all stack data contained in a frame; this option can be
3198         used to determine the arguments passed to each function; on ia64,
3199         the argument register contents are dumped.
3200     -F[F] similar to -f, except that the stack data is displayed symbolically
3201         when appropriate; if the stack data references a slab cache object,
3202         the name of the slab cache will be displayed in brackets; on ia64,
3203         the substitution is done to the argument register contents.  If -F
3204         is entered twice, and the stack data references a slab cache object,
3205         both the address and the name of the slab cache will be displayed
3206         in brackets.

```

```

3207     -v check the kernel stack of all tasks for evidence of stack overflows.
3208     It does so by verifying the thread_info.task pointer, ensuring that
3209     the thread_info.cpu is a valid cpu number, and checking the end of
3210     the stack for the STACK_END_MAGIC value.
3211     -o arm64: use optional backtrace method; not supported on Linux 4.14 or
3212     later kernels.
3213     x86: use old backtrace method, permissible only on kernels that were
3214     compiled without the -fomit-frame_pointer.
3215     x86_64: use old backtrace method, which dumps potentially stale
3216     kernel text return addresses found on the stack.
3217     -O arm64: use optional backtrace method by default; subsequent usage
3218     of this option toggles the backtrace method.
3219     x86: use old backtrace method by default, permissible only on kernels
3220     that were compiled without the -fomit-frame_pointer; subsequent usage
3221     of this option toggles the backtrace method.
3222     x86_64: use old backtrace method by default; subsequent usage of this
3223     option toggles the backtrace method.
3224     -R ref display stack trace only if there is a reference to this symbol
3225     or text address.
3226     -s display the symbol name plus its offset.
3227     -x when displaying a symbol offset with the -s option, override the
3228     default output format with hexadecimal format.
3229     -d when displaying a symbol offset with the -s option, override the
3230     default output format with decimal format.
3231     -I ip use ip as the starting text location.
3232     -S sp use sp as the starting stack frame address.
3233     pid displays the stack trace(s) of this pid.
3234     taskp displays the stack trace the the task referenced by this hexadecimal
3235     task_struct pointer.
3236

```

Multiple pid and taskp arguments may be specified.

Note that all examples below are for x86 only. The output format will differ for other architectures. x86 backtraces from kernels that were compiled with the --fomit-frame-pointer CFLAG occasionally will drop stack frames, or display a stale frame reference. When in doubt as to the accuracy of a backtrace, the -t or -T options may help fill in the blanks.

EXAMPLES

Display the stack trace of the active task(s) when the kernel panicked:

```

3246 crash> bt -a
3247
3248 PID: 286 TASK: c0b3a000 CPU: 0 COMMAND: "in.rlogind"
3249 #0 [c0b3be90] crash_save_current_state at c011aed0
3250 #1 [c0b3bea4] panic at c011367c
3251 #2 [c0b3bee8] tulip_interrupt at c01bc820
3252 #3 [c0b3bf08] handle_IRQ_event at c010a551
3253 #4 [c0b3bf2c] do_8259A_IRQ at c010a319
3254 #5 [c0b3bf3c] do_IRQ at c010a653
3255 #6 [c0b3bfbc] ret_from_intr at c0109634
3256 EAX: 00000000 EBX: c0e68280 ECX: 00000000 EDX: 00000004 EBP: c0b3bfbc
3257 DS: 0018 ESI: 00000004 ES: 0018 EDI: c0e68284
3258 CS: 0010 EIP: c012f803 ERR: ffffffff09 EFLAGS: 00000246
3259 #7 [c0b3bfbc] sys_select at c012f803
3260 #8 [c0b3bfc0] system_call at c0109598
3261 EAX: 0000008e EBX: 00000004 ECX: bfffc9a0 EDX: 00000000
3262 DS: 002b ESI: bfffc8a0 ES: 002b EDI: 00000000
3263 SS: 002b ESP: bfffc82c EBP: bfffd224
3264 CS: 0023 EIP: 400d032e ERR: 0000008e EFLAGS: 00000246
3265

```

Display the stack trace of the active task(s) when the kernel panicked, and filter out the stack of the idle tasks:

```

3267 crash> bt -a -n idle
3268 ...
3269 PID: 0 TASK: ffff889ff8c35a00 CPU: 11 COMMAND: "swapper/11"
3270
3271 PID: 0 TASK: ffff889ff8c3c380 CPU: 12 COMMAND: "swapper/12"
3272
3273 PID: 150773 TASK: ffff889fe85a1680 CPU: 13 COMMAND: "bash"
3274 #0 [ffffc9000d35bcd0] machine_kexec at ffffffff8105a407
3275 #1 [ffffc9000d35bd28] __crash_kexec at ffffffff8113033d
3276 #2 [ffffc9000d35bdf0] panic at ffffffff81081930
3277

```

```

3280 #3 [fffffc9000d35be70] sysrq_handle_crash at ffffffff814e38d1
3281 #4 [fffffc9000d35be78] __handle_sysrq.cold.12 at ffffffff814e4175
3282 #5 [fffffc9000d35bea8] write_sysrq_trigger at ffffffff814e404b
3283 #6 [fffffc9000d35beb8] proc_reg_write at ffffffff81330d86
3284 #7 [fffffc9000d35bed0] vfs_write at ffffffff812a72d5
3285 #8 [fffffc9000d35bf00] ksys_write at ffffffff812a7579
3286 #9 [fffffc9000d35bf38] do_syscall_64 at ffffffff81004259
3287 RIP: 00007fa7abdc274 RSP: 00007ffa731f678 RFLAGS: 00000246
3288 RAX: ffffffff81000000 RBX: 0000000000000002 RCX: 00007fa7abdc274
3289 RDX: 0000000000000002 RSI: 0000563ca51ee6d0 RDI: 0000000000000001
3290 RBP: 0000563ca51ee6d0 R8: 000000000000000a R9: 00007fa7abd6be80
3291 R10: 000000000000000a R11: 0000000000000246 R12: 00007fa7abdad760
3292 R13: 0000000000000002 R14: 00007fa7abda8760 R15: 0000000000000002
3293 ORIG_RAX: 0000000000000001 CS: 0033 SS: 002b
3294 ...
3295

```

Display the stack trace of the active task on CPU 0 and 1:

```

3297 crash> bt -c 0,1
3298
3299 PID: 0      TASK: ffffffff81a8d020 CPU: 0      COMMAND: "swapper"
3300 #0 [ffff880002207e90] crash_nmi_callback at ffffffff8102fee6
3301 #1 [ffff880002207ea0] notifier_call_chain at ffffffff8152d525
3302 #2 [ffff880002207ee0] atomic_notifier_call_chain at ffffffff8152d58a
3303 #3 [ffff880002207ef0] notify_die at ffffffff810a155e
3304 #4 [ffff880002207f20] do_nmi at ffffffff8152b1eb
3305 #5 [ffff880002207f50] nmi at ffffffff8152aab0
3306 [exception RIP: native_safe_halt+0xb]
3307 RIP: ffffffff8103each RSP: ffffffff81a01ea8 RFLAGS: 00000296
3308 RAX: 0000000000000000 RBX: 0000000000000000 RCX: 0000000000000000
3309 RDX: 0000000000000000 RSI: 0000000000000001 RDI: ffffffff81de5228
3310 RBP: ffffffff81a01ea8 R8: 0000000000000000 R9: 0000000000000000
3311 R10: 0012099429a6bea3 R11: 0000000000000000 R12: ffffffff81c066c0
3312 R13: 0000000000000000 R14: ffffffff81de1000 R15: ffffffff81de1000
3313 ORIG_RAX: ffffffff81de1000 CS: 0010 SS: 0018
3314 --- <NMI exception stack> ---
3315 #6 [fffffff81a01ea8] native_safe_halt at ffffffff8103each
3316 #7 [fffffff81a01eb0] default_idle at ffffffff810167bd
3317 #8 [fffffff81a01ed0] cpu_idle at ffffffff81009fc6
3318
3319 PID: 38     TASK: ffff88003eaae040 CPU: 1     COMMAND: "khungtaskd"
3320 #0 [ffff88003ad97ce8] machine_kexec at ffffffff81038f3b
3321 #1 [ffff88003ad97d48] crash_kexec at ffffffff810c5da2
3322 #2 [ffff88003ad97e18] panic at ffffffff8152721a
3323 #3 [ffff88003ad97e98] watchdog at ffffffff810e6346
3324 #4 [ffff88003ad97ee8] kthread at ffffffff8109af06
3325 #5 [ffff88003ad97f48] kernel_thread at ffffffff8100c20a
3326

```

Display the stack traces of task f2814000 and PID 1592:

```

3327 crash> bt f2814000 1592
3328
3329 PID: 1018   TASK: f2814000 CPU: 1     COMMAND: "java"
3330 #0 [f2815db4] schedule at c011af85
3331 #1 [f2815de4] __down at c010600f
3332 #2 [f2815e14] __down_failed at c01061b3
3333 #3 [f2815e24] stext_lock (via drain_cpu_caches) at c025fa55
3334 #4 [f2815ec8] kmem_cache_shrink_nr at c013a53e
3335 #5 [f2815ed8] do_try_to_free_pages at c013f402
3336 #6 [f2815f04] try_to_free_pages at c013f8d2
3337 #7 [f2815f1c] _wrapped_alloc_pages at c01406bd
3338 #8 [f2815f40] __alloc_pages at c014079d
3339 #9 [f2815f60] __get_free_pages at c014083e
3340 #10 [f2815f68] do_fork at c011cebb
3341 #11 [f2815fa4] sys_clone at c0105ceb
3342 #12 [f2815fc0] system_call at c010740c
3343 EAX: 00000078 EBX: 00000f21 ECX: bc1ffbd8 EDX: bc1ffbe0
3344 DS: 002b      ESI: 00000000 ES: 002b      EDI: bc1ffd04
3345 SS: 002b      ESP: 0807316c EBP: 080731bc
3346 CS: 0023      EIP: 4012881e ERR: 00000078 EFLAGS: 00000296
3347
3348
3349 PID: 1592   TASK: c0cec000 CPU: 3     COMMAND: "httpd"
3350 #0 [c0ceded4] schedule at c011af85
3351 #1 [c0cedf04] pipe_wait at c0153083
3352 #2 [c0cedf58] pipe_read at c015317f

```

```

3353 #3 [c0cedf7c] sys_read at c0148be6
3354 #4 [c0cedfc0] system_call at c010740c
3355 EAX: 00000003 EBX: 00000004 ECX: bffed4a3 EDX: 00000001
3356 DS: 002b ESI: 00000001 ES: 002b EDI: bffed4a3
3357 SS: 002b ESP: bffed458 EBP: bffed488
3358 CS: 0023 EIP: 4024f1d4 ERR: 00000003 EFLAGS: 00000286
3359

```

In order to examine each stack frame's contents use the bt -f option. From the extra frame data that is displayed, the arguments passed to each function can be determined. Re-examining the PID 1592 trace above:

```

3360 crash> bt -f 1592
3361 PID: 1592 TASK: c0cec000 CPU: 3 COMMAND: "httpd"
3362 #0 [c0ceded4] schedule at c011af85
3363 [RA: c0153088 SP: c0ceded4 FP: c0cedf04 SIZE: 52]
3364 c0ceded4: c0cedf00 c0cec000 cela6000 00000003
3365 c0cedee4: c0cec000 f26152c0 cfafc8c0 c0cec000
3366 c0cedef4: ef70a0a0 c0cec000 c0cedf28 c0cedf54
3367 c0cedf04: c0153088
3368 #1 [c0cedf04] pipe_wait at c0153083
3369 [RA: c0153184 SP: c0cedf08 FP: c0cedf58 SIZE: 84]
3370 c0cedf08: 00000000 c0cec000 00000000 00000000
3371 c0cedf18: 00000000 c0a41fa0 c011d38b c0394120
3372 c0cedf28: 00000000 c0cec000 ceeebf30 ce4adf30
3373 c0cedf38: 00000000 d4b60ce0 00000000 c0cedf58
3374 c0cedf48: e204f820 ef70a040 00000001 c0cedf78
3375 c0cedf58: c0153184
3376 #2 [c0cedf58] pipe_read at c015317f
3377 [RA: c0148be8 SP: c0cedf5c FP: c0cedf7c SIZE: 36]
3378 c0cedf5c: ef70a040 c0cec000 00000000 00000000
3379 c0cedf6c: 00000001 f27ae680 ffffffff c0cedfbc
3380 c0cedf7c: c0148be8
3381 #3 [c0cedf7c] sys_read at c0148be6
3382 [RA: c0107413 SP: c0cedf80 FP: c0cedfc0 SIZE: 68]
3383 c0cedf80: f27ae680 bffed4a3 00000001 f27ae6a0
3384 c0cedf90: 40160370 24000000 4019ba28 00000000
3385 c0cedfa0: 00000000 ffffffff bffba207 ffffffff
3386 c0cedfb0: c0cec000 00000001 bffed4a3 bffed488
3387 c0cedfc0: c0107413
3388 #4 [c0cedfc0] system_call at c010740c
3389 EAX: 00000003 EBX: 00000004 ECX: bffed4a3 EDX: 00000001
3390 DS: 002b ESI: 00000001 ES: 002b EDI: bffed4a3
3391 SS: 002b ESP: bffed458 EBP: bffed488
3392 CS: 0023 EIP: 4024f1d4 ERR: 00000003 EFLAGS: 00000286
3393 [RA: 4024f1d4 SP: c0cedfc4 FP: c0cedffc SIZE: 60]
3394 c0cedfc4: 00000004 bffed4a3 00000001 00000001
3395 c0cedfd4: bffed4a3 bffed488 00000003 0000002b
3396 c0cedfe4: 0000002b 00000003 4024f1d4 00000023
3397 c0cedff4: 00000286 bffed458 0000002b
3402

```

Typically the arguments passed to a function will be the last values that were pushed onto the stack by the next higher-numbered function, i.e., the lowest stack addresses in the frame above the called function's stack frame. That can be verified by disassembling the calling function. For example, the arguments passed from sys_read() to pipe_read() above are the file pointer, the user buffer address, the count, and a pointer to the file structure's f_pos field. Looking at the frame #3 data for sys_read(), the last four items pushed onto the stack (lowest addresses) are f27ae680, bffed4a3, 00000001, and f27ae6a0 -- which are the 4 arguments above, in that order. Note that the first (highest address) stack content in frame #2 data for pipe_read() is c0148be8, which is the return address back to sys_read().

Dump the text symbols found in the current context's stack:

```

3416 crash> bt -t
3417 PID: 1357 TASK: claa0000 CPU: 0 COMMAND: "lockd"
3418 START: schedule at c01190e0
3419 [claa1f28] dput at c0157dbc
3420 [claa1f4c] schedule_timeout at c0124cd4
3421 [claa1f78] svc_rcv at cb22c4d8 [sunrpc]
3422 [claa1f98] put_files_struct at c011eb21
3423 [claa1fcc] nlmclnt_proc at cb237bef [lockd]
3425

```

```
3426 [c1aalfff0] kernel_thread at c0105826
3427 [c1aalfff8] nlmclnt_proc at cb237a60 [lockd]
3428
```

Search the current stack for possible exception frames:

```
3430
3431 crash> bt -e
3432 PID: 286      TASK: c0b3a000  CPU: 0    COMMAND: "in.rlogind"
3433
3434     KERNEL-MODE EXCEPTION FRAME AT c0b3bf44:
3435     EAX: 00000000  EBX: c0e68280  ECX: 00000000  EDX: 00000004  EBP: c0b3bfbc
3436     DS:  0018      ESI: 00000004  ES:  0018      EDI: c0e68284
3437     CS:  0010      EIP: c012f803  ERR: ffffffff09  EFLAGS: 00000246
3438
3439     USER-MODE EXCEPTION FRAME AT c0b3bfc4:
3440     EAX: 0000008e  EBX: 00000004  ECX: bfffc9a0  EDX: 00000000
3441     DS:  002b      ESI: bfffc8a0  ES:  002b      EDI: 00000000
3442     SS:  002b      ESP: bfffc82c  EBP: bfffd224
3443     CS:  0023      EIP: 400d032e  ERR: 0000008e  EFLAGS: 00000246
3444
```

Display the back trace from a dumpfile that resulted from the execution of the crash utility's "sys-panic" command:

```
3447
3448 crash> bt
3449 PID: 12523  TASK: c610c000  CPU: 0    COMMAND: "crash"
3450 #0 [c610de64] die at c01076ec
3451 #1 [c610de74] do_invalid_op at c01079bc
3452 #2 [c610df2c] error_code (via invalid_op) at c0107256
3453     EAX: 0000001d  EBX: c024a4c0  ECX: c02f13c4  EDX: 000026f6  EBP: c610c000
3454     DS:  0018      ESI: 401de2e0  ES:  0018      EDI: c610c000
3455     CS:  0010      EIP: c011bbb4  ERR: ffffffff  EFLAGS: 00010296
3456 #3 [c610df68] panic at c011bbb4
3457 #4 [c610df78] do_exit at c011f1fe
3458 #5 [c610dfc0] system_call at c0107154
3459     EAX: 00000001  EBX: 00000000  ECX: 00001000  EDX: 401df154
3460     DS:  002b      ESI: 401de2e0  ES:  002b      EDI: 00000000
3461     SS:  002b      ESP: bffebf0c  EBP: bffebf38
3462     CS:  0023      EIP: 40163afd  ERR: 00000001  EFLAGS: 00000246
3463
```

Display the back trace from a dumpfile that resulted from an attempt to insmod the sample "crash.c" kernel module that comes as part of the Red Hat netdump package:

```
3467
3468 crash> bt
3469 PID: 1696  TASK: c74de000  CPU: 0    COMMAND: "insmod"
3470 #0 [c74dfdcc] die at c01076ec
3471 #1 [c74dfddc] do_page_fault at c0117bbc
3472 #2 [c74dfee0] error_code (via page_fault) at c0107256
3473     EAX: 00000013  EBX: cb297000  ECX: 00000000  EDX: c5962000  EBP: c74dff28
3474     DS:  0018      ESI: 00000000  ES:  0018      EDI: 00000000
3475     CS:  0010      EIP: cb297076  ERR: ffffffff  EFLAGS: 00010282
3476 #3 [c74dff1c] crash_init at cb297076 [crash]
3477 #4 [c74dff2c] sys_init_module at c011d233
3478 #5 [c74dffc0] system_call at c0107154
3479     EAX: 00000080  EBX: 08060528  ECX: 08076450  EDX: 0000000a
3480     DS:  002b      ESI: 0804b305  ES:  002b      EDI: 08074ed0
3481     SS:  002b      ESP: bffe9a90  EBP: bffe9ac8
3482     CS:  0023      EIP: 4012066e  ERR: 00000080  EFLAGS: 00000246
3483
```

Display the symbol name plus its offset in each frame, overriding the current output format with hexadecimal:

```
3487
3488 crash> bt -sx
3489 PID: 1499  TASK: ffff88006af43cc0  CPU: 2    COMMAND: "su"
3490 #0 [ffff8800664a1c90] machine_kexec+0x167 at ffffffff810327b7
3491 #1 [ffff8800664a1ce0] crash_kexec+0x60 at ffffffff810a9ec0
3492 #2 [ffff8800664a1db0] oops_end+0xb0 at ffffffff81504160
3493 #3 [ffff8800664a1dd0] general_protection+0x25 at ffffffff81503435
3494     [exception RIP: kmem_cache_alloc+120]
3495     RIP: ffffffff81113cf88  RSP: ffff8800664a1e88  RFLAGS: 00010086
3496     RAX: 0000000000000000  RBX: ff88006ef56840ff  RCX: ffffffff8114e9e4
3497     RDY: 0000000000000000  RSI: 000000000000080d0  RDI: ffffffff81796020
3498     RBP: ffffffff81796020  R8: ffff88000a3137a0  R9: 0000000000000000
3499     R10: ffff88007ac97300  R11: 0000000000000400  R12: 000000000000080d0
```



```

3499 R13: 00000000000000292 R14: 000000000000080d0 R15: 00000000000000c0
3500 ORIG_RAX: ffffffff8114e9e4 CS: 0010 SS: 0018
3501 #4 [ffff8800664aled0] get_empty_filp+0x74 at ffffffff8114e9e4
3502 #5 [ffff8800664alef0] sock_alloc_fd+0x23 at ffffffff81142f553
3503 #6 [ffff8800664alf10] sock_map_fd+0x23 at ffffffff81142f693
3504 #7 [ffff8800664alf50] sys_socket+0x43 at ffffffff8114302a3
3505 #8 [ffff8800664alf80] system_call_fastpath+0x16 at ffffffff811013042
3506 RIP: 00007f5720b368e7 RSP: 00007fff52b629a8 RFLAGS: 00010206
3507 RAX: 0000000000000029 RBX: ffffffff81013042 RCX: 0000000000000000
3508 RDX: 0000000000000009 RSI: 0000000000000003 RDI: 0000000000000010
3509 RBP: 00000000000066f320 R8: 0000000000000001 R9: 0000000000000000
3510 R10: 0000000000000000 R11: 0000000000000202 R12: ffff88007ac97300
3511 R13: 0000000000000000 R14: 00007f571e104a80 R15: 00007f571e305048
3512 ORIG_RAX: 0000000000000029 CS: 0033 SS: 002b
3513

```

The following three examples show the difference in the display of the same stack frame's contents using -f, -F, and -FF:

```

3514 crash> bt -f
3515 ...
3516 #4 [ffff810072b47f10] vfs_write at ffffffff800789d8
3517 ffff810072b47f18: ffff81007e020380 ffff81007e2c2880
3518 ffff810072b47f28: 0000000000000002 ffffffff81142f553
3519 ffff810072b47f38: 00002b141825d000 ffffffff80078f75
3520 #5 [ffff810072b47f40] sys_write at ffffffff80078f75
3521 ...
3522 crash> bt -F
3523 ...
3524 #4 [ffff810072b47f10] vfs_write at ffffffff800789d8
3525 ffff810072b47f18: [files_cache] [filp]
3526 ffff810072b47f28: 0000000000000002 ffffffff81142f553
3527 ffff810072b47f38: 00002b141825d000 sys_write+69
3528 #5 [ffff810072b47f40] sys_write at ffffffff80078f75
3529 ...
3530 crash> bt -FF
3531 ...
3532 #4 [ffff810072b47f10] vfs_write at ffffffff800789d8
3533 ffff810072b47f18: [ffff81007e020380:files_cache] [ffff81007e2c2880:filp]
3534 ffff810072b47f28: 0000000000000002 ffffffff81142f553
3535 ffff810072b47f38: 00002b141825d000 sys_write+69
3536 #5 [ffff810072b47f40] sys_write at ffffffff80078f75
3537 ...

```

Check the kernel stack of all tasks for evidence of a stack overflow:

```

3544 crash> bt -v
3545 PID: 5823 TASK: ffff88102aae0040 CPU: 1 COMMAND: "flush-253:0"
3546 possible stack overflow: thread_info.task: 102efb5adc0 != ffff88102aae0040
3547 possible stack overflow: 40ffffff != STACK_END_MAGIC
3548

```

NAME

help - get help

SYNOPSIS

help [command | all] [-<option>]

DESCRIPTION

When entered with no argument, a list of all currently available crash commands is listed. If a name of a crash command is entered, a man-like page for the command is displayed. If "all" is entered, help pages for all commands will be displayed. If neither of the above is entered, the argument string will be passed on to the gdb help command.

A number of internal debug, statistical, and other dumpfile related data is available with the following options:

```

3566 -a - alias data
3567 -b - shared buffer data
3568 -B - build data
3569 -c - numargs cache
3570 -d - device table
3571 -D - dumpfile contents/statistics

```

```

3572 -e - extension table data
3573 -f - filesys table
3574 -g - gdb data
3575 -h - hash_table data
3576 -H - hash_table data (verbose)
3577 -k - kernel_table
3578 -K - kernel_table (verbose)
3579 -L - LKCD page cache environment
3580 -M <num> machine specific
3581 -m - machdep_table
3582 -N - net_table
3583 -n - dumpfile contents/statistics
3584 -o - offset_table and size_table
3585 -p - program_context
3586 -r - dump registers from dumpfile header
3587 -s - symbol table data
3588 -t - task_table
3589 -T - task_table plus context_array
3590 -v - vm_table
3591 -V - vm_table (verbose)
3592 -z - help options

```

NAME

```
ps - display process status information
```

SYNOPSIS

```
ps [-k|-u|-G|-y policy] [-s] [-p|-c|-t|-l|m] [-C cpu] [-a|-g|-r|-S|-A]
    [pid | task | command] ...
```

DESCRIPTION

This command displays process status for selected, or all, processes in the system. If no arguments are entered, the process data is displayed for all processes. Specific processes may be selected by using the following identifier formats:

```

pid      a process PID.
task     a hexadecimal task_struct pointer.
command  a command name. If a command name is made up of letters that
         are all numerical values, precede the name string with a "\".
         If the command string is enclosed within '"' characters, then
         the encompassed string must be a POSIX extended regular expression
         that will be used to match task names.

```

The process list may be further restricted by the following options:

```

-k      restrict the output to only kernel threads.
-u      restrict the output to only user tasks.
-G      display only the thread group leader in a thread group.
-y policy restrict the output to tasks having a specified scheduling policy
         expressed by its integer value or by its (case-insensitive) name;
         multiple policies may be entered in a comma-separated list:
         0 or NORMAL
         1 or FIFO
         2 or RR
         3 or BATCH
         4 or ISO
         5 or IDLE
         6 or DEADLINE

```

The process identifier types may be mixed. For each task, the following items are displayed:

1. the process PID.
2. the parent process PID.
3. the CPU number that the task ran on last.
4. the task_struct address or the kernel stack pointer of the process. (see -s option below)
5. the task state (RU, IN, UN, ZO, ST, TR, DE, SW, WA, PA, ID, NE).
6. the percentage of physical memory being used by this task.
7. the virtual address size of this task in kilobytes.
8. the resident set size of this task in kilobytes.
9. the command name.

The default output shows the task_struct address of each process under a column titled "TASK". This can be changed to show the kernel stack pointer under a column titled "KSTACKP".

-s replace the TASK column with the KSTACKP column.

On SMP machines, the active task on each CPU will be highlighted by an angle bracket (">") preceding its information. If the crash variable "offline" is set to "hide", the active task on an offline CPU will be highlighted by a "-" preceding its information.

Alternatively, information regarding parent-child relationships, per-task time usage data, argument/environment data, thread groups, or resource limits may be displayed:

-p display the parental hierarchy of selected, or all, tasks.
-c display the children of selected, or all, tasks.
-t display the task run time, start time, and cumulative user and system times.
-l display the task's last-run timestamp value, using either the task_struct's last_run value, the task_struct's timestamp value or the task_struct's sched_entity last_arrival value, whichever applies, of selected, or all, tasks; the list is sorted with the most recently-run task (with the largest timestamp) shown first, followed by the task's current state.
-m similar to -l, but the timestamp value is translated into days, hours, minutes, seconds, and milliseconds since the task was last run on a cpu.
-C cpus only usable with the -l or -m options, dump the timestamp data in per-cpu blocks, where the cpu[s] can be specified as "1,3,5", "1-3", "1,3,5-7,10", "all", or "a" (shortcut for "all").
-a display the command line arguments and environment strings of selected, or all, user-mode tasks.
-g display tasks by thread group, of selected, or all, tasks.
-r display resource limits (rlimits) of selected, or all, tasks.
-S display a summary consisting of the number of tasks in a task state.
-A display only the active task on each cpu.

EXAMPLES

Show the process status of all current tasks:

crash> ps

	PID	PPID	CPU	TASK	ST	%MEM	VSZ	RSS	COMM
>	0	0	3	c024c000	RU	0.0	0	0	[swapper]
>	0	0	0	c0dce000	RU	0.0	0	0	[swapper]
	0	0	1	c0fa8000	RU	0.0	0	0	[swapper]
>	0	0	2	c009a000	RU	0.0	0	0	[swapper]
	1	0	1	c0098000	IN	0.0	1096	476	init
	2	1	1	c0090000	IN	0.0	0	0	[kflushd]
	3	1	1	c000e000	IN	0.0	0	0	[kpiod]
	4	1	3	c000c000	IN	0.0	0	0	[kswapd]
	5	1	1	c0008000	IN	0.0	0	0	[mdrecoveryd]
	253	1	2	fb4c000	IN	0.0	1088	376	portmap
	268	1	2	fb482000	IN	0.1	1232	504	yplibind
	274	268	2	fa984000	IN	0.1	1260	556	yplibind
	321	1	1	fabf6000	IN	0.1	1264	608	syslogd
	332	1	1	fa9be000	RU	0.1	1364	736	klogd
	346	1	2	fae88000	IN	0.0	1112	472	atd
	360	1	2	faeb2000	IN	0.1	1284	592	crond
	378	1	2	fafd6000	IN	0.1	1236	560	inetd
	392	1	0	fb710000	IN	0.1	2264	1468	named
	406	1	3	fb768000	IN	0.1	1284	560	lpd
	423	1	1	fb8ac000	IN	0.1	1128	528	rpc.statd
	434	1	2	fb75a000	IN	0.0	1072	376	rpc.rquotad
	445	1	2	fb4a4000	IN	0.0	1132	456	rpc.mountd
	460	1	1	fa938000	IN	0.0	0	0	[nfsd]
	461	1	1	faa86000	IN	0.0	0	0	[nfsd]
	462	1	0	fac48000	IN	0.0	0	0	[nfsd]
	463	1	0	fb4ca000	IN	0.0	0	0	[nfsd]
	464	1	0	fb4c8000	IN	0.0	0	0	[nfsd]
	465	1	2	fb46e000	IN	0.0	0	0	[nfsd]
	466	1	1	fb46c000	IN	0.0	0	0	[nfsd]

```

3718      467      1      2      fac04000      IN      0.0      0      0      [nfsd]
3719      468      461      2      fa93a000      IN      0.0      0      0      [lockd]
3720      469      468      2      fa93e000      IN      0.0      0      0      [rpciod]
3721      486      1      0      fab54000      IN      0.1      1596      880      amd
3722      523      1      2      fa84e000      IN      0.1      1884      1128      sendmail
3723      538      1      0      fa82c000      IN      0.0      1112      416      gpm
3724      552      1      3      fa70a000      IN      0.1      2384      1220      httpd
3725      556      552      3      fa776000      IN      0.1      2572      1352      httpd
3726      557      552      2      faba4000      IN      0.1      2572      1352      httpd
3727      558      552      1      fa802000      IN      0.1      2572      1352      httpd
3728      559      552      3      fa6ee000      IN      0.1      2572      1352      httpd
3729      560      552      3      fa700000      IN      0.1      2572      1352      httpd
3730      561      552      0      fa6f0000      IN      0.1      2572      1352      httpd
3731      562      552      3      fa6ea000      IN      0.1      2572      1352      httpd
3732      563      552      0      fa67c000      IN      0.1      2572      1352      httpd
3733      564      552      3      fa674000      IN      0.1      2572      1352      httpd
3734      565      552      3      fa66a000      IN      0.1      2572      1352      httpd
3735      582      1      2      fa402000      IN      0.2      2968      1916      xfs
3736      633      1      2      falec000      IN      0.2      5512      2248      innd
3737      636      1      3      fa088000      IN      0.1      2536      804      actived
3738      676      1      0      fa840000      IN      0.0      1060      384      mingetty
3739      677      1      1      fa590000      IN      0.0      1060      384      mingetty
3740      678      1      2      fa3b8000      IN      0.0      1060      384      mingetty
3741      679      1      0      fa5b8000      IN      0.0      1060      384      mingetty
3742      680      1      1      fa3a4000      IN      0.0      1060      384      mingetty
3743      681      1      2      fa30a000      IN      0.0      1060      384      mingetty
3744      683      1      3      fa5d8000      IN      0.0      1052      280      update
3745      686      378      1      fa3aa000      IN      0.1      2320      1136      in.rlogind
3746      687      686      2      f9e52000      IN      0.1      2136      1000      login
3747      688      687      0      f9dec000      IN      0.1      1732      976      bash
3748      >      700      688      1      f9d62000      RU      0.0      1048      256      gen12
3749

```

Display the parental hierarchy of the "crash" process on a live system:

```

3750
3751      crash> ps -p 4249
3752
3753      PID: 0      TASK: c0252000      CPU: 0      COMMAND: "swapper"
3754      PID: 1      TASK: c009a000      CPU: 1      COMMAND: "init"
3755      PID: 632     TASK: c73b6000      CPU: 1      COMMAND: "prefdm"
3756      PID: 637     TASK: c5a4a000      CPU: 1      COMMAND: "prefdm"
3757      PID: 649     TASK: c179a000      CPU: 0      COMMAND: "kwm"
3758      PID: 683     TASK: c1164000      CPU: 0      COMMAND: "kfm"
3759      PID: 1186    TASK: c165a000      CPU: 0      COMMAND: "xterm"
3760      PID: 1188    TASK: c705e000      CPU: 1      COMMAND: "bash"
3761      PID: 4249    TASK: c6b9a000      CPU: 0      COMMAND: "crash"
3762

```

Display all children of the "kwm" window manager:

```

3763
3764      crash> ps -c kwm
3765
3766      PID: 649     TASK: c179a000      CPU: 0      COMMAND: "kwm"
3767      PID: 682     TASK: c2d58000      CPU: 1      COMMAND: "kwmsound"
3768      PID: 683     TASK: c1164000      CPU: 1      COMMAND: "kfm"
3769      PID: 685     TASK: c053c000      CPU: 0      COMMAND: "krootwm"
3770      PID: 686     TASK: c13fa000      CPU: 0      COMMAND: "kpanel"
3771      PID: 687     TASK: c13f0000      CPU: 1      COMMAND: "kbgndwm"
3772

```

Display all threads in a firefox session:

```

3773
3774      crash> ps firefox
3775
3776      PID      PPID      CPU      TASK      ST      %MEM      VSZ      RSS      COMM
3777      21273     21256     6      ffff81003ec15080      IN      46.3      1138276      484364      firefox
3778      21276     21256     6      ffff81003f49e7e0      IN      46.3      1138276      484364      firefox
3779      21280     21256     0      ffff81003ec1d7e0      IN      46.3      1138276      484364      firefox
3780      21286     21256     6      ffff81000b0d1820      IN      46.3      1138276      484364      firefox
3781      21287     21256     2      ffff81000b0d10c0      IN      46.3      1138276      484364      firefox
3782      26975     21256     5      ffff81003b5c1820      IN      46.3      1138276      484364      firefox
3783      26976     21256     5      ffff810023232820      IN      46.3      1138276      484364      firefox
3784      26977     21256     4      ffff810021a11820      IN      46.3      1138276      484364      firefox
3785      26978     21256     5      ffff810003159040      IN      46.3      1138276      484364      firefox
3786      26979     21256     5      ffff81003a058820      IN      46.3      1138276      484364      firefox
3787

```

Display only the thread group leader in the firefox session:

```

3788
3789      crash> ps -G firefox
3790

```

```
3791      PID      PPID    CPU      TASK      ST  %MEM      VSZ      RSS      COMM
3792      21273    21256     0  ffff81003ec15080  IN  46.3 1138276 484364  firefox
3793
```

3794 Show the time usage data for pid 10318:

```
3795
3796 crash> ps -t 10318
3797 PID: 10318 TASK: f7b85550 CPU: 5  COMMAND: "bash"
3798 RUN TIME: 1 days, 01:35:32
3799 START TIME: 5209
3800 UTIME: 95
3801 STIME: 57
3802
```

3803 Show the process status of PID 1, task f9dec000, and all nfsd tasks:

```
3804
3805 crash> ps 1 f9dec000 nfsd
3806      PID      PPID    CPU      TASK      ST  %MEM      VSZ      RSS      COMM
3807      1         0      1  c0098000  IN   0.0    1096     476    init
3808     688      687     0  f9dec000  IN   0.1    1732     976    bash
3809     460        1      1  fa938000  IN   0.0        0        0  [nfsd]
3810     461        1      1  faa86000  IN   0.0        0        0  [nfsd]
3811     462        1      0  fac48000  IN   0.0        0        0  [nfsd]
3812     463        1      0  fb4ca000  IN   0.0        0        0  [nfsd]
3813     464        1      0  fb4c8000  IN   0.0        0        0  [nfsd]
3814     465        1      2  fba6e000  IN   0.0        0        0  [nfsd]
3815     466        1      1  fba6c000  IN   0.0        0        0  [nfsd]
3816     467        1      2  fac04000  IN   0.0        0        0  [nfsd]
3817
```

3818 Show all kernel threads:

```
3819
3820 crash> ps -k
3821      PID      PPID    CPU      TASK      ST  %MEM      VSZ      RSS      COMM
3822        0         0      1  c0fac000  RU   0.0        0        0  [swapper]
3823        0         0      0  c0252000  RU   0.0        0        0  [swapper]
3824        2         1      1  c0fa0000  IN   0.0        0        0  [kflushd]
3825        3         1      1  c03de000  IN   0.0        0        0  [kpiod]
3826        4         1      1  c03dc000  IN   0.0        0        0  [kswapd]
3827        5         1      0  c0092000  IN   0.0        0        0  [mdrecoveryd]
3828     336        1      0  c4a9a000  IN   0.0        0        0  [rpciod]
3829     337        1      0  c4830000  IN   0.0        0        0  [lockd]
3830     487        1      1  c4ba6000  IN   0.0        0        0  [nfsd]
3831     488        1      0  c18c6000  IN   0.0        0        0  [nfsd]
3832     489        1      0  c0cac000  IN   0.0        0        0  [nfsd]
3833     490        1      0  c056a000  IN   0.0        0        0  [nfsd]
3834     491        1      0  c0860000  IN   0.0        0        0  [nfsd]
3835     492        1      1  c0254000  IN   0.0        0        0  [nfsd]
3836     493        1      0  c0a86000  IN   0.0        0        0  [nfsd]
3837     494        1      0  c0968000  IN   0.0        0        0  [nfsd]
3838
```

3839 Display a summary consisting of the number of tasks in a task state:

```
3840
3841 crash> ps -S
3842 RU: 5
3843 IN: 259
3844 UN: 31
3845 ZO: 1
3846
```

3847 Display only the active task, on each cpu:

```
3848
3849 crash> ps -A
3850      PID      PPID    CPU      TASK      ST  %MEM      VSZ      RSS      COMM
3851 >    10         2      1  ffff880212969710  IN  0.0        0        0  [migration/1]
3852 >      0         0      3  ffff884026d43520  RU  0.0        0        0  [swapper]
3853 >   6582        1      2  ffff880f49c52040  RU  0.0 42202472 33368  oracle
3854 >   9497        1      0  ffff880549ec2ab0  RU  0.0 42314692 138664  oracle
3855
```

3856 Show all tasks sorted by their task_struct's last_run, timestamp, or
3857 sched_entity last_arrival timestamp value, whichever applies:

```
3858
3859 crash> ps -l
3860 [20811245123] [IN] PID: 37  TASK: f7153030 CPU: 2  COMMAND: "events/2"
3861 [20811229959] [IN] PID: 1756 TASK: f2a5a570 CPU: 2  COMMAND: "ntpd"
3862 [20800696644] [IN] PID: 1456 TASK: f2b1f030 CPU: 4  COMMAND: "irqbalance"
3863 [20617047229] [IN] PID: 2324 TASK: f57f9570 CPU: 5  COMMAND: "flush-253:0"

```

```

3864 [20617029209] [IN] PID: 49 TASK: f7167030 CPU: 4 COMMAND: "bdi-default"
3865 [20438025365] [IN] PID: 345 TASK: f55c7ab0 CPU: 3 COMMAND: "mpt_poll_0"
3866 [20103026046] [IN] PID: 728 TASK: f72ba570 CPU: 3 COMMAND: "edac-poller"
3867 [20000189409] [IN] PID: 35 TASK: f7153ab0 CPU: 0 COMMAND: "events/0"
3868 [20000179905] [IN] PID: 48 TASK: f7167570 CPU: 0 COMMAND: "sync_supers"
3869 [19997120354] [IN] PID: 36 TASK: f7153570 CPU: 1 COMMAND: "events/1"
3870 [19991059209] [IN] PID: 38 TASK: f715fab0 CPU: 3 COMMAND: "events/3"
3871 [19988091608] [IN] PID: 39 TASK: f715f570 CPU: 4 COMMAND: "events/4"
3872 [19985076530] [IN] PID: 40 TASK: f715f030 CPU: 5 COMMAND: "events/5"
3873 [19982019106] [IN] PID: 41 TASK: f7161ab0 CPU: 6 COMMAND: "events/6"
3874 [19982016294] [IN] PID: 29 TASK: f7109ab0 CPU: 6 COMMAND: "ksoftirqd/6"
3875 [19838402345] [RU] PID: 2331 TASK: f297f570 CPU: 7 COMMAND: "bash"
3876 [19837129436] [IN] PID: 2326 TASK: f2ad5030 CPU: 6 COMMAND: "sshd"
3877 [19289476417] [IN] PID: 1772 TASK: f5665570 CPU: 5 COMMAND: "sendmail"
3878 ...
3879

```

Show the most-recently run tasks on cpu 0 using both the -l and the -m options:

```

3883 crash> ps -m -C0
3884 CPU: 0
3885 [ 0 00:00:00.003] [RU] PID: 1205 TASK: dee03f20 CPU: 0 COMMAND: "insmod"
3886 [ 0 00:00:00.006] [RU] PID: 770 TASK: df9e9940 CPU: 0 COMMAND: "rsyslogd"
3887 [ 0 00:00:00.009] [IN] PID: 603 TASK: df9bcbcb0 CPU: 0 COMMAND: "udev"
3888 [ 0 00:00:00.010] [IN] PID: 348 TASK: df9ecbcb0 CPU: 0 COMMAND: "udev"
3889 [ 0 00:00:00.013] [IN] PID: 934 TASK: df9171a0 CPU: 0 COMMAND: "hald"
3890 [ 0 00:00:00.023] [IN] PID: 6 TASK: df443f20 CPU: 0 COMMAND: "events/0"
3891 [ 0 00:00:00.029] [IN] PID: 15 TASK: df46b280 CPU: 0 COMMAND: "kblockd/0"
3892 [ 0 00:00:00.101] [IN] PID: 1168 TASK: dee01940 CPU: 0 COMMAND: "bash"
3893 [ 0 00:00:01.404] [IN] PID: 272 TASK: dfa48ca0 CPU: 0 COMMAND: "flush-8:0"
3894 ...
3895

```

```

3896 crash> ps -l -C0
3897 CPU: 0
3898 [137146164748] [RU] PID: 1205 TASK: dee03f20 CPU: 0 COMMAND: "insmod"
3899 [137142534372] [RU] PID: 770 TASK: df9e9940 CPU: 0 COMMAND: "rsyslogd"
3900 [137140168469] [IN] PID: 603 TASK: df9bcbcb0 CPU: 0 COMMAND: "udev"
3901 [137138826427] [IN] PID: 348 TASK: df9ecbcb0 CPU: 0 COMMAND: "udev"
3902 [137135214599] [IN] PID: 934 TASK: df9171a0 CPU: 0 COMMAND: "hald"
3903 [137125651275] [IN] PID: 6 TASK: df443f20 CPU: 0 COMMAND: "events/0"
3904 [137119564815] [IN] PID: 15 TASK: df46b280 CPU: 0 COMMAND: "kblockd/0"
3905 [137047715027] [IN] PID: 1168 TASK: dee01940 CPU: 0 COMMAND: "bash"
3906 [135744209052] [IN] PID: 272 TASK: dfa48ca0 CPU: 0 COMMAND: "flush-8:0"
3907 ...
3908

```

Show the kernel stack pointer of each user task:

```

3911 crash> ps -us
3912 PID PPID CPU KSTACKP ST %MEM VSZ RSS COMM
3913 1 0 0 c009bedc IN 0.0 1096 52 init
3914 239 1 0 c15e7ed8 IN 0.2 1332 224 pump
3915 280 1 1 c7cbdedc IN 0.2 1092 208 portmap
3916 295 1 0 c7481edc IN 0.0 1232 0 ypbind
3917 301 295 0 c7c7bf28 IN 0.1 1260 124 ypbind
3918 376 1 1 c5053f28 IN 0.0 1316 40 automount
3919 381 1 0 c34ddf28 IN 0.2 1316 224 automount
3920 391 1 1 c2777f28 IN 0.2 1316 224 automount
3921 ...
3922

```

Display the argument and environment data for the automount task:

```

3925 crash> ps -a automount
3926 PID: 3948 TASK: f722ee30 CPU: 0 COMMAND: "automount"
3927 ARG: /usr/sbin/automount --timeout=60 /net program /etc/auto.net
3928 ENV: SELINUX_INIT=YES
3929     CONSOLE=/dev/console
3930     TERM=linux
3931     INIT_VERSION=sysvinit-2.85
3932     PATH=/sbin:/usr/sbin:/bin:/usr/bin
3933     LC_MESSAGES=en_US
3934     RUNLEVEL=3
3935     runlevel=3
3936     PWD=/

```

```
3937     LANG=ja_JP.UTF-8
3938     PREVLEVEL=N
3939     previous=N
3940     HOME=/
3941     SHLVL=2
3942     _=/usr/sbin/automount
3943
```

3944 Display the tasks in the thread group containing task c20ab0b0:

```
3945
3946 crash> ps -g c20ab0b0
3947 PID: 6425   TASK: f72f50b0  CPU: 0   COMMAND: "firefox-bin"
3948 PID: 6516   TASK: f71bf1b0  CPU: 0   COMMAND: "firefox-bin"
3949 PID: 6518   TASK: d394b930  CPU: 0   COMMAND: "firefox-bin"
3950 PID: 6520   TASK: c20aa030  CPU: 0   COMMAND: "firefox-bin"
3951 PID: 6523   TASK: c20ab0b0  CPU: 0   COMMAND: "firefox-bin"
3952 PID: 6614   TASK: f1f181b0  CPU: 0   COMMAND: "firefox-bin"
3953
```

3954 Display the tasks in the thread group for each instance of the
3955 program named "multi-thread":

```
3956
3957 crash> ps -g multi-thread
3958 PID: 2522   TASK: 1003f0dc7f0    CPU: 1   COMMAND: "multi-thread"
3959 PID: 2523   TASK: 10037b13030    CPU: 1   COMMAND: "multi-thread"
3960 PID: 2524   TASK: 1003e064030    CPU: 1   COMMAND: "multi-thread"
3961 PID: 2525   TASK: 1003e13a7f0    CPU: 1   COMMAND: "multi-thread"
3962
3963 PID: 2526   TASK: 1002f82b7f0    CPU: 1   COMMAND: "multi-thread"
3964 PID: 2527   TASK: 1003e1737f0    CPU: 1   COMMAND: "multi-thread"
3965 PID: 2528   TASK: 10035b4b7f0    CPU: 1   COMMAND: "multi-thread"
3966 PID: 2529   TASK: 1003f0c37f0    CPU: 1   COMMAND: "multi-thread"
3967 PID: 2530   TASK: 10035597030    CPU: 1   COMMAND: "multi-thread"
3968 PID: 2531   TASK: 100184be7f0    CPU: 1   COMMAND: "multi-thread"
3969
```

3970 Display the resource limits of "bash" task 13896:

```
3971
3972 crash> ps -r 13896
3973 PID: 13896  TASK: cf402000  CPU: 0   COMMAND: "bash"
3974 RLIMIT      CURRENT      MAXIMUM
3975 CPU         (unlimited)    (unlimited)
3976 FSIZE      (unlimited)    (unlimited)
3977 DATA      (unlimited)    (unlimited)
3978 STACK      10485760      (unlimited)
3979 CORE       (unlimited)    (unlimited)
3980 RSS        (unlimited)    (unlimited)
3981 NPROC      4091          4091
3982 NOFILE     1024          1024
3983 MEMLOCK    4096          4096
3984 AS         (unlimited)    (unlimited)
3985 LOCKS      (unlimited)    (unlimited)
3986
```

3987 Search for task names matching a POSIX regular expression:

```
3988
3989 crash> ps 'migration*'
3990 PID  PPID  CPU  TASK          ST  %MEM  VSZ   RSS  COMM
3991 8     2    0   ffff8802128a2e20  IN   0.0    0     0   [migration/0]
3992 10    2    1   ffff880212969710  IN   0.0    0     0   [migration/1]
3993 15    2    2   ffff880212989710  IN   0.0    0     0   [migration/2]
3994 20    2    3   ffff8802129a9710  IN   0.0    0     0   [migration/3]
3995
```

3997 NAME

3998 struct - structure contents

4000 SYNOPSIS

```
4001 struct struct_name[.member[,member]] [-o] [-l offset] [-rfuxdp]
4002 [address | symbol][:cpuspec] [count | -c count]
```

4004 DESCRIPTION

4005 This command displays either a structure definition, or a formatted display
4006 of the contents of a structure at a specified address. When no address is
4007 specified, the structure definition is shown along with the structure size.
4008 A structure member may be appended to the structure name in order to limit
4009 the scope of the data displayed to that particular member; when no address

is specified, the member's offset and definition are shown.

struct_name name of a C-code structure used by the kernel.
.member name of a structure member; to display multiple members of a structure, use a comma-separated list of members. If any member contains an embedded structure, or the member is an array, the output may be restricted to just the embedded structure or an array element by expressing the member argument as "member.member" or "member[index]"; embedded member specifications may extend beyond one level deep, by expressing the member argument as "member.member.member...".
-o show member offsets when displaying structure definitions; if used with an address or symbol argument, each member will be preceded by its virtual address.
-l offset if the address argument is a pointer to a structure member that is contained by the target data structure, typically a pointer to an embedded list_head, the offset to the embedded member may be entered in either of the following manners:
 1. in "structure.member" format.
 2. a number of bytes.
-r raw dump of structure data.
-f address argument is a dumpfile offset.
-u address argument is a user virtual address in the current context.
-x override default output format with hexadecimal format.
-d override default output format with decimal format.
-p if a structure member is a pointer value, show the member's data type on the output line; and on the subsequent line(s), dereference the pointer, display the pointer target's symbol value in brackets if appropriate, and if possible, display the target data; requires an address argument.
address hexadecimal address of a structure; if the address points to an embedded list_head structure contained within the target data structure, then the "-l" option must be used.
symbol symbolic reference to the address of a structure.
:cpuspec CPU specification for a per-cpu address or symbol:
 : CPU of the currently selected task.
 :a[ll] all CPUs.
 :#[-#][,...] CPU list(s), e.g. "1,3,5", "1-3", or "1,3,5-7,10".
count count of structures to dump from an array of structures; if used, this must be the last argument entered.
-c count "-c" is only required if "count" is not the last argument entered or if a negative number is entered; if a negative value is entered, the (positive) "count" structures that lead up to and include the target structure will be displayed.

Structure data, sizes, and member offsets are shown in the current output radix unless the -x or -d option is specified.

Please note that in the vast majority of cases, the "struct" command name may be dropped; if the structure name does not conflict with any crash or gdb command name, then the "struct_name[.member]" argument will be recognized as a structure name, and this command automatically executed. See the NOTE below.

EXAMPLES

Display the vm_area_struct at address cle44f10:

```
crash> struct vm_area_struct cle44f10
struct vm_area_struct {
    vm_mm = 0xc2857750,
    vm_start = 0x8048000,
    vm_end = 0x80a5000,
    vm_next = 0xcle44a10,
    vm_page_prot = {
        pgprot = 0x25
    },
    vm_flags = 0x1875,
    vm_avl_height = 0x2,
    vm_avl_left = 0xc30fe200,
    vm_avl_right = 0xc30fed00,
    vm_next_share = 0x0,
```



```

4083     vm_pprev_share = 0xc1e44a30,
4084     vm_ops = 0xc0215ca0,
4085     vm_offset = 0x0,
4086     vm_file = 0xc0bfdc70,
4087     vm_pte = 0
4088 }
4089

```

Display the definition and size of a vm_area_struct structure. This first example below displays just the structure and size. The second example uses the -o option to also display member offsets. Both examples were run with the output radix set to 10 (decimal):

```

4094
4095 crash> struct vm_area_struct
4096 struct vm_area_struct {
4097     struct mm_struct *vm_mm;
4098     long unsigned int vm_start;
4099     long unsigned int vm_end;
4100     struct vm_area_struct *vm_next;
4101     pgprot_t vm_page_prot;
4102     short unsigned int vm_flags;
4103     short int vm_avl_height;
4104     struct vm_area_struct *vm_avl_left;
4105     struct vm_area_struct *vm_avl_right;
4106     struct vm_area_struct *vm_next_share;
4107     struct vm_area_struct **vm_pprev_share;
4108     struct vm_operations_struct *vm_ops;
4109     long unsigned int vm_offset;
4110     struct file *vm_file;
4111     long unsigned int vm_pte;
4112 }
4113 SIZE: 56
4114
4115 crash> struct vm_area_struct -o
4116 struct vm_area_struct {
4117     [0] struct mm_struct *vm_mm;
4118     [4] long unsigned int vm_start;
4119     [8] long unsigned int vm_end;
4120     [12] struct vm_area_struct *vm_next;
4121     [16] pgprot_t vm_page_prot;
4122     [20] short unsigned int vm_flags;
4123     [22] short int vm_avl_height;
4124     [24] struct vm_area_struct *vm_avl_left;
4125     [28] struct vm_area_struct *vm_avl_right;
4126     [32] struct vm_area_struct *vm_next_share;
4127     [36] struct vm_area_struct **vm_pprev_share;
4128     [40] struct vm_operations_struct *vm_ops;
4129     [44] long unsigned int vm_offset;
4130     [48] struct file *vm_file;
4131     [52] long unsigned int vm_pte;
4132 }
4133 SIZE: 56
4134

```

Display the definition and offset of the pgd member of an mm_struct:

```

4135
4136 crash> struct mm_struct.pgd
4137 struct mm_struct {
4138     [80] pgd_t *pgd;
4139 }
4140

```

Display the pgd member of the mm_struct at address ffff810022e7d080:

```

4141
4142 crash> struct mm_struct.pgd ffff810022e7d080
4143 pgd = 0xffff81000e3ac000
4144

```

Display the pgd_t pointed to by the mm_struct.pgd pointer above, forcing the output to be expressed in hexadecimal:

```

4145
4146 crash> mm_struct.pgd ffff810022e7d080 -px
4147 pgd_t *pgd = 0xffff81000e3ac000
4148 -> {
4149     pgd = 0x2c0a6067
4150 }
4151

```

4156 Display the thread_info structure pointed to by the thread_info
4157 member of the task_struct at ffff8100181190c0:
4158

```
4159 crash> task_struct.thread_info ffff8100181190c0 -p
4160 struct thread_info *thread_info = 0xffff810023c06000
4161 -> {
4162     task = 0xffff8100181190c0,
4163     exec_domain = 0xffffffff802f78e0,
4164     flags = 128,
4165     status = 1,
4166     cpu = 3,
4167     preempt_count = 0,
4168     addr_limit = {
4169         seg = 18446604435732824064
4170     },
4171     restart_block = {
4172         fn = 0xffffffff80095a52 <do_no_restart_syscall>,
4173         arg0 = 0,
4174         arg1 = 0,
4175         arg2 = 0,
4176         arg3 = 0
4177     }
4178 }
```

4179
4180 Display the flags and virtual members of 4 contiguous page structures
4181 in the mem_map page structure array:
4182

```
4183 crash> page.flags,virtual c101196c 4
4184 flags = 0x8000,
4185 virtual = 0xc04b0000
4186
4187 flags = 0x8000,
4188 virtual = 0xc04b1000
4189
4190 flags = 0x8000,
4191 virtual = 0xc04b2000
4192
4193 flags = 0x8000,
4194 virtual = 0xc04b3000
4195
```

4196 Display the array of tcp_sl_timer structures declared by tcp_sl_array[]:
4197

```
4198 crash> struct tcp_sl_timer tcp_sl_array 4
4199 struct tcp_sl_timer {
4200     count = {
4201         counter = 0x0
4202     },
4203     period = 0x32,
4204     last = 0x1419e4,
4205     handler = 0xc0164854 <tcp_syn_recv_timer>
4206 }
4207 struct tcp_sl_timer {
4208     count = {
4209         counter = 0x2
4210     },
4211     period = 0x753,
4212     last = 0x14a6df,
4213     handler = 0xc01645b0 <tcp_keepalive>
4214 }
4215 struct tcp_sl_timer {
4216     count = {
4217         counter = 0x0
4218     },
4219     period = 0x2ee,
4220     last = 0x143134,
4221     handler = 0xc016447c <tcp_twkill>
4222 }
4223 struct tcp_sl_timer {
4224     count = {
4225         counter = 0x0
4226     },
4227     period = 0x64,
4228     last = 0x143198,
```

```
4229     handler = 0xc0164404 <tcp_bucketgc>
4230 }
4231
```

Without using the "struct" command name, display the the "d_child" list_head member from a dentry structure:

```
4234
4235 crash> dentry.d_child 0xe813cb4
4236     d_child = {
4237         next = 0x3661344,
4238         prev = 0xdea4bc4
4239     },
4240
```

Display the child dentry structure referenced by the "next" pointer above. Since the "next" address of 0x3661344 above is a pointer to an embedded list_head structure within the child dentry structure, the -l option is required:

```
4245
4246 crash> dentry -l dentry.d_child 0x3661344
4247 struct dentry {
4248     d_count = {
4249         counter = 1
4250     },
4251     d_flags = 0,
4252     d_inode = 0xf9aa604,
4253     d_parent = 0x11152b1c,
4254     d_hash = {
4255         next = 0x11fb3fc0,
4256         prev = 0x11fb3fc0
4257     },
4258     d_lru = {
4259         next = 0x366133c,
4260         prev = 0x366133c
4261     },
4262     d_child = {
4263         next = 0x36613cc,
4264         prev = 0xe813cd4
4265     },
4266     d_subdirs = {
4267         next = 0x366134c,
4268         prev = 0x366134c
4269     },
4270     d_alias = {
4271         next = 0xf9aa614,
4272         prev = 0xf9aa614
4273     },
4274     d_mounted = 0,
4275     d_name = {
4276         name = 0x3661384 "boot.log",
4277         len = 8,
4278         hash = 1935169207
4279     },
4280     d_time = 1515870810,
4281     d_op = 0x0,
4282     d_sb = 0x11fc9c00,
4283     d_vfs_flags = 0,
4284     d_fsdata = 0x0,
4285     d_extra_attributes = 0x0,
4286     d_iname = "boot.log\000"
4287 }
4288
```

Display the virtual address of each member of the task_struct at ffff8100145d2080:

```
4291
4292 crash> task_struct -o ffff8100145d2080
4293 struct task_struct {
4294     [ffff8100145d2080] volatile long int state;
4295     [ffff8100145d2088] struct thread_info *thread_info;
4296     [ffff8100145d2090] atomic_t usage;
4297     [ffff8100145d2098] long unsigned int flags;
4298     [ffff8100145d20a0] int lock_depth;
4299     [ffff8100145d20a4] int load_weight;
4300     [ffff8100145d20a8] int prio;
4301     [ffff8100145d20ac] int static_prio;

```

```

4302     [ffff8100145d20b0] int normal_prio;
4303     [ffff8100145d20b8] struct list_head run_list;
4304     [ffff8100145d20c8] struct prio_array *array;
4305     ...
4306
4307 Display the embedded sched_entity structure's on_rq member and
4308 the third pid_link structure in the embedded pids[] array of the
4309 task_struct at ffff88011653e250:
4310
4311 crash> task_struct.se.on_rq,pids[2] ffff88011653e250
4312     se.on_rq = 1,
4313     pids[2] = {
4314         node = {
4315             next = 0xffff88011653aff0,
4316             pprev = 0xffff88011653a860
4317         },
4318         pid = 0xffff88010d07ed00
4319     }
4320
4321 For an example of displaying per-cpu variables, consider the
4322 struct hd_struct.dkstats member, which is a percpu pointer to
4323 a disk_stats structure:
4324
4325 crash> struct hd_struct.dkstats
4326 struct hd_struct {
4327     [1232] struct disk_stats *dkstats;
4328 }
4329
4330 Taking an hd_struct at address ffff8802450e2848, display all
4331 of the per-cpu disk_stats structures that it references:
4332
4333 crash> struct hd_struct.dkstats ffff8802450e2848
4334     dkstats = 0x60fdb48026c8
4335 crash> struct disk_stats 0x60fdb48026c8:a
4336 [0]: ffffe8fefe6026c8
4337 struct disk_stats {
4338     sectors = {451376, 80468},
4339     ios = {6041, 971},
4340     merges = {386, 390},
4341     ticks = {194877, 56131},
4342     io_ticks = 12371,
4343     time_in_queue = 309163
4344 }
4345 [1]: ffffe8fefe8026c8
4346 struct disk_stats {
4347     sectors = {0, 0},
4348     ios = {0, 0},
4349     merges = {7, 242},
4350     ticks = {0, 0},
4351     io_ticks = 23,
4352     time_in_queue = 581
4353 }
4354 [2]: ffffe8fefe026c8
4355 struct disk_stats {
4356     sectors = {0, 0},
4357     ios = {0, 0},
4358     merges = {4, 112},
4359     ticks = {0, 0},
4360     io_ticks = 11,
4361     time_in_queue = 305
4362 }
4363 [3]: ffffe8fefec026c8
4364 struct disk_stats {
4365     sectors = {0, 0},
4366     ios = {0, 0},
4367     merges = {5, 54},
4368     ticks = {0, 0},
4369     io_ticks = 17,
4370     time_in_queue = 41
4371 }
4372
4373
4374 NOTE

```

If the structure name does not conflict with any crash command name, the "struct" command may be dropped. Accordingly, the examples above could also have been accomplished like so:

```
crash> vm_area_struct c1e44f10
crash> vm_area_struct
crash> vm_area_struct -o
crash> mm_struct.pgd fffff810022e7d080
crash> mm_struct.pgd
crash> tcp_sl_timer tcp_sl_array 4
```

Lastly, the short-cut "*" pointer-to command may also be used to negate the need to enter the "struct" command name (enter "help *" for details).

NAME

waitq - list tasks queued on a wait queue

SYNOPSIS

waitq [symbol] | [struct.member struct_addr] | [address]

DESCRIPTION

This command walks the wait queue list displaying the tasks which are blocked on the specified wait queue. The command differentiates between the old- and new-style wait queue structures used by the kernel. It can be invoked with the following argument types:

symbol	a global symbol of a wait queue.
struct.member struct_addr	a structure name and wait queue member combination followed by the structure's hexadecimal address.
address	a hexadecimal wait queue pointer.

EXAMPLES

Find out if any tasks are blocked on the "buffer_wait" wait queue:

```
crash> waitq buffer_wait
wait queue "buffer_wait" (c02927f0) is empty
```

See who is blocked on the "wait_chldexit" queue of task c5496000:

```
crash> waitq task_struct.wait_chldexit c5496000
PID: 30879 TASK: c5496000 CPU: 0 COMMAND: "bash"
```

Display the task list waiting on a known task queue:

```
crash> waitq c3534098
PID: 13691 TASK: c3534000 CPU: 1 COMMAND: "bash"
```

NAME

btop - bytes to page

SYNOPSIS

btop address ...

DESCRIPTION

This command translates a hexadecimal address to its page number.

EXAMPLES

```
crash> btop 512a000
512a000: 512a
```

NAME

ipcs - System V IPC facilities

SYNOPSIS

ipcs [-smMq] [-n pid|task] [id | addr]

DESCRIPTION

This command provides information on the System V IPC facilities. With no arguments, the command will display kernel usage of all three facilities.

```

4448         -s show semaphore arrays.
4449         -m show shared memory segments.
4450         -M show shared memory segments with additional details.
4451         -q show message queues.
4452         id show the data associated with this resource ID.
4453         addr show the data associated with this virtual address of a
4454             shm, sem_array or msg_queue.
4455
4456

```

For kernels supporting namespaces, the -n option may be used to display the IPC facilities with respect to the namespace of a specified task:

```

4461     -n pid    a process PID.
4462     -n task   a hexadecimal task_struct pointer.
4463

```

EXAMPLES

Display all IPC facilities:

```

4467 crash> ipcs
4468 SHMID_KERNEL      KEY          SHMID      UID      PERMS  BYTES      NATTCH STATUS
4469 ffff880473a28310 00000000 0          0        666    90000      1
4470 ffff880473a28490 00000001 32769      0        666    90000      1
4471 ffff880473a28250 00000002 65538      0        666    90000      1
4472
4473 SEM_ARRAY         KEY          SEMID      UID      PERMS  NSEMS
4474 ffff88047200f9d0 00000000 0          0        600    1
4475 ffff88046f826910 00000000 32769      0        600    1
4476
4477 MSG_QUEUE         KEY          MSQID      UID      PERMS  USED-BYTES  MESSAGES
4478 ffff8100036bb8d0 000079d7 0          3369    666    16640      104
4479 ffff8100036bb3d0 000079d8 32769      3369    666    12960      81
4480 ffff810026d751d0 000079d9 65538      3369    666    10880      68
4481

```

Display shared memory usage with detailed information:

```

4484 crash> ipcs -M
4485 SHMID_KERNEL      KEY          SHMID      UID      PERMS  BYTES      NATTCH STATUS
4486 ffff880473a28310 00000000 0          0        666    90000      1
4487 PAGES ALLOCATED/RESIDENT/SWAPPED: 22/1/0
4488 INODE: ffff88047239cd98
4489
4490 SHMID_KERNEL      KEY          SHMID      UID      PERMS  BYTES      NATTCH STATUS
4491 ffff880473a28490 00000001 32769      0        666    90000      1
4492 PAGES ALLOCATED/RESIDENT/SWAPPED: 22/1/0
4493 INODE: ffff88047239c118
4494
4495 SHMID_KERNEL      KEY          SHMID      UID      PERMS  BYTES      NATTCH STATUS
4496 ffff880473a28250 00000002 65538      0        666    90000      1
4497 PAGES ALLOCATED/RESIDENT/SWAPPED: 22/1/0
4498 INODE: ffff880470503758
4499

```

Display the shared memory data associated with shm, kernel ffff880473a28250:

```

4502 crash> ipcs -M ffff880473a28250
4503 SHMID_KERNEL      KEY          SHMID      UID      PERMS  BYTES      NATTCH STATUS
4504 ffff880473a28250 00000002 65538      0        666    90000      1
4505 PAGES ALLOCATED/RESIDENT/SWAPPED: 22/1/0
4506 INODE: ffff880470503758
4507

```

NAME

pte - translate a page table entry

SYNOPSIS

pte contents ...

DESCRIPTION

This command translates the hexadecimal contents of a PTE into its physical page address and page bit settings. If the PTE references a swap location, the swap device and offset are displayed.

EXAMPLES

```

4521
4522     crash> pte d8e067
4523     PTE      PHYSICAL  FLAGS
4524     d8e067    d8e000    (PRESENT|RW|USER|ACCESSED|DIRTY)
4525
4526     crash> pte 13f600
4527     PTE      SWAP      OFFSET
4528     13f600    /dev/hda2  5104
4529
4530
4531 NAME
4532     swap - swap device information
4533
4534 SYNOPSIS
4535     swap
4536
4537 DESCRIPTION
4538     This command displays information for each configured swap device.
4539
4540 EXAMPLE
4541     crash> swap
4542     SWAP_INFO_STRUCT  TYPE      SIZE      USED      PCT  PRI  FILENAME
4543     ffff880153d45f40  PARTITION  7192568k  1200580k  16%  -1  /dev/dm-1
4544
4545
4546 NAME
4547     whatis - search symbol table for data or type information
4548
4549 SYNOPSIS
4550     whatis [[-o] [struct | union | typedef | symbol]] |
4551           [[-r [size|range]] [-m member]]
4552
4553 DESCRIPTION
4554     This command displays the definition of structures, unions, typedefs or
4555     text/data symbols:
4556
4557     struct    a structure name. The output is the same as if the "struct"
4558                command was used.
4559     union     a union name. The output is the same as if the "union" command
4560                was used.
4561     -o        display the offsets of structure/union members.
4562     typedef   a typedef name. If the typedef translates to a structure or union
4563                the output is the same as if the "struct" or "union" command
4564                was used. If the typedef is a primitive datatype, the one-line
4565                declaration is displayed.
4566     symbol    a kernel symbol.
4567
4568     Alternatively, a search can be made for data structures of a given size or
4569     size range, that contain a member of a given type, or contain a pointer to
4570     given type. The -r and -m options may be used alone or in conjunction with
4571     one another:
4572
4573     -r size    search for structures of this exact size.
4574     -r range   search for structures of a range of sizes, expressed as "low-high".
4575     -m member  search for structures that contain a member of this data type, or
4576                that contain a pointer to this data type; if a structure contains
4577                another structure, the members of the embedded structure will also
4578                be subject to the search. The member argument may also be expressed
4579                as a substring of a member's data type.
4580
4581 EXAMPLES
4582     Display the definition of a linux_binfmt structure:
4583
4584     crash> whatis linux_binfmt
4585     struct linux_binfmt {
4586         struct list_head lh;
4587         struct module *module;
4588         int (*load_binary)(struct linux_binprm *);
4589         int (*load_shlib)(struct file *);
4590         int (*core_dump)(struct coredump_params *);
4591         unsigned long min_coredump;
4592     }
4593     SIZE: 56

```

Display the same structure with member offsets:

```
crash> whatis -o linux_binfmt
struct linux_binfmt {
    [0] struct list_head lh;
    [16] struct module *module;
    [24] int (*load_binary)(struct linux_binprm *);
    [32] int (*load_shlib)(struct file *);
    [40] int (*core_dump)(struct coredump_params *);
    [48] unsigned long min_coredump;
}
SIZE: 56
```

Since a `kmem_bufctl_t` is typedef'd to be a `kmem_bufctl_s` structure, the output of the following two commands is identical:

```
crash> whatis kmem_bufctl_s
struct kmem_bufctl_s {
    union {
        struct kmem_bufctl_s *buf_nextp;
        kmem_slab_t *buf_slabp;
        void *buf_objp;
    } u;
};

crash> whatis kmem_bufctl_t
struct kmem_bufctl_s {
    union {
        struct kmem_bufctl_s *buf_nextp;
        kmem_slab_t *buf_slabp;
        void *buf_objp;
    } u;
};
SIZE: 4 (0x4)
```

Display the type data of `sys_read()` and jiffies text and data symbols:

```
crash> whatis sys_read
ssize_t sys_read(unsigned int, char *, size_t);

crash> whatis jiffies
long unsigned int jiffies;
```

Display definition of a `kdev_t` typedef:

```
crash> whatis kdev_t
typedef short unsigned int kdev_t;
SIZE: 2 (0x2)
```

Display all structures which have a size of 192 bytes:

```
crash> whatis -r 192
SIZE  TYPE
192   _intel_private
192   blkcg_gq
192   clock_event_device
192   cper_sec_proc_generic
192   dentry
192   dst_ops
192   ehci_itd
192   ethtool_rxnfc
192   fb_ops
192   file_lock
192   inode_operations
192   input_device_id
192   ip_vs_stats
192   numa_group
192   parallel_data
192   pcie_port_service_driver
192   pebs_record_hsw
192   pnp_driver
192   regmap_config
```



```
4667      192 sched_entity
4668      192 tcp_timewait_sock
4669      192 timerfd_ctx
4670      192 tpm_vendor_specific
4671      192 urb
```

```
4672
4673 Display all structures that contain members that point to
4674 an mm_struct:
```

```
4675
4676 crash> whatis -m mm_struct
4677 SIZE  TYPE
4678    16  tlb_state
4679    24  flush_tlb_info
4680    24  ftrace_raw_xen_mmu_pgd
4681    24  futex_key
4682    24  map_info
4683    32  ftrace_raw_xen_mmu_alloc_ptpage
4684    32  ftrace_raw_xen_mmu_pte_clear
4685    40  ftrace_raw_xen_mmu_flush_tlb_others
4686    40  ftrace_raw_xen_mmu_ptep_modify_prot
4687    40  ftrace_raw_xen_mmu_set_pte_at
4688    40  mm_slot
4689    64  mm_walk
4690    64  rmap_item
4691   104  userfaultfd_ctx
4692   128  mmu_gather
4693   216  vm_area_struct
4694   256  linux_binprm
4695  2616  rq
4696  2936  task_struct
```

```
4697
4698 Display all structures sized from 256 to 512 bytes that
4699 contain members that point to a task_struct:
```

```
4700
4701 crash> whatis -r 256-512 -m task_struct
4702 SIZE  TYPE
4703   256  file
4704   256  od_cpu_dbs_info_s
4705   264  srcu_notifier_head
4706   272  protection_domain
4707   288  clk_notifier
4708   288  fsnotify_group
4709   296  quota_info
4710   312  tty_port
4711   320  workqueue_struct
4712   344  trace_array
4713   344  uart_state
4714   352  cpufreq_policy
4715   352  elf_thread_core_info
4716   376  perf_event_context
4717   384  rcu_data
4718   400  cgroup
4719   408  subsys_private
4720   424  hvc_struct
4721   496  psmouse
```

```
4722
4723 NAME
```

```
4724     dev - device data
```

```
4725
4726 SYNOPSIS
```

```
4727     dev [-i | -p | -d | -D ] [-V | -v index [file]]
```

```
4728
4729 DESCRIPTION
```

```
4730     If no argument is entered, this command dumps character and block
4731     device data.
```

```
4732
4733
4734     -i  display I/O port usage; on 2.4 kernels, also display I/O memory usage.
4735     -p  display PCI device data.
4736     -d  display disk I/O statistics:
4737         TOTAL: total number of allocated in-progress I/O requests
4738         SYNC: I/O requests that are synchronous
4739         ASYNC: I/O requests that are asynchronous
```

4740 READ: I/O requests that are reads (older kernels)
 4741 WRITE: I/O requests that are writes (older kernels)
 4742 DRV: I/O requests that are in-flight in the device driver.
 4743 If the device driver uses blk-mq interface, this field
 4744 shows N/A(MQ). If not available, this column is not shown.
 4745 -D same as -d, but filter out disks with no in-progress I/O requests.
 4746

4747 If the dumpfile contains device dumps:

- 4748 -V display an indexed list of all device dumps present in the vmcore,
- 4749 showing their file offset, size and name.
- 4750 -v index select and display one device dump based upon an index value
- 4751 shown by the -V option, shown in a default human-readable format;
- 4752 alternatively, the "rd -f" option along with its various format
- 4753 options may be used to further tailor the output.
- 4754 file only used with -v, copy the device dump data to a file.
- 4755

4756 EXAMPLES

4757 Display character and block device data:

```
4759 crash> dev
4760 CHRDEV NAME CDEV OPERATIONS
4761 1 mem f79b83c0 memory_fops
4762 4 /dev/vc/0 c07bc560 console_fops
4763 4 tty f7af5004 tty_fops
4764 4 ttyS f7b02204 tty_fops
4765 5 /dev/tty c07bc440 tty_fops
4766 5 /dev/console c07bc4a0 console_fops
4767 5 /dev/ptmx c07bc500 ptmx_fops
4768 6 lp c5797e40 lp_fops
4769 7 vcs f7b03d40 vcs_fops
4770 10 misc f7f68640 misc_fops
4771 13 input f79b8840 input_fops
4772 21 sg f7f12840 sg_fops
4773 29 fb f7f8c640 fb_fops
4774 128 ptm f7b02604 tty_fops
4775 136 pts f7b02404 tty_fops
4776 162 raw c0693e40 raw_fops
4777 180 usb f79b8bc0 usb_fops
4778 189 usb_device c06a0300 usbfs_device_file_operations
4779 216 rfcomm f5961a04 tty_fops
4780 254 pcmcia f79b82c0 ds_fops
```

```
4782 BLKDEV NAME GENDISK OPERATIONS
4783 1 ramdisk f7b23480 rd_bd_op
4784 8 sd f7cab280 sd_fops
4785 9 md f7829b80 md_fops
4786 11 sr f75c24c0 sr_bdops
4787 65 sd (none)
4788 66 sd (none)
4789 67 sd (none)
4790 68 sd (none)
4791 69 sd (none)
4792 70 sd (none)
4793 71 sd (none)
4794 128 sd (none)
4795 129 sd (none)
4796 130 sd (none)
4797 131 sd (none)
4798 132 sd (none)
4799 133 sd (none)
4800 134 sd (none)
4801 135 sd (none)
4802 253 device-mapper c57a0ac0 dm_blk_dops
4803 254 mdp (none)
```

4805 Display PCI data:

```
4807 crash> dev -p
4808 PCI_DEV BU:SL.FN CLASS: VENDOR-DEVICE
4809 c00051c0 00:00.0 Host bridge: Intel 440BX - 82443BX Host
4810 c0005250 00:01.0 PCI bridge: Intel 440BX - 82443BX AGP
4811 c00052e0 00:07.0 ISA bridge: Intel 82371AB PIIX4 ISA
4812 c0005370 00:07.1 IDE interface: Intel 82371AB PIIX4 IDE
```

```
4813 c0005400 00:07.2 USB Controller: Intel 82371AB PIIX4 USB
4814 c0005490 00:07.3 Bridge: Intel 82371AB PIIX4 ACPI
4815 c0005520 00:11.0 Ethernet controller: 3Com 3C905B 100bTX
4816 c00055b0 00:13.0 PCI bridge: DEC DC21152
4817 c0005640 01:00.0 VGA compatible controller: NVidia [PCI_DEVICE 28]
4818 c00056d0 02:0a.0 SCSI storage controller: Adaptec AIC-7890/1
4819 c0005760 02:0e.0 SCSI storage controller: Adaptec AIC-7880U
4820
```

4821 Display I/O port and I/O memory usage:

```
4822
4823 crash> dev -i
4824 RESOURCE      RANGE      NAME
4825 c03036d4 0000-ffff PCI IO
4826 c0302594 0000-001f dma1
4827 c03025b0 0020-003f pic1
4828 c03025cc 0040-005f timer
4829 c03025e8 0060-006f keyboard
4830 c0302604 0080-008f dma page reg
4831 c0302620 00a0-00bf pic2
4832 c030263c 00c0-00df dma2
4833 c0302658 00f0-00ff fpu
4834 c122ff20 0170-0177 idel
4835 c122f240 0213-0213 isapnp read
4836 c122ff40 02f8-02ff serial(auto)
4837 c122ff00 0376-0376 idel
4838 c03186e8 03c0-03df vga+
4839 c122ff60 03f8-03ff serial(auto)
4840 c123851c 0800-083f Intel Corporation 82371AB PIIX4 ACPI
4841 c1238538 0840-085f Intel Corporation 82371AB PIIX4 ACPI
4842 c122f220 0a79-0a79 isapnp write
4843 c122f200 0cf8-0cff PCI conf1
4844 c1238858 dc00-dc7f 3Com Corporation 3c905B 100BaseTX [Cyclone]
4845 c122fc00 dc00-dc7f 00:11.0
4846 c12380c8 dce0-dcff Intel Corporation 82371AB PIIX4 USB
4847 c1238d1c e000-ffff PCI Bus #02
4848 c1237858 e800-e8ff Adaptec AIC-7880U
4849 c1237458 ec00-ecff Adaptec AHA-2940U2/W / 7890
4850 c1239cc8 ffa0-ffaf Intel Corporation 82371AB PIIX4 IDE
4851
4852 RESOURCE      RANGE      NAME
4853 c03036f0 00000000-ffffffff PCI mem
4854 c0004000 00000000-0009ffff System RAM
4855 c03026ac 000a0000-000bffff Video RAM area
4856 c03026fc 000c0000-000c7fff Video ROM
4857 c0302718 000c9800-000cdfff Extension ROM
4858 c0302734 000ce000-000ce7ff Extension ROM
4859 c0302750 000ce800-000cffff Extension ROM
4860 c03026e0 000f0000-000fffff System ROM
4861 c0004040 00100000-07ffdcfff System RAM
4862 c0302674 00100000-0028682b Kernel code
4863 c0302690 0028682c-0031c63f Kernel data
4864 c0004060 07ffe000-07ffffff reserved
4865 c1239058 ec000000-efffffff Intel Corporation 440BX/ZX - 82443BX/ZX Host
4866 bridge
4867 c1238d54 f1000000-f1ffffff PCI Bus #02
4868 c1239554 f2000000-f5ffffff PCI Bus #01
4869 c1237074 f4000000-f5ffffff nVidia Corporation Riva TnT2 [NV5]
4870 c1238d38 fa000000-fbffffff PCI Bus #02
4871 c1237874 faffe000-faffefff Adaptec AIC-7880U
4872 c127ec40 faffe000-faffefff aic7xxx
4873 c1237474 fafff000-faffffff Adaptec AHA-2940U2/W / 7890
4874 c127eec0 fafff000-faffffff aic7xxx
4875 c1239538 fc000000-fdffffff PCI Bus #01
4876 c1237058 fc000000-fcffffff nVidia Corporation Riva TnT2 [NV5]
4877 c1238874 fe000000-fe00007f 3Com Corporation 3c905B 100BaseTX [Cyclone]
4878 c0004080 fec00000-fec0ffff reserved
4879 c00040a0 fee00000-fee0ffff reserved
4880 c00040c0 ffe00000-ffffffff reserved
4881
```

4882 Display disk I/O statistics:

```
4883
4884 crash> dev -d
4885 MAJOR GENDISK      NAME      REQUEST_QUEUE     TOTAL  READ WRITE  DRV
```

4886	2	ffff81012d8a5000	fd0	ffff81012dc053c0	12	0	12	0
4887	22	ffff81012dc6b000	hdc	ffff81012d8ae340	2	2	0	0
4888	8	ffff81012dd71000	sda	ffff81012d8af040	6	0	6	6
4889	8	ffff81012dc77000	sdb	ffff81012d8b5740	0	0	0	0
4890	8	ffff81012d8d0c00	sd0	ffff81012d8ae9c0	0	0	0	0

Display the available device dumps:

```
crash> dev -V
INDEX  OFFSET          SIZE          NAME
0      0x240         33558464     cxgb4_0000:02:00.4
1      0x2001240    33558464     cxgb4_0000:03:00.4
```

Extract a specified device dump to file:

```
crash> dev -v 0 device_dump_0.bin
DEVICE: cxgb4_0000:02:00.4
33558464 bytes copied from 0x240 to device_dump_0.bin
```

Format and display a device's dump data to the screen using the "rd" command:

```
crash> rd -f 0x240 -32 8
240: 040b69e2 00000038 000e0001 00675fd4 .i..8....._g.
250: 00000000 21600047 00000000 00000000 ....G.`!.....
```

Display a device's dump data to the screen using the default format:

```
crash> dev -v 1
DEVICE: cxgb4_0000:03:00.4
2001240: 00000038040b69e2 00af985c000e0001 .i..8.....\...
2001250: 2150004700000000 0000000000000000 ....G.P!.....
2001260: 0000000000000000 0000000000000000 .....
2001270: 0000000000000000 0002fccc00000001 .....
2001280: 000000000000027b0 0000000000000000 .'.....
...
```

NAME

irq - IRQ data

SYNOPSIS

```
irq [[[index ...] | -u ] | -d | -b | -a | -s [-c cpu]]
```

DESCRIPTION

This command collaborates the data in an `irq_desc_t`, along with its associated `hw_interrupt_type` and `irqaction` structure data, into a consolidated per-IRQ display. For kernel versions 2.6.37 and later the display consists of the `irq_desc/irq_data` address, its `irqaction` address(es), and the `irqaction` name strings. Alternatively, the intel interrupt descriptor table, bottom half data, cpu affinity for in-use irqs, or kernel irq stats may be displayed. If no index value argument(s) nor any options are entered, the IRQ data for all IRQs will be displayed.

```
index    a valid IRQ index.
-u       dump data for in-use IRQs only.
-d       dump the intel interrupt descriptor table.
-b       dump bottom half data.
-a       dump cpu affinity for in-use IRQs.
-s       dump the kernel irq stats; if no cpu specified with -c, the
         irq stats of all cpus will be displayed.
-c cpu   only usable with the -s option, dump the irq stats of the
         specified cpu[s]; cpu can be specified as "1,3,5", "1-3",
         "1,3,5-7,10", "all", or "a" (shortcut for "all").
```

EXAMPLES

Display the relevant data for IRQ 18 from a pre-2.6.37 kernel:

```
crash> irq 18
IRQ: 18
STATUS: 0
HANDLER: c02301e0 <ioapic_level_irq_type>
typename: c01f9e0c "IO-APIC-level"
```

```

4959         startup: c0110234 <unmask_IO_APIC_irq>
4960         shutdown: c01101cc <mask_IO_APIC_irq>
4961         handle: c0110518 <do_level_ioapic_IRQ>
4962         enable: c0110234 <unmask_IO_APIC_irq>
4963         disable: c01101cc <mask_IO_APIC_irq>
4964     ACTION: c009c6b0
4965         handler: c01ce818 <do_aic7xxx_isr>
4966         flags: 4000000 (SA_SHIRQ)
4967         mask: 0
4968         name: c0217780 "aic7xxx"
4969         dev_id: c0090078
4970         next: c009c770
4971     ACTION: c009c770
4972         handler: c01ce818 <do_aic7xxx_isr>
4973         flags: 4000000 (SA_SHIRQ)
4974         mask: 0
4975         name: c0217780 "aic7xxx"
4976         dev_id: c0091078
4977         next: 0
4978     DEPTH: 0

```

Display the relevant data for IRQ 21 from a 2.6.37 kernel:

```

4981     crash> irq 21
4982     IRQ    IRQ_DESC/ DATA    IRQACTION    NAME
4983     21     ffff88003787f780    ffff8800379a8b40    "ehci_hcd:usb2"
4984                                     ffff8800379cbac0    "uhci_hcd:usb5"
4985                                     ffff8800379cb140    "uhci_hcd:usb7"

```

Display the intel interrupt descriptor table entries:

```

4990     crash> irq -d
4991     [0] divide_error
4992     [1] debug
4993     [2] nmi
4994     [3] int3
4995     [4] overflow
4996     [5] bounds
4997     [6] invalid_op
4998     [7] device_not_available
4999     [8] double_fault
5000     [9] coprocessor_segment_overrun
5001     [10] invalid_TSS
5002     [11] segment_not_present
5003     [12] stack_segment
5004     [13] general_protection
5005     [14] page_fault
5006     [15] spurious_interrupt_bug
5007     [16] coprocessor_error
5008     [17] alignment_check
5009     [18] ignore_int
5010     [19] ignore_int
5011     [20] ignore_int
5012     [21] ignore_int
5013     ...
5014
5015     [250] IRQ0xda_interrupt
5016     [251] IRQ0xdb_interrupt
5017     [252] IRQ0xdc_interrupt
5018     [253] IRQ0xdd_interrupt
5019     [254] IRQ0xde_interrupt
5020     [255] spurious_interrupt

```

Display the bottom half data:

```

5024     crash> irq -b
5025     SOFTIRQ_VEC    ACTION
5026     [0]            ffffffff81068f60    <tasklet_hi_action>
5027     [1]            ffffffff81071b80    <run_timer_softirq>
5028     [2]            ffffffff813e6f30    <net_tx_action>
5029     [3]            ffffffff813ee370    <net_rx_action>
5030     [4]            ffffffff81211a60    <blk_done_softirq>
5031     [5]            ffffffff812122f0    <blk_iopoll_softirq>

```

```

5032         [6]      ffffffff81069090 <tasklet_action>
5033         [7]      ffffffff81058830 <run_rebalance_domains>
5034         [8]      ffffffff81087f00 <run_hrtimer_softirq>
5035         [9]      ffffffff810ca7a0 <rcu_process_callbacks>
5036
5037 Display the cpu affinity for in-use IRQs:
5038
5039 crash> irq -a
5040      IRQ NAME                                AFFINITY
5041      0 timer                                0-23
5042      1 i8042                                0-23
5043      8 rtc0                                  0-23
5044      9 acpi                                  0-23
5045      16 ehci_hcd:usb2,uhci_hcd:usb3,uhci_hcd:usb6 0,6,18
5046      17 uhci_hcd:usb4,uhci_hcd:usb7 0-23
5047      18 ehci_hcd:usb1,uhci_hcd:usb5,uhci_hcd:usb8,ioc0 0,11,23
5048      24 dmar0                                0
5049      35 pciehp                                0-23
5050      36 pciehp                                0-23
5051      37 pciehp                                0-23
5052      38 pciehp                                0-23
5053      39 megasas                                0-5,12-17
5054      40 lpfc:sp                                0-5,12-17
5055      41 lpfc:fp                                0,6-11,18-23
5056      42 lpfc:sp                                0,6-11,18-23
5057      43 lpfc:fp                                0,6-11,18-23
5058      ...
5059
5060      80 ioat-msix                                0-23
5061      81 ioat-msix                                0-23
5062      82 ioat-msix                                0-23
5063      83 ioat-msix                                0-23
5064      84 ioat-msix                                0-23
5065      85 ioat-msix                                0-23
5066      86 ioat-msix                                0-23
5067      87 ioat-msix                                0-23
5068      88 eth4                                    0,17
5069
5070 Display the kernel irq stats:
5071
5072 crash>irq -c 0,2 -s
5073      CPU0      CPU2
5074      0: 2068161471      0 IR-IO-APIC-edge      timer
5075      1:          9      0 IR-IO-APIC-edge      i8042
5076      8:          1      0 IR-IO-APIC-edge      rtc0
5077      9:          0      0 IR-IO-APIC-fasteoi    acpi
5078     16:         36      0 IR-IO-APIC-fasteoi    ehci_hcd:usb2
5079      ...
5080
5081     85:          3      0 IR-PCI-MSI-edge      ioat-msix
5082     86:          3      0 IR-PCI-MSI-edge      ioat-msix
5083     87:          3      0 IR-PCI-MSI-edge      ioat-msix
5084     88:         24     295 IR-PCI-MSI-edge      eth4
5085
5086 NAME
5087 ptob - page to bytes
5088
5089 SYNOPSIS
5090 ptob page_number ...
5091
5092 DESCRIPTION
5093 This command translates a page frame number to its byte value.
5094
5095 EXAMPLES
5096 crash> ptob 512a
5097 512a: 512a000
5098
5099 NAME
5100 sym - translate a symbol to its virtual address, or vice-versa
5101
5102 SYNOPSIS

```

```
5105     sym [-l] | [-M] | [-m module] | [-p|-n] | [-q string] | [symbol | vaddr]
```

```
5106
```

DESCRIPTION

```
5108     This command translates a symbol to its virtual address, or a static
5109     kernel virtual address to its symbol -- or to a symbol-plus-offset value,
5110     if appropriate.  Additionally, the symbol type is shown in parentheses,
5111     and if the symbol is a known text value, the file and line number are shown.
```

```
5112
```

```
5113         -l     dumps all symbols and their values.
```

```
5114         -M     dumps the current set of module symbols.
```

```
5115     -m module  dumps the current set of symbols for a specified module.
```

```
5116         -p     display the target symbol and the previous symbol.
```

```
5117         -n     display the target symbol and the next symbol.
```

```
5118     -q string  searches for all symbols containing "string".
```

```
5119         symbol  a kernel text or data symbol.
```

```
5120         vaddr   a kernel virtual address.
```

```
5121
```

```
5122     If the "symbol", "vaddr" or "string" argument resolves to a module
5123     symbol, then the module name will be displayed in brackets following the
5124     symbol value.
```

```
5125
```

EXAMPLES

```
5127     Translate data symbol jiffies to its value, and vice-versa:
```

```
5128
```

```
5129     crash> sym jiffies
```

```
5130     c0213ae0 (D) jiffies
```

```
5131
```

```
5132     crash> sym c0213ae0
```

```
5133     c0213ae0 (D) jiffies
```

```
5134
```

```
5135     Translate a text address to its symbolic value and source file:
```

```
5136
```

```
5137     crash> sym c0109944
```

```
5138     c0109944 (T) system_call+0x34  ../linux-2.2.5/arch/i386/kernel/signal.c: 723
```

```
5139
```

```
5140     Dump the whole symbol table:
```

```
5141
```

```
5142     crash> sym -l
```

```
5143     c0100000 (T) _stext
```

```
5144     c0100000 (A) _text
```

```
5145     c0100000 (t) startup_32
```

```
5146     c0100000 (T) stext
```

```
5147     c01000a4 (t) checkCPUtype
```

```
5148     c0100139 (t) is486
```

```
5149     c0100148 (t) is386
```

```
5150     c01001b1 (t) L6
```

```
5151     c01001b3 (t) ready
```

```
5152     c01001b4 (t) check_x87
```

```
5153     c01001da (t) setup_idt
```

```
5154     c01001f7 (t) rp_sidt
```

```
5155     c0100204 (T) stack_start
```

```
5156     c010020c (t) int_msg
```

```
5157     c0100220 (t) ignore_int
```

```
5158     c0100242 (t) idt_descr
```

```
5159     c0100244 (T) idt
```

```
5160     c010024a (t) gdt_descr
```

```
5161     c010024c (T) gdt
```

```
5162     c0101000 (T) swapper_pg_dir
```

```
5163     c0102000 (T) pg0
```

```
5164     c0103000 (T) empty_bad_page
```

```
5165     c0104000 (T) empty_bad_page_table
```

```
5166     c0105000 (T) empty_zero_page
```

```
5167     ...
```

```
5168
```

```
5169     Find all symbols containing the string "pipe":
```

```
5170
```

```
5171     crash> sym -q pipe
```

```
5172     c010ec60 (T) sys_pipe
```

```
5173     c012f660 (t) pipe_read
```

```
5174     c012f7b8 (t) pipe_write
```

```
5175     c012f9c0 (t) pipe_lseek
```

```
5176     c012f9d0 (t) bad_pipe_r
```

```
5177     c012f9dc (t) bad_pipe_w
```

```

5178 c012f9e8 (t) pipe_ioctl
5179 c012fa18 (t) pipe_poll
5180 c012fb00 (t) pipe_release
5181 c012fb48 (t) pipe_read_release
5182 c012fb5c (t) pipe_write_release
5183 c012fb70 (t) pipe_rdwr_release
5184 c012fba0 (t) pipe_read_open
5185 c012fbb0 (t) pipe_write_open
5186 c012fbc0 (t) pipe_rdwr_open
5187 c012fbec (t) get_pipe_inode
5188 c012fcc4 (T) do_pipe
5189 c023a920 (D) read_pipe_fops
5190 c023a960 (D) write_pipe_fops
5191 c023a9a0 (D) rdwr_pipe_fops
5192 c023a9e0 (D) pipe_inode_operations
5193

```

Dump the symbols of the uart401 module, both before, and then after, the complete set of symbols are loaded with the "mod -s" command:

```

5194
5195
5196
5197 crash> sym -m uart401
5198 c8032000 MODULE START: uart401
5199 c8032138 (?) uart40lintr
5200 c803235c (?) attach_uart401
5201 c8032638 (?) probe_uart401
5202 c80326d4 (?) unload_uart401
5203 c8033770 MODULE END: uart401
5204 crash> mod -s uart401
5205 MODULE NAME SIZE OBJECT FILE
5206 c8032000 uart401 6000 /lib/modules/2.2.14/misc/uart401.o
5207 crash> sym -m uart401
5208 c8032000 MODULE START: uart401
5209 c8032050 (t) my_notifier_call
5210 c8032084 (t) uart401_status
5211 c8032098 (t) uart401_cmd
5212 c80320a8 (t) uart401_read
5213 c80320bc (t) uart401_write
5214 c80320cc (t) uart401_input_loop
5215 c8032138 (T) uart40lintr
5216 c8032168 (t) uart401_open
5217 c80321c8 (t) uart401_close
5218 c80321f4 (t) uart401_out
5219 c80322ac (t) uart401_start_read
5220 c80322b4 (t) uart401_end_read
5221 c80322bc (t) uart401_kick
5222 c80322c4 (t) uart401_buffer_status
5223 c80322cc (t) enter_uart_mode
5224 c803235c (T) attach_uart401
5225 c803259c (t) reset_uart401
5226 c8032638 (T) probe_uart401
5227 c80326d4 (T) unload_uart401
5228 c8032760 (T) init_module
5229 c80327cc (T) cleanup_module
5230 c8032b00 (d) sound_notifier
5231 c8032b0c (d) detected_devc
5232 c8032b20 (d) std_synth_info
5233 c8032bc0 (d) std_midi_synth
5234 c8033600 (d) uart401_operations
5235 c80336c4 (D) io
5236 c80336c8 (D) irq
5237 c80336e0 (b) hw_info.508
5238 c8033770 MODULE END: uart401
5239

```

Display the value of jiffies, along with the next and previous symbols:

```

5240
5241
5242 crash> sym -np jiffies
5243 c023027c (D) prof_shift
5244 c0230280 (D) jiffies
5245 c02302a0 (D) task
5246

```

Translate a symbol value to its name and module:

```

5247
5248
5249 crash> sym f88878d1
5250 f88878d1 (t) ext3_readdir [ext3]

```



```

5251     crash>
5252
5253
5254 NAME
5255     wr - write memory
5256
5257 SYNOPSIS
5258     wr [-u|-k|-p] [-8|-16|-32|-64] [address|symbol] value
5259
5260 DESCRIPTION
5261     This command modifies the contents of memory. The starting address may be
5262     entered either symbolically or by address. The default modification size
5263     is the size of a long data type. Write permission must exist on the
5264     /dev/mem. When writing to memory on a live system, this command should
5265     obviously be used with great care.
5266
5267     -u address argument is a user virtual address.
5268     -k address argument is a kernel virtual address.
5269     -p address argument is a physical address.
5270     -8 write data in an 8-bit value.
5271     -16 write data in a 16-bit value.
5272     -32 write data in a 32-bit values (default on 32-bit machines).
5273     -64 write data in a 64-bit values (default on 64-bit machines).
5274     address address to write. The address is considered virtual unless the
5275     -p option is used. If a virtual address is specified, the
5276     -u or -k options are necessary only if the address space cannot
5277     be determined from the address value itself. If a user virtual
5278     address is specified, the address space of the current context
5279     implied. The address must be expressed in hexadecimal format.
5280     symbol symbol of starting address to write.
5281     value the value of the data to write.
5282
5283 EXAMPLES
5284     Turn on a debug flag:
5285
5286     crash> wr my_debug_flag 1
5287
5288
5289 NAME
5290     dis - disassemble
5291
5292 SYNOPSIS
5293     dis [-rfludxs][-b [num]] [address | symbol | (expression)] [count]
5294
5295 DESCRIPTION
5296     This command disassembles source code instructions starting (or ending) at
5297     a text address that may be expressed by value, symbol or expression:
5298
5299     -r (reverse) displays all instructions from the start of the
5300     routine up to and including the designated address.
5301     -f (forward) displays all instructions from the given address
5302     to the end of the routine.
5303     -l displays source code line number data in addition to the
5304     disassembly output.
5305     -u address is a user virtual address in the current context;
5306     otherwise the address is assumed to be a kernel virtual address.
5307     If this option is used, then -r and -l are ignored.
5308     -x override default output format with hexadecimal format.
5309     -d override default output format with decimal format.
5310     -s displays the filename and line number of the source code that
5311     is associated with the specified text location, followed by a
5312     source code listing if it is available on the host machine.
5313     The line associated with the text location will be marked with
5314     an asterisk; depending upon gdb's internal "listsize" variable,
5315     several lines will precede the marked location. If a "count"
5316     argument is entered, it specifies the number of source code
5317     lines to be displayed after the marked location; otherwise
5318     the remaining source code of the containing function will be
5319     displayed.
5320     -b [num] modify the pre-calculated number of encoded bytes to skip after
5321     a kernel BUG ("ud2a") instruction; with no argument, displays
5322     the current number of bytes being skipped. (x86 and x86_64 only)
5323     address starting hexadecimal text address.

```

```

5324         symbol    symbol of starting text address.  On ppc64, the symbol
5325                  preceded by '.' is used.
5326 (expression) expression evaluating to a starting text address.
5327         count    the number of instructions to be disassembled (default is 1).
5328                  If no count argument is entered, and the starting address
5329                  is entered as a text symbol, then the whole routine will be
5330                  disassembled.  The count argument is supported when used with
5331                  the -r and -f options.
5332

```

5333 EXAMPLES

5334 Disassemble the sys_signal() routine without, and then with, line numbers:

```

5335
5336 crash> dis sys_signal
5337 0xc0112c88 <sys_signal>:      push    %ebp
5338 0xc0112c89 <sys_signal+1>:    mov     %esp,%ebp
5339 0xc0112c8b <sys_signal+3>:    sub     $0x28,%esp
5340 0xc0112c8e <sys_signal+6>:    mov     0xc(%ebp),%eax
5341 0xc0112c91 <sys_signal+9>:    mov     %eax,0xffffffffec(%ebp)
5342 0xc0112c94 <sys_signal+12>:   movl    $0xc0000000,0xfffffffff0(%ebp)
5343 0xc0112c9b <sys_signal+19>:   lea     0xffffffffd8(%ebp),%eax
5344 0xc0112c9e <sys_signal+22>:   push    %eax
5345 0xc0112c9f <sys_signal+23>:   lea     0xffffffffec(%ebp),%eax
5346 0xc0112ca2 <sys_signal+26>:   push    %eax
5347 0xc0112ca3 <sys_signal+27>:   pushl   0x8(%ebp)
5348 0xc0112ca6 <sys_signal+30>:   call    0xc01124b8 <do_sigaction>
5349 0xc0112cab <sys_signal+35>:   test    %eax,%eax
5350 0xc0112cad <sys_signal+37>:   jne     0xc0112cb2 <sys_signal+42>
5351 0xc0112caf <sys_signal+39>:   mov     0xffffffffd8(%ebp),%eax
5352 0xc0112cb2 <sys_signal+42>:   leave
5353 0xc0112cb3 <sys_signal+43>:   ret
5354
5355 crash> dis -l sys_signal
5356 /usr/src/linux-2.2.5/kernel/signal.c: 1074
5357 0xc0112c88 <sys_signal>:      push    %ebp
5358 0xc0112c89 <sys_signal+1>:    mov     %esp,%ebp
5359 0xc0112c8b <sys_signal+3>:    sub     $0x28,%esp
5360 0xc0112c8e <sys_signal+6>:    mov     0xc(%ebp),%eax
5361 /usr/src/linux-2.2.5/kernel/signal.c: 1078
5362 0xc0112c91 <sys_signal+9>:    mov     %eax,0xffffffffec(%ebp)
5363 /usr/src/linux-2.2.5/kernel/signal.c: 1079
5364 0xc0112c94 <sys_signal+12>:   movl    $0xc0000000,0xfffffffff0(%ebp)
5365 /usr/src/linux-2.2.5/kernel/signal.c: 1081
5366 0xc0112c9b <sys_signal+19>:   lea     0xffffffffd8(%ebp),%eax
5367 0xc0112c9e <sys_signal+22>:   push    %eax
5368 0xc0112c9f <sys_signal+23>:   lea     0xffffffffec(%ebp),%eax
5369 0xc0112ca2 <sys_signal+26>:   push    %eax
5370 0xc0112ca3 <sys_signal+27>:   pushl   0x8(%ebp)
5371 0xc0112ca6 <sys_signal+30>:   call    0xc01124b8 <do_sigaction>
5372 /usr/src/linux-2.2.5/kernel/signal.c: 1083
5373 0xc0112cab <sys_signal+35>:   test    %eax,%eax
5374 0xc0112cad <sys_signal+37>:   jne     0xc0112cb2 <sys_signal+42>
5375 0xc0112caf <sys_signal+39>:   mov     0xffffffffd8(%ebp),%eax
5376 /usr/src/linux-2.2.5/kernel/signal.c: 1084
5377 0xc0112cb2 <sys_signal+42>:   leave
5378 0xc0112cb3 <sys_signal+43>:   ret
5379

```

5380 Given a return address expression of "do_no_page+65", find out the
5381 function that do_no_page() calls by using the reverse flag:

```

5382
5383 crash> dis -r (do_no_page+65)
5384 0xc011ea68 <do_no_page>:      push    %ebp
5385 0xc011ea69 <do_no_page+1>:    mov     %esp,%ebp
5386 0xc011ea6b <do_no_page+3>:    push    %edi
5387 0xc011ea6c <do_no_page+4>:    push    %esi
5388 0xc011ea6d <do_no_page+5>:    push    %ebx
5389 0xc011ea6e <do_no_page+6>:    mov     0xc(%ebp),%ebx
5390 0xc011ea71 <do_no_page+9>:    mov     0x10(%ebp),%edx
5391 0xc011ea74 <do_no_page+12>:   mov     0x14(%ebp),%edi
5392 0xc011ea77 <do_no_page+15>:   mov     0x28(%ebx),%eax
5393 0xc011ea7a <do_no_page+18>:   test    %eax,%eax
5394 0xc011ea7c <do_no_page+20>:   je      0xc011ea85 <do_no_page+29>
5395 0xc011ea7e <do_no_page+22>:   mov     0x18(%eax),%ecx
5396 0xc011ea81 <do_no_page+25>:   test    %ecx,%ecx

```

```

5397 0xc011ea83 <do_no_page+27>: jne 0xc011eab0 <do_no_page+72>
5398 0xc011ea85 <do_no_page+29>: mov $0xffffe000,%eax
5399 0xc011ea8a <do_no_page+34>: and %esp,%eax
5400 0xc011ea8c <do_no_page+36>: decl 0x30(%eax)
5401 0xc011ea8f <do_no_page+39>: jns 0xc011ea9a <do_no_page+50>
5402 0xc011ea91 <do_no_page+41>: lock btrl $0x0,0xc022fb60
5403 0xc011ea9a <do_no_page+50>: push %edi
5404 0xc011ea9b <do_no_page+51>: mov 0x18(%ebp),%esi
5405 0xc011ea9e <do_no_page+54>: push %esi
5406 0xc011ea9f <do_no_page+55>: push %ebx
5407 0xc011eaa0 <do_no_page+56>: mov 0x8(%ebp),%esi
5408 0xc011eaa3 <do_no_page+59>: push %esi
5409 0xc011eaa4 <do_no_page+60>: call 0xc011e9e4 <do_anonymous_page>
5410 0xc011eaa9 <do_no_page+65>: jmp 0xc011eb47 <do_no_page+223>
5411

```

Disassemble 10 instructions starting at user virtual address 0x81ec624:

```

5412
5413
5414 crash> dis -u 81ec624 10
5415 0x81ec624: push %ebp
5416 0x81ec625: mov %esp,%ebp
5417 0x81ec627: sub $0x18,%esp
5418 0x81ec62a: movl $0x1,0x8(%ebp)
5419 0x81ec631: mov 0x82f9040,%eax
5420 0x81ec636: mov 0x10(%eax),%edx
5421 0x81ec639: and $0x100,%edx
5422 0x81ec63f: mov 0x14(%eax),%ecx
5423 0x81ec642: and $0x0,%ecx
5424 0x81ec645: mov %ecx,%eax
5425

```

Override the current decimal output radix format:

```

5426
5427
5428 crash> dis sys_read 10 -x
5429 0xffffffff8001178f <sys_read>: push %r13
5430 0xffffffff80011791 <sys_read+0x2>: mov %rsi,%r13
5431 0xffffffff80011794 <sys_read+0x5>: push %r12
5432 0xffffffff80011796 <sys_read+0x7>: mov $0xfffffffffffffffff7,%r12
5433 0xffffffff8001179d <sys_read+0xe>: push %rbp
5434 0xffffffff8001179e <sys_read+0xf>: mov %rdx,%rbp
5435 0xffffffff800117a1 <sys_read+0x12>: push %rbx
5436 0xffffffff800117a2 <sys_read+0x13>: sub $0x18,%rsp
5437 0xffffffff800117a6 <sys_read+0x17>: lea 0x14(%rsp),%rsi
5438 0xffffffff800117ab <sys_read+0x1c>: callq 0xffffffff8000b5b4 <fget_light>
5439

```

Disassemble from vfs_read+320 until the end of the function:

```

5440
5441
5442 crash> dis -f vfs_read+320
5443 0xffffffff8119d4e0 <vfs_read+320>: cmpq $0x0,0x20(%rax)
5444 0xffffffff8119d4e5 <vfs_read+325>: jne 0xffffffff8119d3e8 <vfs_read+72>
5445 0xffffffff8119d4eb <vfs_read+331>: mov $0xfffffffffffffffffea,%r12
5446 0xffffffff8119d4f2 <vfs_read+338>: jmp 0xffffffff8119d4c3 <vfs_read+291>
5447 0xffffffff8119d4f4 <vfs_read+340>: nopl 0x0(%rax)
5448 0xffffffff8119d4f8 <vfs_read+344>: callq 0xffffffff8119cc40 <do_sync_read>
5449 0xffffffff8119d4fd <vfs_read+349>: mov %rax,%r12
5450 0xffffffff8119d500 <vfs_read+352>: jmpq 0xffffffff8119d44c <vfs_read+172>
5451 0xffffffff8119d505 <vfs_read+357>: nopl (%rax)
5452 0xffffffff8119d508 <vfs_read+360>: mov $0xfffffffffffffffff7,%r12
5453 0xffffffff8119d50f <vfs_read+367>: jmp 0xffffffff8119d4c3 <vfs_read+291>
5454 0xffffffff8119d511 <vfs_read+369>: mov $0xfffffffffffffffff2,%r12
5455 0xffffffff8119d518 <vfs_read+376>: jmp 0xffffffff8119d4c3 <vfs_read+291>
5456 0xffffffff8119d51a <vfs_read+378>: nopw 0x0(%rax,%rax,1)
5457

```

Display the source code listing of the mmput() function:

```

5458
5459
5460 crash> dis -s mmput
5461 FILE: kernel/fork.c
5462 LINE: 617
5463
5464 612
5465 613 /*
5466 614 * Decrement the use count and release all resources for an mm.
5467 615 */
5468 616 void mmput(struct mm_struct *mm)
5469 * 617 {

```

```

5470         618             might_sleep();
5471         619
5472         620             if (atomic_dec_and_test(&mm->mm_users)) {
5473         621                 uprobe_clear_state(mm);
5474         622                 exit_aio(mm);
5475         623                 ksm_exit(mm);
5476         624                 khugepaged_exit(mm); /* must run before exit_mmap */
5477         625                 exit_mmap(mm);
5478         626                 set_mm_exe_file(mm, NULL);
5479         627                 if (!list_empty(&mm->mmlist)) {
5480         628                     spin_lock(&mmlist_lock);
5481         629                     list_del(&mm->mmlist);
5482         630                     spin_unlock(&mmlist_lock);
5483         631                 }
5484         632                 if (mm->binfmt)
5485         633                     module_put(mm->binfmt->module);
5486         634                 mmdrop(mm);
5487         635             }
5488         636     }
5489

```

The disassembly of `dentry_kill()` shows an indirect call to a function whose address is contained within a register. Display the source code associated with the indirect function call:

```

5490
5491
5492
5493
5494     crash> dis dentry_kill
5495     ...
5496     0xfffffffff811dcfb4 <dentry_kill+324>:    callq  %rax
5497     ...
5498     crash> dis -s 0xfffffffff811dcfb4
5499     FILE: fs/dcache.c
5500     LINE: 276
5501
5502         271             spin_unlock(&dentry->d_lock);
5503         272             spin_unlock(&inode->i_lock);
5504         273             if (!inode->i_nlink)
5505         274                 fsnotify_inoderemove(inode);
5506         275             if (dentry->d_op && dentry->d_op->d_iput)
5507     * 276                 dentry->d_op->d_iput(dentry, inode);
5508         277             else
5509         278                 iput(inode);
5510         279             } else {
5511         280                 spin_unlock(&dentry->d_lock);
5512         281             }
5513         282     }
5514
5515
5516
5517     NAME
5518     kmem - kernel memory
5519
5520     SYNOPSIS
5521     kmem [-f|-F|-c|-C|-i|-v|-V|-n|-z|-o|-h] [-p | -m member[,member]]
5522           [[-s|-S|-S=cpu[s]|-r] [slab] [-I slab[,slab]]] [-g [flags]] [[-P] address]]
5523
5524     DESCRIPTION
5525     This command displays information about the use of kernel memory.
5526
5527     -f displays the contents of the system free memory headers.
5528         also verifies that the page count equals nr_free_pages.
5529     -F same as -f, but also dumps all pages linked to that header.
5530     -c walks through the page_hash_table and verifies page_cache_size.
5531     -C same as -c, but also dumps all pages in the page_hash_table.
5532     -i displays general memory usage information
5533     -v displays the mapped virtual memory regions allocated by vmalloc().
5534     -V displays the kernel vm_stat table if it exists, or in more recent
5535         kernels, the vm_zone_stat, vm_node_stat and vm_numa_stat tables,
5536         the cumulative page_states counter values if they exist, and/or
5537         the cumulative, vm event_states counter values if they exist.
5538     -n display memory node, memory section, memory block data and state;
5539         the state of each memory section is shown as the following flags:
5540         "P": SECTION_MARKED_PRESENT
5541         "M": SECTION_HAS_MEM_MAP
5542         "O": SECTION_IS_ONLINE

```

```

5543         "E": SECTION_IS_EARLY
5544         "D": SECTION_TAINT_ZONE_DEVICE
5545     -z displays per-zone memory statistics.
5546     -o displays each cpu's offset value that is added to per-cpu symbol
5547         values to translate them into kernel virtual addresses.
5548     -h display the address of hugepage hstate array entries, along with
5549         their hugepage size, total and free counts, and name.
5550     -p displays basic information about each page structure in the system
5551         mem_map[] array, made up of the page struct address, its associated
5552         physical address, the page.mapping, page.index, page._count and
5553         page.flags fields.
5554     -m member similar to -p, but displays page structure contents specified by
5555         a comma-separated list of one or more struct page members. The
5556         "flags" member will always be expressed in hexadecimal format, and
5557         the "_count" and "_mapcount" members will always be expressed
5558         in decimal format. Otherwise, all other members will be displayed
5559         in hexadecimal format unless the output radix is 10 and the member
5560         is a signed/unsigned integer. Members that are data structures may
5561         be specified either by the data structure's member name, or expanded
5562         to specify a member of the data structure. For example, "-m lru"
5563         refers to a list_head data structure, and both the list_head.next
5564         and list_head.prev pointer values will be displayed, whereas if
5565         "-m lru.next" is specified, just the list_head.next value will
5566         be displayed.
5567     -s displays basic kmalloc() slab data.
5568     -S displays all kmalloc() slab data, including all slab objects,
5569         and whether each object is in use or is free. If CONFIG_SLUB,
5570         slab data for each per-cpu slab is displayed, along with the
5571         address of each kmem_cache_node, its count of full and partial
5572         slabs, and a list of all tracked slabs.
5573         Note: one can specify the per-cpu slab data to be displayed;
5574         the cpu[s] can be given as "1,3,5", "1-3", "1,3,5-7,10",
5575         "all", or "a" (shortcut for "all").
5576     -r displays the accumulated basic kmalloc() slab data of each
5577         root slab cache and its children. The kernel must contain the
5578         "slab_root_caches" list_head. (currently only available if
5579         CONFIG_SLUB)
5580     slab when used with -s, -S or -r, limits the command to only the slab
5581         cache of name "slab". If the slab argument is "list", then
5582         all slab cache names and addresses are listed.
5583     -I slab when used with -s, -S or -r, one or more slab cache names in a
5584         comma-separated list may be specified as slab caches to ignore.
5585     -g displays the enumerator value of all bits in the page structure's
5586         "flags" field.
5587     flags when used with -g, translates all bits in this hexadecimal page
5588         structure flags value into its enumerator values.
5589     -P declares that the following address argument is a physical address.
5590     address when used without any flag, the address can be a kernel virtual,
5591         or physical address; a search is made through the symbol table,
5592         the kmalloc() slab subsystem, the free list, the page_hash_table,
5593         the vmalloc() region subsystem, the current set of task_structs
5594         and kernel stacks, and the mem_map array. If found in any of
5595         those areas, the information will be dumped in the same manner as
5596         if the location-specific flags were used; if contained within a
5597         current task_struct or kernel stack, that task's context will be
5598         displayed.
5599     address when used with -s or -S, searches the kmalloc() slab subsystem
5600         for the slab containing of this virtual address, showing whether
5601         it is in use or free.
5602     address when used with -f, the address can be either a page pointer,
5603         a physical address, or a kernel virtual address; the free_area
5604         header containing the page (if any) is displayed.
5605     address when used with -p, the address can be either a page pointer, a
5606         physical address, or a kernel virtual address; its basic mem_map
5607         page information is displayed.
5608     address when used with -m, the address can be either a page pointer, a
5609         physical address, or a kernel virtual address; the specified
5610         members of the associated page struct are displayed.
5611     address when used with -c, the address must be a page pointer address;
5612         the page_hash_table entry containing the page is displayed.
5613     address when used with -l, the address must be a page pointer address;
5614         the page address is displayed if it is contained with the list.
5615     address when used with -v, the address can be a mapped kernel virtual

```

5616 address or physical address; the mapped region containing the
5617 address is displayed.
5618

5619 All address arguments above must be expressed in hexadecimal format.
5620

5621 EXAMPLES

5622 Display memory usage information:
5623

5624 crash> kmem -i

```
5625          PAGES          TOTAL          PERCENTAGE
5626 TOTAL MEM 1974231          7.5 GB          ----
5627     FREE  208962      816.3 MB      10% of TOTAL MEM
5628     USED  1765269        6.7 GB      89% of TOTAL MEM
5629   SHARED  365066        1.4 GB      18% of TOTAL MEM
5630   BUFFERS  111376      435.1 MB       5% of TOTAL MEM
5631   CACHED  1276196        4.9 GB      64% of TOTAL MEM
5632     SLAB  120410      470.4 MB       6% of TOTAL MEM
5633
5634 TOTAL HUGE  524288          2 GB          ----
5635   HUGE FREE  524288          2 GB     100% of TOTAL HUGE
5636
5637 TOTAL SWAP  2498559          9.5 GB          ----
5638   SWAP USED   81978      320.2 MB       3% of TOTAL SWAP
5639   SWAP FREE  2416581          9.2 GB      96% of TOTAL SWAP
5640
5641 COMMIT LIMIT 3485674      13.3 GB          ----
5642   COMMITTED  850651        3.2 GB      24% of TOTAL LIMIT
5643
```

5644 Display and verify free memory data:
5645

5646 crash> kmem -f

```
5647 NODE
5648 0
5649 ZONE NAME          SIZE    FREE  MEM_MAP  START_PADDR  START_MAPNR
5650 0 DMA             4096    3372  c4000040      0          0
5651 AREA SIZE  FREE AREA_STRUCT  BLOCKS  PAGES
5652 0 4k      c02eb004          2      2
5653 1 8k      c02eb010          3      6
5654 2 16k     c02eb01c          5     20
5655 3 32k     c02eb028          4     32
5656 4 64k     c02eb034          5     80
5657 5 128k    c02eb040          3     96
5658 6 256k    c02eb04c          3    192
5659 7 512k    c02eb058          1    128
5660 8 1024k   c02eb064          1    256
5661 9 2048k   c02eb070          5   2560
5662
5663 ZONE NAME          SIZE    FREE  MEM_MAP  START_PADDR  START_MAPNR
5664 1 Normal     225280   202269  c4044040    1000000      4096
5665 AREA SIZE  FREE AREA_STRUCT  BLOCKS  PAGES
5666 0 4k      c02eb0b8          1      1
5667 1 8k      c02eb0c4          2      4
5668 2 16k     c02eb0d0          0      0
5669 3 32k     c02eb0dc          1      8
5670 4 64k     c02eb0e8          1     16
5671 5 128k    c02eb0f4          0      0
5672 6 256k    c02eb100          0      0
5673 7 512k    c02eb10c          0      0
5674 8 1024k   c02eb118          0      0
5675 9 2048k   c02eb124        395 202240
5676
5677 ZONE NAME          SIZE    FREE  MEM_MAP  START_PADDR  START_MAPNR
5678 2 HighMem    819200   748686  c4ee0040    3800000     229376
5679 AREA SIZE  FREE AREA_STRUCT  BLOCKS  PAGES
5680 0 4k      c02eb16c          10     10
5681 1 8k      c02eb178          2      4
5682 2 16k     c02eb184          0      0
5683 3 32k     c02eb190          2     16
5684 4 64k     c02eb19c          1     16
5685 5 128k    c02eb1a8          1     32
5686 6 256k    c02eb1b4          1     64
5687 7 512k    c02eb1c0          0      0
5688 8 1024k   c02eb1cc          0      0
```

5689 9 2048k c02eb1d8 1462 748544

5690

5691 nr_free_pages: 954327 (verified)

5692

5693 Dump all the base addresses of each free memory area from above:

5694

5695 crash> kmem -F

5696 NODE

5697 0

ZONE	NAME	SIZE	FREE	MEM_MAP	START_PADDR	START_MAPNR
0	DMA	4096	3372	c4000040	0	0

AREA	SIZE	FREE	AREA_STRUCT
0	4k	c02eb004	

5701 c400ded8

5702 c4042528

AREA	SIZE	FREE	AREA_STRUCT
1	8k	c02eb010	

5704 c400de50

5705 c400cee8

5706 c40424a0

AREA	SIZE	FREE	AREA_STRUCT
2	16k	c02eb01c	

5709 c400dd40

5710 c400cf70

5711 c40425b0

5712 c400f7d0

5713 c40028a0

AREA	SIZE	FREE	AREA_STRUCT
3	32k	c02eb028	

5716 c4042280

5717 c400f8e0

5718 c4002680

5719 c4000260

AREA	SIZE	FREE	AREA_STRUCT
4	64k	c02eb034	

5722 c400d080

5723 c4041e40

5724 ...

5725

5726 Dump the mem_map[] array:

5727

5728 crash> kmem -p

PAGE	PHYSICAL	MAPPING	INDEX	CNT	FLAGS
f5c51200	10000	0	0	1	80 slab
f5c51220	11000	0	0	1	80 slab
f5c51240	12000	0	0	1	80 slab
f5c51260	13000	0	0	1	80 slab
f5c51280	14000	0	0	1	80 slab
f5c512a0	15000	0	0	1	80 slab
f5c512c0	16000	0	0	1	80 slab
f5c512e0	17000	0	0	1	80 slab
f5c51300	18000	0	0	1	80 slab
f5c51320	19000	0	0	1	80 slab
f5c51340	1a000	0	0	1	80 slab
f5c51360	1b000	0	0	1	80 slab
f5c51380	1c000	e6c6a754	13b67	2	868 uptodate,lru,active,private
f5c513a0	1d000	0	0	1	80 slab
f5c513c0	1e000	0	0	1	80 slab
f5c513e0	1f000	0	0	1	80 slab
f5c51400	20000	e6c6a754	13bbb	2	868 uptodate,lru,active,private
f5c51420	21000	0	0	1	80 slab
f5c51440	22000	0	0	1	80 slab

5751 ...

5752

5753 Display the "page.lru" list_head structure member in each page:

5754

5755 crash> kmem -m lru

PAGE	lru
ffffea0000000000	0000000000000000,0000000000000000
ffffea0000000040	ffffea0000000060,ffffea0000000060
ffffea0000000080	ffffea00000000a0,ffffea00000000a0
ffffea00000000c0	ffffea00000000e0,ffffea00000000e0
ffffea0000000100	ffffea0000000120,ffffea0000000120

```

5762     fffffea00000000140 fffffea00000000160,fffffea00000000160
5763     fffffea00000000180 fffffea000000001a0,fffffea000000001a0
5764     fffffea000000001c0 fffffea000000001e0,fffffea000000001e0
5765     fffffea00000000200 fffffea00000000220,fffffea00000000220
5766     fffffea00000000240 fffffea00000000260,fffffea00000000260
5767     fffffea00000000280 fffffea000000002a0,fffffea000000002a0
5768     fffffea000000002c0 fffffea000000002e0,fffffea000000002e0
5769     fffffea00000000300 fffffea00000000320,fffffea00000000320
5770     fffffea00000000340 fffffea00000000360,fffffea00000000360
5771     fffffea00000000380 fffffea000000003a0,fffffea000000003a0
5772     fffffea000000003c0 fffffea000000003e0,fffffea000000003e0
5773     fffffea00000000400 fffff88021e5e41e8,fffffea000000002020
5774     fffffea00000000440 dead0000000100100,dead0000000200200
5775     fffffea00000000480 dead0000000100100,dead0000000200200
5776     fffffea000000004c0 dead0000000100100,dead0000000200200
5777     ...
5778

```

Find the two pages that link to the page at fffffea0001dafb20 via their page.lru list_head's next and prev pointers:

```

5781
5782     crash> kmem -m lru | grep fffffea0001dafb20
5783     fffffea0000006b500 fffffea0001dafb20,fffffea0001eb4520
5784     fffffea00000127d80 fffffea000152b620,fffffea0001dafb20
5785

```

Find all of the combined slab/page structures that are used by the kcalloc-8192 slab cache:

```

5789     crash> kmem -s kcalloc-8192
5790     CACHE          OBJSIZE  ALLOCATED      TOTAL   SLABS  SSIZE  NAME
5791     fffff880215802e00      8192           65         80    20    32k  kcalloc-8192
5792     crash> kmem -m slab_cache | grep fffff880215802e00
5793     fffffea0004117800 fffff880215802e00
5794     fffffea00041ca600 fffff880215802e00
5795     fffffea00044ab200 fffff880215802e00
5796     fffffea0004524000 fffff880215802e00
5797     fffffea0004591600 fffff880215802e00
5798     fffffea00047eac00 fffff880215802e00
5799     fffffea0004875800 fffff880215802e00
5800     fffffea0008357a00 fffff880215802e00
5801     fffffea0008362a00 fffff880215802e00
5802     fffffea00083b9400 fffff880215802e00
5803     fffffea00083c1000 fffff880215802e00
5804     fffffea00083c1e00 fffff880215802e00
5805     fffffea00083c2000 fffff880215802e00
5806     fffffea00083c2a00 fffff880215802e00
5807     fffffea00083d2000 fffff880215802e00
5808     fffffea00083d3e00 fffff880215802e00
5809     fffffea0008407c00 fffff880215802e00
5810     fffffea000848ce00 fffff880215802e00
5811     fffffea0008491800 fffff880215802e00
5812     fffffea00084bf800 fffff880215802e00
5813

```

Use the commands above with a page pointer or a physical address argument:

```

5814
5815
5816     crash> kmem -f c40425b0
5817     NODE
5818     0
5819     ZONE  NAME          SIZE    FREE  MEM_MAP  START_PADDR  START_MAPNR
5820     0    DMA            4096    3372  c4000040      0             0
5821     AREA  SIZE  FREE AREA_STRUCT
5822     2    16k      c02eb01c
5823     c40425b0 (c40425b0 is 1st of 4 pages)
5824
5825     crash> kmem -p c25a9c00
5826     PAGE    PHYSICAL    MAPPING    INDEX CNT  FLAGS
5827     c25a9c00    1fe0000  f429d2e4   21fe3eb  2 800828 uptodate,lru,private
5828
5829     crash> kmem -p 1fe0000
5830     PAGE    PHYSICAL    MAPPING    INDEX CNT  FLAGS
5831     c25a9c00    1fe0000  f429d2e4   21fe3eb  2 800828 uptodate,lru,private
5832

```

Display the mapped memory regions allocated by vmalloc():

```

5834

```



```

5835 crash> kmem -v
5836 VMAP_AREA VM_STRUCT ADDRESS RANGE SIZE
5837 f7048e00 f7048e40 f7dfe000 - f7e00000 8192
5838 f7048ec0 f7048f00 f7e00000 - f7e05000 20480
5839 f7151fc0 f7159540 f7e06000 - f7e08000 8192
5840 f704da80 f704dac0 f7e0a000 - f7e0c000 8192
5841 f704d980 f704d9c0 f7e0e000 - f7e10000 8192
5842 f724f1c0 f724f200 f7e12000 - f7e14000 8192
5843 f704d840 f704d880 f7e14000 - f7e17000 12288
5844 f704d400 f704d440 f7e18000 - f7e1d000 20480
5845 f73f5840 f73f5880 f7e1e000 - f7e2a000 49152
5846 f6334480 f63344c0 f7e2c000 - f7e2e000 8192
5847 f635d600 f635d640 f7e4a000 - f7e5b000 69632
5848 f41b4700 f5771a40 f7e6e000 - f7e70000 8192
5849 f622f6c0 f622f700 f7e71000 - f7e79000 32768
5850 f63a9f00 f63a9f40 f7e84000 - f7e87000 12288
5851 f63a9d00 f63a9d40 f7e8f000 - f7e91000 8192
5852 f5546480 f39db800 f7eb8000 - f7ec2000 40960
5853 f5ce9640 f5777e80 f7ec6000 - f7ed1000 45056
5854 f63a9b00 f63a9b40 f7ed1000 - f7efd000 180224
5855 f63a9800 f63a9840 f7f1d000 - f7f26000 36864
5856 f63a9640 f63a9880 f7f43000 - f7f52000 61440
5857 f5771f00 f4183840 f7f53000 - f7f64000 69632
5858 f5ce9a00 f30c4a00 f7fcf000 - f801e000 323584
5859 f63a93c0 f63a9400 f805d000 - f8132000 872448
5860 f63a91c0 f63a95c0 f814b000 - f8150000 20480
5861 f63a9140 f63a9180 f8151000 - f8352000 2101248
5862 f624eb00 f624eb40 f8353000 - f8355000 8192
5863 f563eb40 f563eb80 f8356000 - f835e000 32768
5864 f63d5ec0 f63d5f00 f8360000 - f8371000 69632
5865 f63d5cc0 f6287b80 f83c2000 - f84c3000 1052672
5866 ...
5867

```

5868 Dump the virtual memory statistics:

```

5869 crash> kmem -V
5870 VM_ZONE_STAT:
5871 NR_FREE_PAGES: 30085
5872 NR_ZONE_INACTIVE_ANON: 1985
5873 NR_ZONE_ACTIVE_ANON: 338275
5874 NR_ZONE_INACTIVE_FILE: 19760
5875 NR_ZONE_ACTIVE_FILE: 12018
5876 NR_ZONE_UNEVICTABLE: 0
5877 NR_ZONE_WRITE_PENDING: 4
5878 NR_MLOCK: 0
5879 NR_PAGETABLE: 1562
5880 NR_KERNEL_STACK_KB: 1728
5881 NR_BOUNCE: 0
5882 NR_FREE_CMA_PAGES: 0
5883
5884 VM_NODE_STAT:
5885 NR_INACTIVE_ANON: 1985
5886 NR_ACTIVE_ANON: 338275
5887 NR_INACTIVE_FILE: 19760
5888 NR_ACTIVE_FILE: 12018
5889 NR_UNEVICTABLE: 0
5890 NR_SLAB_RECLAIMABLE: 3111
5891 NR_SLAB_UNRECLAIMABLE: 3039
5892 NR_ISOLATED_ANON: 0
5893 NR_ISOLATED_FILE: 0
5894 WORKINGSET_REFAULT: 0
5895 WORKINGSET_ACTIVATE: 0
5896 WORKINGSET_NODERECLAIM: 0
5897 NR_ANON_MAPPED: 338089
5898 NR_FILE_MAPPED: 8102
5899 NR_FILE_PAGES: 33949
5900 NR_FILE_DIRTY: 4
5901 NR_WRITEBACK: 0
5902 NR_WRITEBACK_TEMP: 0
5903 NR_SHMEM: 2171
5904 NR_SHMEM_THPS: 0
5905 NR_SHMEM_PMDMAPPED: 0
5906 NR_ANON_THPS: 86
5907

```

```

5908         NR_UNSTABLE_NFS: 0
5909         NR_VMSCAN_WRITE: 0
5910     NR_VMSCAN_IMMEDIATE: 0
5911         NR_DIRTIED: 155
5912         NR_WRITTEN: 75
5913
5914     VM_NUMA_STAT:
5915         NUMA_HIT: 575409
5916         NUMA_MISS: 0
5917         NUMA_FOREIGN: 0
5918         NUMA_INTERLEAVE_HIT: 12930
5919         NUMA_LOCAL: 575409
5920         NUMA_OTHER: 0
5921
5922     VM_EVENT_STATES:
5923         PGPGIN: 282492
5924         PGPGOUT: 6773
5925         PSWPIN: 0
5926         PSWPOUT: 0
5927         PGALLOC_DMA: 0
5928         PGALLOC_DMA32: 693092
5929         PGALLOC_NORMAL: 0
5930     ...
5931
5932 Display hugepage hstate information:
5933
5934     crash> kmem -h
5935         HSTATE      SIZE      FREE      TOTAL  NAME
5936     ffffffff81f7a800      2MB          10          64  hugepages-2048kB
5937
5938 Determine (and verify) the page cache size:
5939
5940     crash> kmem -c
5941     page_cache_size: 18431 (verified)
5942
5943 Dump all pages in the page_hash_table:
5944
5945     crash> kmem -C
5946     page_hash_table[0]
5947     c0325b40
5948     c03a0598
5949     c03b4070
5950     c0364c28
5951     c0357690
5952     c02ef338
5953     c02d7c60
5954     c02c11e0
5955     c02a3d70
5956     page_hash_table[1]
5957     c0394ce8
5958     c03c4218
5959     c03b4048
5960     c0364c00
5961     c0357668
5962     c02d6e50
5963     c02d7dc8
5964     c02c0cb8
5965     c02db630
5966     c02ebad0
5967     page_hash_table[2]
5968     c037e808
5969     c034e248
5970     c03b4020
5971     c02ec868
5972     c03baa60
5973     ...
5974     page_hash_table[2047]
5975     c033a798
5976     c0390b48
5977     c03b4098
5978     c0364890
5979     c03576b8
5980     c02d2c38

```

5981 c02d7c88
5982 c02de5d8
5983
5984 page_cache_size: 18437 (verified)
5985

5986 Find the page_hash_table entry containing page c03576b8:
5987

5988 crash> kmem -c c03576b8
5989 page_hash_table[2047]
5990 c03576b8
5991

5992 Display kmallocc() slab data:
5993

5994 crash> kmem -s
5995 CACHE OBJSIZE ALLOCATED TOTAL SLABS SSIZE NAME
5996 c02eadc0 232 58 68 4 4k kmem_cache
5997 f79c2888 128 0 0 0 4k ip_vs_conn
5998 f79c2970 96 0 0 0 4k tcp_tw_bucket
5999 f79c2a58 32 12 565 5 4k tcp_bind_bucket
6000 f79c2b40 64 0 59 1 4k tcp_open_request
6001 f79c2c28 64 1 59 1 4k inet_peer_cache
6002 f79c2d10 32 11 339 3 4k ip_fib_hash
6003 f79c2df8 160 8 120 5 4k ip_dst_cache
6004 f79c2ee0 128 1 30 1 4k arp_cache
6005 c8402970 96 30208 37800 945 4k blkdev_requests
6006 c8402a58 384 0 0 0 4k nfs_read_data
6007 c8402b40 384 0 0 0 4k nfs_write_data
6008 c8402c28 96 0 0 0 4k nfs_page
6009 c8402d10 20 0 0 0 4k dnotify cache
6010 c8402df8 92 3 336 8 4k file lock cache
6011 c8402ee0 16 0 0 0 4k fasync cache
6012 c84027a0 32 3 339 3 4k uid_cache
6013 c84026b8 160 320 624 26 4k skbuff_head_cache
6014 c84025d0 832 32 180 20 8k sock
6015 c84024e8 132 0 203 7 4k sigqueue
6016 c8402400 64 19 472 8 4k cdev_cache
6017 c8402318 64 8 236 4 4k bdev_cache
6018 c8402230 96 11 120 3 4k mnt_cache
6019 c8402148 480 817 848 106 4k inode_cache
6020 c8402060 128 1352 1470 49 4k dentry_cache
6021 c8403ee0 96 244 440 11 4k filp
6022 c8403df8 4096 0 12 12 4k names_cache
6023 c8403d10 96 14936 16000 400 4k buffer_head
6024 c8403c28 128 25 240 8 4k mm_struct
6025 c8403b40 64 393 1298 22 4k vm_area_struct
6026 c8403a58 64 30 472 8 4k fs_cache
6027 c8403970 416 30 135 15 4k files_cache
6028 c8403888 1312 32 99 33 4k signal_act
6029 c84037a0 131072 0 0 0 128k size-131072 (DMA)
6030 c84036b8 131072 1 1 1 128k size-131072
6031 c84035d0 65536 0 0 0 64k size-65536 (DMA)
6032 c84034e8 65536 0 0 0 64k size-65536
6033 c8403400 32768 0 0 0 32k size-32768 (DMA)
6034 c8403318 32768 0 1 1 32k size-32768
6035 c8403230 16384 0 0 0 16k size-16384 (DMA)
6036 c8403148 16384 0 0 0 16k size-16384
6037 c8403060 8192 0 0 0 8k size-8192 (DMA)
6038 c8401ee0 8192 1 2 2 8k size-8192
6039 c8401df8 4096 0 0 0 4k size-4096 (DMA)
6040 c8401d10 4096 30 30 30 4k size-4096
6041 c8401c28 2048 0 0 0 4k size-2048 (DMA)
6042 c8401b40 2048 37 132 66 4k size-2048
6043 c8401a58 1024 0 0 0 4k size-1024 (DMA)
6044 c8401970 1024 301 328 82 4k size-1024
6045 c8401888 512 0 0 0 4k size-512 (DMA)
6046 c84017a0 512 141 168 21 4k size-512
6047 c84016b8 256 0 0 0 4k size-256 (DMA)
6048 c84015d0 256 80 435 29 4k size-256
6049 c84014e8 128 0 0 0 4k size-128 (DMA)
6050 c8401400 128 508 840 28 4k size-128
6051 c8401318 64 0 0 0 4k size-64 (DMA)
6052 c8401230 64 978 1357 23 4k size-64
6053 c8401148 32 0 0 0 4k size-32 (DMA)

6054 c8401060 32 1244 1808 16 4k size-32

6055

6056 Display all slab data in the "arp_cache" cache:

6057

6058 crash> kmem -S arp_cache

CACHE	OBJSIZE	ALLOCATED	TOTAL	SLABS	SSIZE	NAME
f79c2ee0	128	1	30	1	4k	arp_cache

SLAB	MEMORY	TOTAL	ALLOCATED	FREE
f729d000	f729d0a0	30	1	29

6062 FREE / [ALLOCATED]

6063 f729d0a0 (cpu 7 cache)

6064 f729d120 (cpu 7 cache)

6065 f729d1a0 (cpu 7 cache)

6066 f729d220 (cpu 7 cache)

6067 f729d2a0 (cpu 7 cache)

6068 f729d320 (cpu 7 cache)

6069 f729d3a0 (cpu 7 cache)

6070 f729d420 (cpu 7 cache)

6071 f729d4a0 (cpu 7 cache)

6072 f729d520 (cpu 7 cache)

6073 f729d5a0 (cpu 7 cache)

6074 f729d620 (cpu 7 cache)

6075 f729d6a0 (cpu 7 cache)

6076 f729d720 (cpu 7 cache)

6077 f729d7a0 (cpu 7 cache)

6078 f729d820 (cpu 7 cache)

6079 f729d8a0 (cpu 7 cache)

6080 f729d920 (cpu 7 cache)

6081 f729d9a0 (cpu 7 cache)

6082 f729da20 (cpu 7 cache)

6083 f729daa0 (cpu 7 cache)

6084 f729db20 (cpu 7 cache)

6085 f729dba0 (cpu 7 cache)

6086 f729dc20 (cpu 7 cache)

6087 f729dca0 (cpu 7 cache)

6088 f729dd20 (cpu 7 cache)

6089 f729dda0 (cpu 7 cache)

6090 f729de20 (cpu 7 cache)

6091 f729dea0 (cpu 3 cache)

6092 [f729df20]

6093

6094 Search the kmalloc() slab subsystem for address c3fbdb60:

6095

6096 crash> kmem -s c3fbdb60

CACHE	OBJSIZE	ALLOCATED	TOTAL	SLABS	SSIZE	NAME
c8402970	96	30208	37800	945	4k	blkdev_requests

SLAB	MEMORY	TOTAL	ALLOCATED	FREE
c3fbd020	c3fbd0e0	40	40	0

6100 FREE / [ALLOCATED]

6101 [c3fbdb60]

6102

6103 Make a generic search (no flags) for the same address c3fbdb60:

6104

6105 crash> kmem c3fbdb60

CACHE	OBJSIZE	ALLOCATED	TOTAL	SLABS	SSIZE	NAME
c8402970	96	30208	37800	945	4k	blkdev_requests

SLAB	MEMORY	TOTAL	ALLOCATED	FREE
c3fbd020	c3fbd0e0	40	40	0

6110 FREE / [ALLOCATED]

6111 [c3fbdb60]

6112

PAGE	PHYSICAL	MAPPING	INDEX	CNT	FLAGS
c410ee74	3fbd000	0	0	1	slab

6117

6118 Display memory node data (if supported):

6119

6120 crash> kmem -n

NODE	SIZE	PGLIST_DATA	BOOTMEM_DATA	NODE_ZONES
0	262095	ffff88003d52a000	----	ffff88003d52a000

6121 ffff88003d52a740

6122 ffff88003d52ae80

6123 ffff88003d52b5c0

MEM_MAP	START_PADDR	START_MAPNR
---------	-------------	-------------

6126

6127612861296130613161326133613461356136

fffffea00000000040	1000	1			
ZONE	NAME	SIZE	MEM_MAP	START_PADDR	START_MAPNR
0	DMA	4095	fffffea00000000040	1000	1
1	DMA32	258000	fffffea00000040000	1000000	4096
2	Normal	0	0	0	0
3	Movable	0	0	0	0

6137613861396140614161426143614461456146

NR	SECTION	CODED_MEM_MAP	MEM_MAP	STATE	PFN
0	fffff88003d4d9000	fffffea00000000000	fffffea00000000000	PM	0
1	fffff88003d4d9020	fffffea00000000000	fffffea00000200000	PM	32768
2	fffff88003d4d9040	fffffea00000000000	fffffea00000400000	PM	65536
3	fffff88003d4d9060	fffffea00000000000	fffffea00000600000	PM	98304
4	fffff88003d4d9080	fffffea00000000000	fffffea00000800000	PM	131072
5	fffff88003d4d90a0	fffffea00000000000	fffffea00000a00000	PM	163840
6	fffff88003d4d90c0	fffffea00000000000	fffffea00000c00000	PM	196608
7	fffff88003d4d90e0	fffffea00000000000	fffffea00000e00000	PM	229376

6147614861496150615161526153615461556156

MEM_BLOCK	NAME	PHYSICAL	RANGE	STATE	START_SECTION_NO
fffff88003a707c00	memory0	0	- 7fffffff	ONLINE	0
fffff88003a6e0000	memory1	8000000	- ffffffff	ONLINE	1
fffff88003a6e1000	memory2	10000000	- 17fffffff	ONLINE	2
fffff88003a6e1400	memory3	18000000	- 1fffffff	ONLINE	3
fffff88003a6e1800	memory4	20000000	- 27fffffff	ONLINE	4
fffff88003a6e0400	memory5	28000000	- 2fffffff	ONLINE	5
fffff88003a6e0800	memory6	30000000	- 37fffffff	ONLINE	6
fffff88003a6e0c00	memory7	38000000	- 3fffffff	ONLINE	7

61576158

Translate a page structure's flags field contents:

6159616061616162616361646165

```
crash> kmem -g 4080
FLAGS: 4080
PAGE-FLAG      BIT  VALUE
PG_slab        7   0000080
PG_head        14   0004000
crash>
```

61666167616861696170

NAME

ptov - physical to virtual

per-cpu to virtual

617161726173

SYNOPSIS

ptov [address | offset:cpuspec]

6174617561766177617861796180618161826183618461856186

DESCRIPTION

This command translates a hexadecimal physical address into a kernel virtual address. Alternatively, a hexadecimal per-cpu offset and cpu specifier will be translated into kernel virtual addresses for each cpu specified.

address a physical address

offset:cpuspec a per-cpu offset with a CPU specifier:

- : CPU of the currently selected task.
- :a[ll] all CPUs.
- :[-#][,...] CPU list(s), e.g. "1,3,5", "1-3", or "1,3,5-7,10".

618761886189

EXAMPLES

Translate physical address 56e000 into a kernel virtual address:

6190619161926193

```
crash> ptov 56e000
VIRTUAL      PHYSICAL
fffff88000056e000 56e000
```

619461956196

Translate per-cpu offset b0c0 into a kernel virtual address for all cpus:

619761986199

```
crash> ptov b0c0:a
PER-CPU OFFSET: b0c0
CPU      VIRTUAL
```

```
6200         [0]  ffff88021e20b0c0
6201         [1]  ffff88021e24b0c0
6202         [2]  ffff88021e28b0c0
6203         [3]  ffff88021e2cb0c0
6204
6205
```

NAME

```
sys - system data
```

SYNOPSIS

```
sys [-c [name|number]] [-t] [-i] config
```

DESCRIPTION

This command displays system-specific data. If no arguments are entered, the same system data shown during crash invocation is shown.

```
-c [name|number]  If no name or number argument is entered, dump all
                  sys_call_table entries. If a name string is entered,
                  search the table for all entries containing the string.
                  If a number is entered, the table entry associated with
                  that number is displayed. If the current output radix
                  has been set to 16, the system call numbers will be
                  displayed in hexadecimal.
config           If the kernel was configured with CONFIG_IKCONFIG, then
                  dump the in-kernel configuration data.
-t              Display kernel taint information. If the "tainted_mask"
                  symbol exists, show its hexadecimal value and translate
                  each bit set to the symbolic letter of the taint type.
                  On older kernels with the "tainted" symbol, only its
                  hexadecimal value is shown. The relevant kernel sources
                  should be consulted for the meaning of the letter(s) or
                  hexadecimal bit value(s).
-panic          Panic a live system. Requires write permission to
                  /dev/mem. Results in the crash context causing an
                  "Attempted to kill the idle task!" panic. (The dump
                  will indicate that the crash context has a PID of 0).
-i             Dump the DMI string data if available in the kernel.
```

EXAMPLES

Display essential system information:

```
crash> sys
        KERNEL: vmlinux.4
        DUMPFILE: lcore.cr.4
        CPUS: 4
        DATE: Mon Oct 11 18:48:55 1999
        UPTIME: 10 days, 14:14:39
        LOAD AVERAGE: 0.74, 0.23, 0.08
        TASKS: 77
        NODENAME: test.mclinux.com
        RELEASE: 2.2.5-15smp
        VERSION: #24 SMP Mon Oct 11 17:41:40 CDT 1999
        MACHINE: i686 (500 MHz)
        MEMORY: 1 GB
```

Dump the system configuration data (if CONFIG_IKCONFIG):

```
crash> sys config
#
# Automatically generated make config: don't edit
# Linux kernel version: 2.6.16
# Mon Apr 10 07:58:06 2006
#
CONFIG_X86_64=y
CONFIG_64BIT=y
CONFIG_X86=y
CONFIG_SEMAPHORE_SLEEPERS=y
CONFIG_MMU=y
CONFIG_RWSEM_GENERIC_SPINLOCK=y
CONFIG_GENERIC_CALIBRATE_DELAY=y
CONFIG_X86_CMPXCHG=y
CONFIG_EARLY_PRINTK=y
CONFIG_GENERIC_ISA_DMA=y
```

```

6273 CONFIG_GENERIC_IOMAP=y
6274 CONFIG_ARCH_MAY_HAVE_PC_FDC=y
6275 CONFIG_DMI=y
6276 ...
6277
6278 Display the kernel taint information, in this case where both the
6279 TAINTE_WARN and TAINTE_PROPRIETARY_MODULE bits have been set:
6280
6281 crash> sys -t
6282 TAINTE_MASK: 201 PW
6283
6284 Dump the system call table:
6285
6286 crash> sys -c
6287 NUM SYSTEM CALL FILE AND LINE NUMBER
6288 0 sys_ni_syscall ../kernel/sys.c: 48
6289 1 sys_exit ../kernel/exit.c: 404
6290 2 sys_fork ../arch/i386/kernel/process.c: 771
6291 3 sys_read ../fs/read_write.c: 117
6292 4 sys_write ../fs/read_write.c: 146
6293 5 sys_open ../fs/open.c: 754
6294 6 sys_close ../fs/open.c: 839
6295 7 sys_waitpid ../kernel/exit.c: 503
6296 8 sys_creat ../fs/open.c: 789
6297 9 sys_link ../fs/namei.c: 1213
6298 10 sys_unlink ../fs/namei.c: 1074
6299 11 sys_execve ../arch/i386/kernel/process.c: 806
6300 ...
6301
6302 Find the system call number of the select system call:
6303
6304 crash> sys -c select
6305 NUM SYSTEM CALL FILE AND LINE NUMBER
6306 65 sys_select ../fs/select.c: 259
6307
6308 If the current output radix has been set to 16, the system call numbers
6309 will be displayed in hexadecimal.
6310
6311 Dump the DMI string data:
6312
6313 crash> sys -i
6314 DMI_BIOS_VENDOR: LENOVO
6315 DMI_BIOS_VERSION: G4ET37WW (1.12 )
6316 DMI_BIOS_DATE: 05/29/2012
6317 DMI_SYS_VENDOR: LENOVO
6318 DMI_PRODUCT_NAME: 2429BQ1
6319 DMI_PRODUCT_VERSION: ThinkPad T530
6320 DMI_PRODUCT_SERIAL: R9R91HZ
6321 DMI_PRODUCT_UUID: 568DFA01-5180-11CB-B851-BD06085ADDB0
6322 DMI_BOARD_VENDOR: LENOVO
6323 DMI_BOARD_NAME: 2429BQ1
6324 DMI_BOARD_VERSION: Not Available
6325 DMI_BOARD_SERIAL: 1ZLV127F17M
6326 DMI_BOARD_ASSET_TAG: Not Available
6327 DMI_CHASSIS_VENDOR: LENOVO
6328 DMI_CHASSIS_TYPE: 10
6329 DMI_CHASSIS_VERSION: Not Available
6330 DMI_CHASSIS_SERIAL: R9R91HZ
6331 DMI_CHASSIS_ASSET_TAG: RH0004111
6332
6333
6334 NAME
6335 q - exit this session
6336
6337 SYNOPSIS
6338 q
6339
6340 DESCRIPTION
6341 Bail out of the current crash session.
6342
6343 NOTE
6344 This command is equivalent to the "exit" command.
6345

```

[illegible]


```

6419         bits set: 36
6420
6421
6422 NAME
6423     list - linked list
6424
6425 SYNOPSIS
6426     list [[-o] offset][-e end][-[s|S] struct[.member[,member] [-l offset]] -[x|d]]
6427         [-r|-B] [-h [-O head_offset]|-H] start
6428
6429 DESCRIPTION
6430
6431     This command dumps the contents of a linked list. The entries in a linked
6432     list are typically data structures that are tied together in one of two
6433     formats:
6434
6435     1. A starting address points to a data structure; that structure contains
6436         a member that is a pointer to the next structure, and so on. This type
6437         of a singly-linked list typically ends when a "next" pointer value
6438         contains one of the following:
6439
6440             (a) a NULL pointer.
6441             (b) a pointer to the start address.
6442             (c) a pointer to the first item pointed to by the start address.
6443             (d) a pointer to its containing structure.
6444
6445     2. Most Linux lists of data structures are doubly-linked using "list_head"
6446         structures that are embedded members of the data structures in the list:
6447
6448         struct list_head {
6449             struct list_head *next, *prev;
6450         };
6451
6452     The linked list is typically headed by an external, standalone list_head,
6453     which is simply initialized to point to itself, signifying that the list
6454     is empty:
6455
6456         #define LIST_HEAD_INIT(name) { &(name), &(name) }
6457         #define LIST_HEAD(name) struct list_head name = LIST_HEAD_INIT(name)
6458
6459     In the case of list_head-linked lists, the "list_head.next" pointer is
6460     the address of a list_head structure that is embedded in the next data
6461     structure in the list, and not the address of the next data structure
6462     itself. The starting point of the list may be:
6463
6464         (a) an external, standalone, LIST_HEAD().
6465         (b) a list_head that is embedded within a data structure of the same
6466             type as the whole linked list.
6467         (c) a list_head that is embedded within a data structure that is
6468             different than the type of structures in the the linked list.
6469
6470     The list typically ends when the embedded "list_head.next" pointer of
6471     a data structure in the linked list points back to the LIST_HEAD()
6472     address. However, some list_head-linked lists have no defined starting
6473     point, but just link back onto themselves in a circular manner.
6474
6475     This command can handle both types of linked list; in both cases the list
6476     of addresses that are dumped are the addresses of the data structures
6477     themselves.
6478
6479     Alternatively, the address of a list_head, or other similar list linkage
6480     structure whose first member points to the next linkage structure, may be
6481     used as the starting address. The caveat with this type of usage is that
6482     the list may pass through, and display the address of, an external standalone
6483     list head which is not an address of a list linkage structure that is embedded
6484     within the data structure of interest.
6485
6486     The arguments are as follows:
6487
6488     [-o] offset    The offset within the structure to the "next" pointer
6489                    (default is 0). If non-zero, the offset may be entered
6490                    in either of two manners:
6491

```

1. In "structure.member" format; the "-o" is not necessary.
2. A number of bytes; the "-o" is only necessary on processors where the offset value could be misconstrued as a kernel virtual address.

- e end If the list ends in a manner unlike the typical manners that are described above, an explicit ending address value may be entered.
- s struct For each address in list, format and print as this type of structure; use the "struct.member" format in order to display a particular member of the structure. To display multiple members of a structure, use a comma-separated list of members. If any structure member contains an embedded structure or is an array, the output may be restricted to the embedded structure or an array element by expressing the struct argument as "struct.member.member" or "struct.member[index]"; embedded member specifications may extend beyond one level deep by expressing the argument as "struct.member.member.member...".
- S struct Similar to -s, but instead of parsing gdb output, member values are read directly from memory, so the command works much faster for 1-, 2-, 4-, and 8-byte members.
- O offset Only used in conjunction with -h; it specifies the offset of head node list_head embedded within a data structure which is different than the offset of list_head of other nodes embedded within a data structure. The offset may be entered in either of the following manners:
 1. in "structure.member" format.
 2. a number of bytes.
- l offset Only used in conjunction with -s, if the start address argument is a pointer to an embedded list head (or any other similar list linkage structure whose first member points to the next linkage structure), the offset to the embedded member may be entered in either of the following manners:
 1. in "structure.member" format.
 2. a number of bytes.
- x Override the default output format with hexadecimal format.
- d Override the default output format with decimal format.
- r For a list linked with list_head structures, traverse the list in the reverse order by using the "prev" pointer instead of "next".
- B Use the algorithm from R. P. Brent to detect loops instead of using a hash table. This algorithm uses a tiny fixed amount of memory and so is especially helpful for longer lists. The output is slightly different than the normal list output as it will print the length of the loop, the start of the loop, and the first duplicate in the list.

The meaning of the "start" argument, which can be expressed symbolically, in hexadecimal format, or an expression evaluating to an address, depends upon whether the -h or -H option is pre-pended:

- start The address of the first data structure in the list.
- start When both the -s and -l options are used, the address of an embedded list_head or similar linkage structure whose first member points to the next linkage structure.
- H start The address of a list_head structure, typically that of an external, standalone LIST_HEAD(). The list typically ends when the embedded "list_head.next" of a data structure in the linked list points back to this "start" address.
- h start The address of a data structure which contains an embedded list_head. The list typically ends when the embedded "list_head.next" of a data structure in the linked list points back to the embedded list_head contained in the data structure whose address is this "start" argument.

WARNING

When the "-h start" option is used, it is possible that the list_head-linked list will:

1. pass through an external standalone LIST_HEAD(), or
2. pass through a list_head that is the actual starting list_head, but is contained within a data structure that is not the same type as all of the other data structures in the list.

When that occurs, the data structure address displayed for that list_head will be incorrect, because the "-h start" option presumes that all list_head structures in the list are contained within the same type of data structure. Furthermore, if the "-s struct[.member[,member]]" option is used, it will display bogus data for that particular list_head.

A similar issue may be encountered when the "start" address is an embedded list_head or similar linkage structure whose first member points to the next linkage structure. When that occurs, the address of any external list head will not be distinguishable from the addresses that are embedded in the data structure of interest. Furthermore, if the "-s" and "-l" options are used, it will display bogus structure data when passing through any external list head structure that is not embedded in the specified data structure type.

EXAMPLES

Note that each task_struct is linked to its parent's task_struct via the p_pptr member:

```
crash> struct task_struct.p_pptr
struct task_struct {
    [136] struct task_struct *p_pptr;
}
```

That being the case, given a task_struct pointer of c169a000, show its parental hierarchy back to the "init_task" (the "swapper" task):

```
crash> list task_struct.p_pptr c169a000
c169a000
c0440000
c50d0000
c0562000
c0d28000
c7894000
c6a98000
c009a000
c0252000
```

Given that the "task_struct.p_pptr" offset is 136 bytes, the same result could be accomplished like so:

```
crash> list 136 c169a000
c169a000
c0440000
c50d0000
c0562000
c0d28000
c7894000
c6a98000
c009a000
c0252000
```

The list of currently-registered file system types are headed up by a struct file_system_type pointer named "file_systems", and linked by the "next" field in each file_system_type structure. The following sequence displays the structure address followed by the name and fs_flags members of each registered file system type:

```
crash> p file_systems
file_systems = $1 = (struct file_system_type *) 0xc03adc90
crash> list file_system_type.next -s file_system_type.name,fs_flags c03adc90
c03adc90
    name = 0xc02c05c8 "rootfs",
    fs_flags = 0x30,
c03abf94
    name = 0xc02c0319 "bdev",
    fs_flags = 0x10,
c03acb40
    name = 0xc02c07c4 "proc",
```

```

6638         fs_flags = 0x8,
6639     c03e9834
6640         name = 0xc02cfc83 "sockfs",
6641         fs_flags = 0x10,
6642     c03ab8e4
6643         name = 0xc02bf512 "tmpfs",
6644         fs_flags = 0x20,
6645     c03ab8c8
6646         name = 0xc02c3d6b "shm",
6647         fs_flags = 0x20,
6648     c03ac394
6649         name = 0xc02c03cf "pipefs",
6650         fs_flags = 0x10,
6651     c03ada74
6652         name = 0xc02c0e6b "ext2",
6653         fs_flags = 0x1,
6654     c03adc74
6655         name = 0xc02c0e70 "ramfs",
6656         fs_flags = 0x20,
6657     c03ade74
6658         name = 0xc02c0e76 "hugetlbfs",
6659         fs_flags = 0x20,
6660     c03adf8c
6661         name = 0xc02c0f84 "iso9660",
6662         fs_flags = 0x1,
6663     c03aec14
6664         name = 0xc02c0ffd "devpts",
6665         fs_flags = 0x8,
6666     c03e93f4
6667         name = 0xc02cf1b9 "pcihpfs",
6668         fs_flags = 0x28,
6669     e0831a14
6670         name = 0xe082f89f "ext3",
6671         fs_flags = 0x1,
6672     e0846af4
6673         name = 0xe0841ac6 "usbdevfs",
6674         fs_flags = 0x8,
6675     e0846b10
6676         name = 0xe0841acf "usbfs",
6677         fs_flags = 0x8,
6678     e0992370
6679         name = 0xe099176c "autofs",
6680         fs_flags = 0x0,
6681     e2dcc030
6682         name = 0xe2dc8849 "nfs",
6683         fs_flags = 0x48000,
6684

```

In some kernels, the system run queue is a linked list headed up by the "runqueue_head", which is defined like so:

```
static LIST_HEAD(runqueue_head);
```

The run queue linking is done with the "run_list" member of the task_struct:

```

crash> struct task_struct.run_list
struct task_struct {
    [60] struct list_head run_list;
}

```

Therefore, to view the list of task_struct addresses in the run queue, either of the following commands will work:

```

crash> list task_struct.run_list -H runqueue_head
f79ac000
f7254000
f7004000
crash> list 60 -H runqueue_head
f79ac000
f7254000
f7004000

```

In some kernel versions, the vfsmount structures of the mounted filesystems are linked by the LIST_HEAD "vfsmntlist", which uses the

mnt_list list_head of each vfsmount structure in the list. To dump each vfsmount structure in the list, append the -s option:

```
crash> list -H vfsmntlist vfsmount.mnt_list -s vfsmount
c3fc9e60
struct vfsmount {
    mnt_hash = {
        next = 0xc3fc9e60,
        prev = 0xc3fc9e60
    },
    mnt_parent = 0xc3fc9e60,
    mnt_mountpoint = 0xc3fc5dc0,
    mnt_root = 0xc3fc5dc0,
    mnt_instances = {
        next = 0xc3f60a74,
        prev = 0xc3f60a74
    },
    mnt_sb = 0xc3f60a00,
    mnt_mounts = {
        next = 0xf7445e08,
        prev = 0xf7445f88
    },
    mnt_child = {
        next = 0xc3fc9e88,
        prev = 0xc3fc9e88
    },
    mnt_count = {
        counter = 209
    },
    mnt_flags = 0,
    mnt_devname = 0xc8465b20 "/dev/root",
    mnt_list = {
        next = 0xf7445f9c,
        prev = 0xc02eb828
    },
    mnt_owner = 0
}
f7445f60
struct vfsmount {
...
```

The task_struct of every task in the system is linked into a circular list by its embedded "tasks" list_head. Show the task_struct addresses and the pids of all tasks in the system using "-h" option, starting with the task_struct at ffff88012b98e040:

```
crash> list task_struct.tasks -s task_struct.pid -h ffff88012b98e040
ffff88012b98e040
    pid = 14187
ffff8801277be0c0
    pid = 14248
fffffffff81a2d020
    pid = 0
ffff88012d7dd4c0
    pid = 1
ffff88012d7dca80
    pid = 2
ffff88012d7dc040
    pid = 3
ffff88012d7e9500
    pid = 4
...
ffff88012961a100
    pid = 14101
ffff880129017580
    pid = 14134
ffff8801269ed540
    pid = 14135
ffff880128256080
    pid = 14138
ffff88012b8f4100
    pid = 14183
```

Similar to the above, display the embedded sched_entity structure's on_rq member from each task_struct in the system:

```
crash> list task_struct.tasks -s task_struct.se.on_rq -h ffff8800b66a0000
ffff8800b66a0000
    se.on_rq = 1,
ffff8800b66a0ad0
    se.on_rq = 0,
ffff8800b66a15a0
    se.on_rq = 0,
ffff8800b66a2070
    se.on_rq = 0,
ffff8800b66a2b40
    se.on_rq = 0,
ffff8800b67315a0
    se.on_rq = 0,
ffff8800b6732b40
    se.on_rq = 0,
...
```

The task_struct.tasks example above requires that the -h option be given the address of a task_struct. Alternatively, the -l option can be given the address of a list_head or similar linkage structure whose first member points to the next linkage structure. Again using the task_struct.tasks embedded list head, dump the "comm" member of all tasks by using -l in conjunction with -s option:

```
crash> task -R tasks.next
PID: 7044   TASK: ffff88005ac10000   CPU: 2   COMMAND: "crash"
tasks.next = 0xffff880109b8e3d0,
crash> list 0xffff880109b8e3d0 -l task_struct.tasks -s task_struct.comm
ffff880109b8e3d0
    comm = "kworker/1:2"
ffff880109b8be00
    comm = "bash"
ffff88019d26c590
    comm = "cscope"
ffff880109b8b670
    comm = "kworker/0:1"
ffff880109b8cd20
    comm = "kworker/1:0"
ffff88005ac15c40
    comm = "vi"
ffff88005ac11fc0
    comm = "sleep"
fffffffff81c135c0
    comm = "swapper/0"
ffff880212828180
    comm = "systemd"
...
ffff8801288d1830
    comm = "chrome"
ffff8801534dd4b0
    comm = "kworker/0:0"
ffff8801534d8180
    comm = "kworker/1:1"
ffff88010902b670
    comm = "kworker/2:2"
ffff880109b8a750
    comm = "sudo"
ffff88005ac10180
    comm = "crash"
```

To display a linked list whose head node and other nodes are embedded within either same or different data structures resulting in different offsets for head node and other nodes, e.g. dentry.d_subdirs and dentry.d_child, the -O option can be used:

```
crash> list -o dentry.d_child -s dentry.d_name.name -O dentry.d_subdirs -h
ffff9c585b81a180
ffff9c585b9cb140
    d_name.name = 0xffff9c585b9cb178 ccc.txt
ffff9c585b9cb980
```

```
6856     d_name.name = 0xffff9c585b9cb9b8 bbb.txt
6857     ffff9c585b9cb740
6858     d_name.name = 0xffff9c585b9cb778 aaa.txt
6859
```

6860 The dentry.d_subdirs example above is equal to the following sequence:

```
6861
6862     crash> struct -o dentry.d_subdirs ffff9c585b81a180
6863     struct dentry {
6864         [ffff9c585b81a220] struct list_head d_subdirs;
6865     }
6866     crash> list -o dentry.d_child -s dentry.d_name.name -H ffff9c585b81a220
6867
```

6868 NAME

6869 rd - read memory

6870 SYNOPSIS

```
6871     rd [-adDsSupxmfNR] [-8|-16|-32|-64] [-o offs] [-e addr] [-r file] [address|symbol]
6872     [count]
```

6873 DESCRIPTION

6874 This command displays the contents of memory, with the output formatted
6875 in several different manners. The starting address may be entered either
6876 symbolically or by address. The default output size is the size of a long
6877 data type, and the default output format is hexadecimal. When hexadecimal
6878 output is used, the output will be accompanied by an ASCII translation.

```
6883     -p address argument is a physical address.
6884     -u address argument is a user virtual address; only required on
6885     processors with common user and kernel virtual address spaces.
6886     -m address argument is a xen host machine address.
6887     -f address argument is a dumpfile offset.
6888     -d display output in signed decimal format (default is hexadecimal).
6889     -D display output in unsigned decimal format (default is hexadecimal).
6890     -s displays output symbolically when appropriate.
6891     -S[S] displays output symbolically when appropriate; if the memory
6892     contents reference a slab cache object, the name of the slab cache
6893     will be displayed in brackets. If -S is entered twice, and the
6894     memory contents reference a slab cache object, both the memory
6895     contents and the name of the slab cache will be displayed in
6896     brackets.
6897     -x do not display ASCII translation at end of each line.
6898     -8 display output in 8-bit values.
6899     -16 display output in 16-bit values.
6900     -32 display output in 32-bit values (default on 32-bit machines).
6901     -64 display output in 64-bit values (default on 64-bit machines).
6902     -a display output in ASCII characters if the memory contains printable
6903     ASCII characters; if no count argument is entered, stop at the first
6904     non-printable character.
6905     -N display output in network byte order (only valid for 16- and 32-bit
6906     values)
6907     -R display memory in reverse order; memory will be displayed up to and
6908     including the address argument, requiring the count argument to be
6909     greater than 1 in order to display memory before the specified
6910     address.
6911     -o offs offset the starting address by offs.
6912     -e addr display memory until reaching specified ending hexadecimal address.
6913     -r file dumps raw data to the specified output file; the number of bytes that
6914     are copied to the file must be specified either by a count argument
6915     or by the -e option.
6916     address starting hexadecimal address:
6917         1 the default presumes a kernel virtual address.
6918         2. -p specifies a physical address.
6919         3. -u specifies a user virtual address, but is only necessary on
6920         processors with common user and kernel virtual address spaces.
6921     symbol symbol of starting address to read.
6922     count number of memory locations to display; if entered, it must be the
6923     last argument on the command line; if not entered, the count defaults
6924     to 1, or unlimited for -a; when used with the -r option, it is the
6925     number of bytes to be written to the file.
```

6926 EXAMPLES

6927 Display the kernel's version string:

```

6929
6930 crash> rd -a linux_banner
6931 c082a020: Linux version 2.6.32-119.el6.i686 (mockbuild@hs20-bc2-4.buil
6932 c082a05c: d.redhat.com) (gcc version 4.4.4 20100726 (Red Hat 4.4.4-13)
6933 c082a098: (GCC) ) #1 SMP Tue Mar 1 18:16:57 EST 2011
6934

```

Display the same block of memory, first without symbols, again with symbols, and then with symbols and slab cache references:

```

6937
6938 crash> rd f6e31f70 28
6939 f6e31f70: f6e31f6c f779c180 c04a4032 00a9dd40 1.....y.2@J.@...
6940 f6e31f80: 00000fff c0472da0 f6e31fa4 f779c180 .....-G.....y.
6941 f6e31f90: ffffffff7 00a9b70f f6e31000 c04731ee .....1G.
6942 f6e31fa0: f6e31fa4 00000000 00000000 00000000 .....
6943 f6e31fb0: 00000000 00a9dd40 c0404f17 00000000 ....@....O@.....
6944 f6e31fc0: 00a9dd40 00000fff 00a9dd40 00a9b70f @.....@.....
6945 f6e31fd0: bf9e2718 ffffffffda c040007b 0000007b .'.....{.@.{...
6946 crash> rd -s f6e31f70 28
6947 f6e31f70: f6e31f6c f779c180 kmsg_read 00a9dd40
6948 f6e31f80: 00000fff vfs_read+159 f6e31fa4 f779c180
6949 f6e31f90: ffffffff7 00a9b70f f6e31000 sys_read+60
6950 f6e31fa0: f6e31fa4 00000000 00000000 00000000
6951 f6e31fb0: 00000000 00a9dd40 syscall_call+7 00000000
6952 f6e31fc0: 00a9dd40 00000fff 00a9dd40 00a9b70f
6953 f6e31fd0: bf9e2718 ffffffffda startup_32+123 0000007b
6954 crash> rd -S f6e31f70 28
6955 f6e31f70: [size-4096] [filp] kmsg_read 00a9dd40
6956 f6e31f80: 00000fff vfs_read+159 [size-4096] [filp]
6957 f6e31f90: ffffffff7 00a9b70f [size-4096] sys_read+60
6958 f6e31fa0: [size-4096] 00000000 00000000 00000000
6959 f6e31fb0: 00000000 00a9dd40 syscall_call+7 00000000
6960 f6e31fc0: 00a9dd40 00000fff 00a9dd40 00a9b70f
6961 f6e31fd0: bf9e2718 ffffffffda startup_32+123 0000007b
6962 crash> rd -SS f6e31f70 28
6963 f6e31f70: [f6e31f6c:size-4096] [f779c180:filp] kmsg_read 00a9dd40
6964 f6e31f80: 00000fff vfs_read+159 [f6e31fa4:size-4096] [f779c180:filp]
6965 f6e31f90: ffffffff7 00a9b70f [f6e31000:size-4096] sys_read+60
6966 f6e31fa0: [f6e31fa4:size-4096] 00000000 00000000 00000000
6967 f6e31fb0: 00000000 00a9dd40 syscall_call+7 00000000
6968 f6e31fc0: 00a9dd40 00000fff 00a9dd40 00a9b70f
6969 f6e31fd0: bf9e2718 ffffffffda startup_32+123 0000007b
6970

```

Read jiffies in hexadecimal and decimal format:

```

6971
6972 crash> rd jiffies
6973 c0213ae0: 0008cc3a :...
6974
6975 crash> rd -d jiffies
6976 c0213ae0: 577376
6977
6978

```

Access the same memory in different sizes:

```

6979
6980 crash> rd -64 kernel_version
6981 c0226a6c: 35312d352e322e32 2.2.5-15
6982
6983 crash> rd -32 kernel_version 2
6984 c0226a6c: 2e322e32 35312d35 2.2.5-15
6985
6986 crash> rd -16 kernel_version 4
6987 c0226a6c: 2e32 2e32 2d35 3531 2.2.5-15
6988
6989 crash> rd -8 kernel_version 8
6990 c0226a6c: 32 2e 32 2e 35 2d 31 35 2.2.5-15
6991
6992

```

Read the range of memory from c009bf2c to c009bf60:

```

6993
6994 crash> rd c009bf2c -e c009bf60
6995 c009bf2c: c009bf64 c01328c3 c009bf64 c0132838 d....(..d...8(..
6996 c009bf3c: 0000002a 00000004 c57d77e8 00000104 *......w}.....
6997 c009bf4c: 0000000b c009a000 7fffffff 00000000 .....
6998 c009bf5c: 00000000 ....
6999
7000
7001

```



```

7002 NAME
7003     task - task_struct and thread_info contents
7004
7005 SYNOPSIS
7006     task [-R member[,member]] [-dx] [pid | taskp] ...
7007
7008 DESCRIPTION
7009     This command dumps a formatted display of the contents of a task's
7010     task_struct and thread_info structures. Multiple task or PID numbers
7011     may be entered; if no arguments are entered, the task_struct and
7012     thread_info structures of the current context are displayed. The -R option,
7013     which may also be invoked indirectly via "foreach task", pares the output
7014     down to one or more structure members.
7015
7016     pid    a process PID.
7017     taskp   a hexadecimal task_struct pointer.
7018     -R member  a comma-separated list of one or more task_struct and/or
7019                thread_info structure members. If any member contains an embedded
7020                structure, or is an array, the output may be restricted to the
7021                embedded structure or an array element by expressing the member
7022                argument as "member.member" or "member[index]"; embedded member
7023                specifications may extend beyond one level deep, by expressing the
7024                member argument as "member.member.member...".
7025     -x      override default output format with hexadecimal format.
7026     -d      override default output format with decimal format.
7027
7028 EXAMPLES
7029     Dump the task_struct and thread_info structures of the current context
7030     in hexadecimal format:
7031
7032     crash> task -x
7033     PID: 3176   TASK: f2451550   CPU: 1   COMMAND: "memtest"
7034     struct task_struct {
7035         state = 0x0,
7036         stack = 0xf05b6000,
7037         usage = {
7038             counter = 0x2
7039         },
7040         flags = 0x402040,
7041         ptrace = 0x0,
7042         lock_depth = 0xffffffff,
7043         prio = 0x78,
7044         static_prio = 0x78,
7045         normal_prio = 0x78,
7046         rt_priority = 0x0,
7047     ...
7048     perf_event_ctxp = {0x0, 0x0},
7049     memcg_batch = {
7050         do_batch = 0x0,
7051         memcg = 0x0,
7052         bytes = 0x0,
7053         memsw_bytes = 0x0
7054     }
7055 }
7056
7057     struct thread_info {
7058         task = 0xf2451550,
7059         exec_domain = 0xc0a60860,
7060         flags = 0x88,
7061         status = 0x0,
7062         cpu = 0x1,
7063         preempt_count = 0x4010000,
7064         addr_limit = {
7065             seg = 0xc0000000
7066         },
7067         restart_block = {
7068     ...
7069
7070     Display the ngroups and groups task_struct members for PID 2958:
7071
7072     crash> task -R ngroups,groups 2958
7073     PID: 2958   TASK: c6718000   CPU: 0   COMMAND: "bash"
7074     ngroups = 6,

```

```
7075     groups = {504, 8, 9, 1000, 1007, 1006, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
7076               0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0},
7077
```

```
7078 Display the embedded sched_entity structure's on_rq member:
```

```
7079
7080 crash> task -R se.on_rq
7081 PID: 6529   TASK: ffff880116538790   CPU: 1   COMMAND: "bash"
7082     se.on_rq = 1,
7083
```

```
7084 Display the 3rd pid_link structure in the embedded pids[] array:
```

```
7085
7086 crash> task -R pids[2]
7087 PID: 6529   TASK: ffff880116538790   CPU: 0   COMMAND: "bash"
7088     pids[2] = {
7089         node = {
7090             next = 0xffff8801165391b0,
7091             pprev = 0xffff880209d011b0
7092         },
7093         pid = 0xffff8801f0876e00
7094     }
7095
```

```
7096 NOTE: When this command is invoked directly (i.e., not from "foreach"), it
7097 is not necessary to include the "-R" before the task_struct/thread_info
7098 member name(s).
7099
```

7101 NAME

```
7102     exit - exit this session
7103
```

7104 SYNOPSIS

```
7105     exit
7106
```

7107 DESCRIPTION

```
7108     Bail out of the current crash session.
7109
```

7110 NOTE

```
7111     This command is equivalent to the "q" command.
7112
7113
```

7114 NAME

```
7115     log - dump system message buffer
7116
```

7117 SYNOPSIS

```
7118     log [-Ttdmas]
7119
```

7120 DESCRIPTION

```
7121     This command dumps the kernel log_buf contents in chronological order. The
7122     command supports the older log_buf formats, which may or may not contain a
7123     timestamp inserted prior to each message, as well as the newer variable-length
7124     record format, where the timestamp is contained in each log entry's header.
7125
```

- 7126 -T Display the message text with human readable timestamp.
7127 (Be aware that the timestamp could be inaccurate! The timestamp is
7128 from local_clock(), which is different from the elapsed wall time.)
- 7129 -t Display the message text without the timestamp; only applicable to the
7130 variable-length record format.
- 7131 -d Display the dictionary of key/value pair properties that are optionally
7132 appended to a message by the kernel's dev_printk() function; only
7133 applicable to the variable-length record format.
- 7134 -m Display the message log level in brackets preceding each message. For
7135 the variable-length record format, the level will be displayed in
7136 hexadecimal. In older kernels, by default, the facility/flag bits
7137 will be stripped to only show the level, but if needed, can still be
7138 shown with 'set debug 1'.
- 7139 -a Dump the audit logs remaining in kernel audit buffers that have not
7140 been copied out to the user-space audit daemon.
- 7141 -s Dump the printk logs remaining in kernel safe per-CPU buffers that
7142 have not been flushed out to log_buf.

7145 EXAMPLES

```
7146     Dump the kernel message buffer:
7147
```

```

7148 crash> log
7149 Linux version 2.2.5-15smp (root@mclinux1) (gcc version egcs-2.91.66 19990
7150 314/Linux (egcs-1.1.2 release)) #1 SMP Thu Aug 26 11:04:37 EDT 1999
7151 Intel MultiProcessor Specification v1.4
7152 Virtual Wire compatibility mode.
7153 OEM ID: DELL Product ID: WS 410 APIC at: 0xFEE00000
7154 Processor #0 Pentium(tm) Pro APIC version 17
7155 Processor #1 Pentium(tm) Pro APIC version 17
7156 I/O APIC #2 Version 17 at 0xFEC00000.
7157 Processors: 2
7158 mapped APIC to fffffe000 (fee000000)
7159 mapped IOAPIC to fffffd000 (fec000000)
7160 Detected 447696347 Hz processor.
7161 Console: colour VGA+ 80x25
7162 Calibrating delay loop... 445.64 BogoMIPS
7163 ...
7164 8K byte-wide RAM 5:3 Rx:Tx split, autoselect/Autonegotiate interface.
7165 MII transceiver found at address 24, status 782d.
7166 Enabling bus-master transmits and whole-frame receives.
7167 Installing knfsd (copyright (C) 1996 okir@monad.swb.de).
7168 nfds_init: initialized fhcache, entries=256
7169 ...

```

Do the same thing, but also show the log level preceding each message:

```

7170
7171 crash> log -m
7172
7173 <4>Linux version 2.2.5-15smp (root@mclinux1) (gcc version egcs-2.91.66 19990
7174 314/Linux (egcs-1.1.2 release)) #1 SMP Thu Aug 26 11:04:37 EDT 1999
7175 <4>Intel MultiProcessor Specification v1.4
7176 <4> Virtual Wire compatibility mode.
7177 <4> OEM ID: DELL Product ID: WS 410 APIC at: 0xFEE00000
7178 <4> Processor #0 Pentium(tm) Pro APIC version 17
7179 <4> Processor #1 Pentium(tm) Pro APIC version 17
7180 <4> I/O APIC #2 Version 17 at 0xFEC00000.
7181 <4> Processors: 2
7182 <4> mapped APIC to fffffe000 (fee000000)
7183 <4> mapped IOAPIC to fffffd000 (fec000000)
7184 <4> Detected 447696347 Hz processor.
7185 <4> Console: colour VGA+ 80x25
7186 <4> Calibrating delay loop... 445.64 BogoMIPS
7187 ...
7188 <6> 8K byte-wide RAM 5:3 Rx:Tx split, autoselect/Autonegotiate interface.
7189 <6> MII transceiver found at address 24, status 782d.
7190 <6> Enabling bus-master transmits and whole-frame receives.
7191 <6> Installing knfsd (copyright (C) 1996 okir@monad.swb.de).
7192 <7> nfds_init: initialized fhcache, entries=256
7193 ...
7194
7195

```

On a system with the variable-length record format, and whose log_buf has been filled and wrapped around, display the log with timestamp data:

```

7196 crash> log
7197
7198 [ 0.467730] pci 0000:ff:02.0: [8086:2c10] type 00 class 0x060000
7199 [ 0.467749] pci 0000:ff:02.1: [8086:2c11] type 00 class 0x060000
7200 [ 0.467769] pci 0000:ff:02.4: [8086:2c14] type 00 class 0x060000
7201 [ 0.467788] pci 0000:ff:02.5: [8086:2c15] type 00 class 0x060000
7202 [ 0.467809] pci 0000:ff:03.0: [8086:2c18] type 00 class 0x060000
7203 [ 0.467828] pci 0000:ff:03.1: [8086:2c19] type 00 class 0x060000
7204 ...
7205
7206

```

Display the same message text as above, without the timestamp data:

```

7207 crash> log -t
7208
7209 pci 0000:ff:02.0: [8086:2c10] type 00 class 0x060000
7210 pci 0000:ff:02.1: [8086:2c11] type 00 class 0x060000
7211 pci 0000:ff:02.4: [8086:2c14] type 00 class 0x060000
7212 pci 0000:ff:02.5: [8086:2c15] type 00 class 0x060000
7213 pci 0000:ff:03.0: [8086:2c18] type 00 class 0x060000
7214 pci 0000:ff:03.1: [8086:2c19] type 00 class 0x060000
7215 ...
7216
7217

```

Display the same message text as above, with appended dictionary data:

```

7218
7219
7220

```

```
7221 crash> log -td
7222 pci 0000:ff:02.0: [8086:2c10] type 00 class 0x060000
7223 SUBSYSTEM=pci
7224 DEVICE+=pci:0000:ff:02.0
7225 pci 0000:ff:02.1: [8086:2c11] type 00 class 0x060000
7226 SUBSYSTEM=pci
7227 DEVICE+=pci:0000:ff:02.1
7228 pci 0000:ff:02.4: [8086:2c14] type 00 class 0x060000
7229 SUBSYSTEM=pci
7230 DEVICE+=pci:0000:ff:02.4
7231 pci 0000:ff:02.5: [8086:2c15] type 00 class 0x060000
7232 SUBSYSTEM=pci
7233 DEVICE+=pci:0000:ff:02.5
7234 pci 0000:ff:03.0: [8086:2c18] type 00 class 0x060000
7235 SUBSYSTEM=pci
7236 DEVICE+=pci:0000:ff:03.0
7237 pci 0000:ff:03.1: [8086:2c19] type 00 class 0x060000
7238 SUBSYSTEM=pci
7239 DEVICE+=pci:0000:ff:03.1
7240 ...
7241
```

7242 Dump the kernel audit logs:

```
7243
7244 crash> log -a
7245 type=1320 audit(1489384479.809:4342):
7246 type=1300 audit(1489384479.809:4343): arch=c000003e syscall=0 success=yes
7247 exit=0 a0=4 a1=7f84154a2000 a2=400 a3=22 items=0 ppid=2560 pid=2591 auid=0
7248 uid=0 gid=0 euid=0 suid=0 fsuid=0 egid=0 sgid=0 fsgid=0 tty=ttyS0 ses=1
7249 comm="pidof" exe="/usr/sbin/killall5"
7250 subj=unconfined_u:unconfined_r:unconfined_t:s0-s0:c0.c1023 key=(null)
7251 type=1320 audit(1489384479.809:4343):
7252 type=1300 audit(1489384479.809:4344): arch=c000003e syscall=3 success=yes
7253 exit=0 a0=4 a1=1 a2=8 a3=0 items=0 ppid=2560 pid=2591 auid=0 uid=0 gid=0
7254 euid=0 suid=0 fsuid=0 egid=0 sgid=0 fsgid=0 tty=ttyS0 ses=1 comm="pidof"
7255 exe="/usr/sbin/killall5"
7256 subj=unconfined_u:unconfined_r:unconfined_t:s0-s0:c0.c1023 key=(null)
7257 type=1320 audit(1489384479.809:4344):
7258 type=1300 audit(1489384479.809:4345): arch=c000003e syscall=11
7259 success=yes exit=0 a0=7f84154a2000 a1=1000 a2=0 a3=0 items=0 ppid=2560
7260 pid=2591 auid=0 uid=0 gid=0 euid=0 suid=0 fsuid=0 egid=0 sgid=0 fsgid=0
7261 tty=ttyS0 ses=1 comm="pidof" exe="/usr/sbin/killall5"
7262 subj=unconfined_u:unconfined_r:unconfined_t:s0-s0:c0.c1023 key=(null)
7263 type=1320 audit(1489384479.809:4345):
7264 type=1300 audit(1489384479.809:4346): arch=c000003e syscall=2 success=yes
7265 exit=4 a0=7ffcfdf20f5a0 a1=0 a2=1b6 a3=24 items=1 ppid=2560 pid=2591 auid=0
7266 uid=0 gid=0 euid=0 suid=0 fsuid=0 egid=0 sgid=0 fsgid=0 tty=ttyS0 ses=1
7267 comm="pidof" exe="/usr/sbin/killall5"
7268 subj=unconfined_u:unconfined_r:unconfined_t:s0-s0:c0.c1023 key=(null)
7269 type=1307 audit(1489384479.809:4346): cwd="/proc"
7270 ...
7271
```

7272 Display the message text with human readable timestamp:

```
7273
7274 crash> log -T
7275 [Sat Apr  4 07:41:09 2020] BIOS-e820: [mem 0x0000000000000000-0x000000000009fbff]
usable
7276 [Sat Apr  4 07:41:09 2020] BIOS-e820: [mem 0x000000000009fc00-0x000000000009ffff]
reserved
7277 [Sat Apr  4 07:41:09 2020] BIOS-e820: [mem 0x00000000000f0000-0x00000000000fffff]
reserved
7278 [Sat Apr  4 07:41:09 2020] BIOS-e820: [mem 0x0000000000100000-0x0000000000dfffff]
usable
7279 [Sat Apr  4 07:41:09 2020] BIOS-e820: [mem 0x000000000dfff000-0x000000000dfffffff]
ACPI data
7280 [Sat Apr  4 07:41:09 2020] BIOS-e820: [mem 0x00000000fec00000-0x00000000fec00fff]
reserved
7281 [Sat Apr  4 07:41:09 2020] BIOS-e820: [mem 0x00000000fee00000-0x00000000fee00fff]
reserved
7282 [Sat Apr  4 07:41:09 2020] BIOS-e820: [mem 0x00000000fffc0000-0x00000000ffffffff]
reserved
7283 [Sat Apr  4 07:41:09 2020] BIOS-e820: [mem 0x0000000010000000-0x0000000011fffffff]
usable
7284 [Sat Apr  4 07:41:09 2020] NX (Execute Disable) protection: active
```

```

7285 [Sat Apr 4 07:41:09 2020] SMBIOS 2.5 present.
7286 [Sat Apr 4 07:41:09 2020] DMI: innotek GmbH VirtualBox/VirtualBox, BIOS
VirtualBox 12/01/2006
7287 [Sat Apr 4 07:41:09 2020] Hypervisor detected: KVM
7288 [Sat Apr 4 07:41:09 2020] kvm-clock: Using msrs 4b564d01 and 4b564d00
7289 [Sat Apr 4 07:41:09 2020] kvm-clock: cpu 0, msr 6de01001, primary cpu clock
7290 [Sat Apr 4 07:41:09 2020] kvm-clock: using sched offset of 11838753697 cycles
7291 [Sat Apr 4 07:41:09 2020] clocksource: kvm-clock: mask: 0xffffffffffffffff
max_cycles: 0x1cd42e4dffb, max_idle_ns: 881590591483 ns
7292 [Sat Apr 4 07:41:09 2020] e820: update [mem 0x00000000-0x00000fff] usable ==>
reserved
7293 [Sat Apr 4 07:41:09 2020] e820: remove [mem 0x000a0000-0x000fffff] usable
7294 [Sat Apr 4 07:41:09 2020] last_pfn = 0x120000 max_arch_pfn = 0x400000000
7295 [Sat Apr 4 07:41:09 2020] MTRR default type: uncachable
7296 [Sat Apr 4 07:41:09 2020] MTRR variable ranges disabled:
7297 ...
7298

```

On a system which has printk_safe_seq_buf buffer,
display its unflushed log with buffer name:

```
crash> log
```

```

...
[nmi_print_seq] Uhhuh. NMI received for unknown reason 30 on CPU 0.
[nmi_print_seq] Do you have a strange power saving mode enabled?
[nmi_print_seq] Dazed and confused, but trying to continue

```

Dump the printk safe buffers:

```
crash> log -s
```

```

PRINTK_SAFE_SEQ_BUF: nmi_print_seq
CPU: 0 ADDR: ffff8ca4fbc19ce0 LEN: 150 MESSAGE_LOST: 0
Uhhuh. NMI received for unknown reason 20 on CPU 0.
Do you have a strange power saving mode enabled?
Dazed and confused, but trying to continue
...
PRINTK_SAFE_SEQ_BUF: safe_print_seq
CPU: 0 ADDR: ffff8ca4fbc1ad00 LEN: 0 MESSAGE_LOST: 0
(empty)
...

```

NAME

repeat - repeat a command

SYNOPSIS

repeat [-seconds] command

DESCRIPTION

This command repeats a command indefinitely, optionally delaying a given number of seconds between each command execution.

```

-seconds    The number of seconds to delay between command executions.
             This option must precede the command name to be executed.

```

Command execution may be stopped with CTRL-C, or if scrolling is in effect, by entering "q". This command is meant for use on a live system; it is hard to conceive of a reason to use it when debugging a crash dump.

EXAMPLES

Display the value of jiffies once per second:

```

crash> repeat -1 p jiffies
jiffies = $1 = 155551079
jiffies = $2 = 155551180
jiffies = $3 = 155551281
jiffies = $4 = 155551382
jiffies = $5 = 155551483
jiffies = $6 = 155551584
jiffies = $7 = 155551685
jiffies = $8 = 155551786
jiffies = $9 = 155551887
jiffies = $10 = 155551988
jiffies = $11 = 155552089

```

```

7355     jiffies = $12 = 155552190
7356     jiffies = $13 = 155552291
7357     jiffies = $14 = 155552392
7358     jiffies = $15 = 155552493
7359     jiffies = $16 = 155552594
7360     jiffies = $17 = 155552695
7361     jiffies = $18 = 155552796
7362     ...
7363
7364
7365 NAME
7366     timer - timer queue data
7367
7368 SYNOPSIS
7369     timer [-r][-C cpu]
7370
7371 DESCRIPTION
7372     This command displays the timer queue entries, both old- and new-style,
7373     in chronological order. In the case of the old-style timers, the
7374     timer_table array index is shown; in the case of the new-style timers,
7375     the timer_list address is shown. On later kernels, the timer data is
7376     per-cpu.
7377
7378     -r Display hrtimer timer queue entries, both old- and new-style, in
7379     chronological order. In the case of the old-style hrtimers, the
7380     expiration time is a single value; in the new-style hrtimers, the
7381     expiration time is a range.
7382     -C cpu Restrict the output to one or more CPUs, where multiple cpu[s] can
7383     be specified, for example, as "1,3,5", "1-3", or "1,3,5-7,10".
7384
7385 EXAMPLES
7386     Display the timer queue on an SMP system:
7387
7388     crash> timer
7389     JIFFIES
7390     4296291038
7391     ...
7392     TIMER BASES[1][BASE_STD]: ffff9801aba5aa00
7393
7394     EXPIRES      TTE      TIMER_LIST      FUNCTION
7395     4296282997   -8041   ffff9801aba55ce0   ffffffff83a3bda0   <mce_timer_fn>
7396     4296283104   -7934   ffff97fd84bd35e0   ffffffff83ac6b70   <delayed_work_timer_fn>
7397     4296291061     23     ffffa6b283967de0   ffffffff83b29880   <process_timeout>
7398     4296291112     74     ffff9800c9b62ad8   ffffffff83e6b550   <cursor_timer_handler>
7399     4296291345    307     ffff980186d5ef88   ffffffff84146b80   <tcp_keepalive_timer>
7400     4296291484    446     ffff9801a7c54740   ffffffff84147f50   <tcp_write_timer>
7401     4296291997    959     ffffffff9c073f880   ffffffff83ac6b70   <delayed_work_timer_fn>
7402     4296296213    5175    ffffa6b28339be18   ffffffff83b29880   <process_timeout>
7403     4296304383   13345    ffff980194ca72a8   ffffffff8412e4e0   <tw_timer_handler>
7404     4296305724   14686    ffff980194ca6918   ffffffff8412e4e0   <tw_timer_handler>
7405     4296306036   14998    ffff980194ca6d58   ffffffff8412e4e0   <tw_timer_handler>
7406     4296306883   15845    ffff980194ca7e58   ffffffff8412e4e0   <tw_timer_handler>
7407     4296307588   16550    ffff9801aaa27e58   ffffffff8412e4e0   <tw_timer_handler>
7408     4296307625   16587    ffff980194ca6a28   ffffffff8412e4e0   <tw_timer_handler>
7409     4296313542   22504    ffff980194ca7c38   ffffffff8412e4e0   <tw_timer_handler>
7410     4296317680   26642    ffff9800c9149c58   ffffffff840da870   <neigh_timer_handler>
7411     4296317744   26706    ffff9801a5354468   ffffffff83ac6b70   <delayed_work_timer_fn>
7412     4296343322   52284    ffff980194ca63c8   ffffffff8412e4e0   <tw_timer_handler>
7413     4296343581   52543    ffff980194ca7088   ffffffff8412e4e0   <tw_timer_handler>
7414     4296343597   52559    ffff9801aaa274c8   ffffffff8412e4e0   <tw_timer_handler>
7415     4296714205   423167   ffffffff84caf3c0   ffffffff83ac6b70   <delayed_work_timer_fn>
7416     TIMER BASES[1][BASE_DEF]: ffff9801aba5bc80
7417
7418     EXPIRES      TTE      TIMER_LIST      FUNCTION
7419     4296291264     226     ffffffff855eb238   ffffffff83c08fb0   <writeout_period>
7420     4296319997    28959    ffffffff9c06ede40   ffffffff83ac6b70   <delayed_work_timer_fn>
7421     4296506084   215046    ffff9801aba629c8   ffffffff83ac5ea0   <idle_worker_timeout>
7422     ...
7423
7424 Display a new-style hrtimer queue:
7425
7426     crash> timer -r
7427     ...
7428     CPU: 2  HRTIMER_CPU_BASE: ffff9801aba9cf00
7429     CLOCK: 0  HRTIMER_CLOCK_BASE: ffff9801aba9cf40  [ktime_get]

```

```

7428         CURRENT
7429         1623742000000
7430         SOFTEXPIRES      EXPIRES      TTE      HRTIMER      FUNCTION
7431         1623741000000    1623741000000    -1000000    ffff9801aba9d540    ffffffff83b3c8e0
<tick_sched_timer>
7432         1624024000000    1624024000000    282000000    ffff9801aba9d720    ffffffff83b7e7a0
<watchdog_timer_fn>
7433         1626000939806    1626010929804    2268929804    ffffa6b28399fa40    ffffffff83b2c1e0
<hrtimer_wakeup>
7434         1627576915615    1627576915615    3834915615    ffff9801a5727978    ffffffff83b365c0
<posix_timer_fn>
7435         1627637194488    1627647194487    3905194487    ffffa6b283977db0    ffffffff83b2c1e0
<hrtimer_wakeup>
7436         1629937423000    1629937423000    6195423000    ffff9801a9af2900    ffffffff83cf3d30
<timerfd_tmrproc>

```

```

7437
7438         CLOCK: 1  HRTIMER_CLOCK_BASE: ffff9801aba9cf80  [ktime_get_real]
7439         CURRENT

```

```

7440         1558362388334558243
7441         SOFTEXPIRES      EXPIRES      TTE      HRTIMER
FUNCTION
7442         1558362389331238000    1558362389331288000      996729757    ffffa6b28574bcf0
ffffffff83b2c1e0  <hrtimer_wakeup>
7443         15583643720000000000    15583643720000000000    1983665441757    ffff9801a3513278
ffffffff83b365c0  <posix_timer_fn>

```

```

7444
7445         CLOCK: 2  HRTIMER_CLOCK_BASE: ffff9801aba9cfc0  [ktime_get_boottime]
7446         (empty)

```

```

7447         ...

```

```

7448
7449
7450

```

```

7451     NAME

```

```

7452         extend - extend the crash command set

```

```

7453

```

```

7454     SYNOPSIS

```

```

7455         extend [shared-object ...] | [-u [shared-object ...]] | -s

```

```

7456

```

```

7457     DESCRIPTION

```

```

7458         This command dynamically loads or unloads crash extension shared object
7459         libraries:

```

```

7460
7461         shared-object      load the specified shared object file; more than one
7462                             one object file may be entered.
7463         -u shared-object    unload the specified shared object file; if no file
7464                             arguments are specified, unload all objects.
7465         -s                  show all available shared object files.

```

```

7466

```

```

7467         If the shared-object filename is not expressed with a fully-qualified
7468         pathname, the following directories will be searched in the order shown,
7469         and the first instance of the file that is found will be selected:

```

```

7470

```

- 7471 1. the current working directory
- 7472 2. the directory specified in the CRASH_EXTENSIONS environment variable
- 7473 3. /usr/lib64/crash/extensions (64-bit architectures)
- 7474 4. /usr/lib/crash/extensions
- 7475 5. the ./extensions subdirectory of the current directory

```

7476

```

```

7477         If no arguments are entered, the current set of shared object files and
7478         a list of their commands will be displayed. The registered commands
7479         contained in each shared object file will appear automatically in the
7480         "help" command screen.

```

```

7481

```

```

7482         An example of a shared object prototype file, and how to compile it
7483         into a shared object, is appended below.

```

```

7484

```

```

7485     EXAMPLES

```

```

7486         Load two shared object files:

```

```

7487

```

```

7488         crash> extend extlib1.so extlib2.so
7489         ./extlib1.so: shared object loaded
7490         ./extlib2.so: shared object loaded

```

```

7491

```

```

7492 Display the current set of shared object files and their commands:
7493
7494 crash> extend
7495 SHARED OBJECT  COMMANDS
7496 ./extlib1.so   echo util bin
7497 ./extlib2.so   smp show
7498
7499 Unload one of the shared object files:
7500
7501 crash> extend -u extlib1.so
7502 ./extlib1.so: shared object unloaded
7503
7504 Unload all currently-loaded object files:
7505
7506 crash> extend -u
7507 ./extlib2.so: shared object unloaded
7508
7509 CREATING A SHARED OBJECT
7510 The extend command loads shared object files using dlopen(3), which in
7511 turn calls the shared object's constructor function. The shared object's
7512 constructor function should register its command set by calling
7513 register_extension(), passing it a pointer to an array of one or more
7514 structures of the following type:
7515
7516 struct command_table_entry {
7517     char *name;
7518     cmd_func_t func;
7519     char **help_data,
7520     ulong flags;
7521 };
7522
7523 Each command_table_entry structure contains the ASCII name of a command,
7524 the command's function address, a pointer to an array of help data strings,
7525 and a flags field. The help_data field is optional; if it is non-NULL, it
7526 should point to an array of character strings used by the "help"
7527 command, and during command failures. The flags field currently has two
7528 available bit settings, REFRESH_TASK_TABLE, which should be set if it is
7529 preferable to reload the current set of running processes just prior to
7530 executing the command (on a live system) and MINIMAL, which should be
7531 set if the command should be available in minimal mode. Terminate the array
7532 of command_table_entry structures with an entry with a NULL command name.
7533
7534 Below is an example shared object file consisting of just one command,
7535 called "echo", which simply echoes back all arguments passed to it.
7536 Note the comments contained within it for further details. Cut and paste
7537 the following output into a file, and call it, for example, "echo.c".
7538 Then compiled in either of two manners. Either manually like so:
7539
7540 gcc -shared -rdynamic -o echo.so echo.c -fPIC -D<machine-type> $(TARGET_CFLAGS)
7541
7542 where <machine-type> must be one of the MACHINE_TYPE #define's in defs.h,
7543 and where $(TARGET_CFLAGS) is the same as it is declared in the top-level
7544 Makefile after a build is completed. Or alternatively, the "echo.c" file
7545 can be copied into the "extensions" subdirectory, and compiled automatically
7546 like so:
7547
7548 make extensions
7549
7550 The echo.so file may be dynamically linked into crash during runtime, or
7551 during initialization by putting "extend echo.so" into a .crashrc file
7552 located in the current directory, or in the user's $HOME directory.
7553
7554 ----- cut here -----
7555
7556 #include "defs.h"          /* From the crash source top-level directory */
7557
7558 void echo_init(void);      /* constructor function */
7559 void echo_fini(void);      /* destructor function (optional) */
7560
7561 void cmd_echo(void);       /* Declare the commands and their help data. */
7562 char *help_echo[];
7563
7564 static struct command_table_entry command_table[] = {

```



```

7565         { "echo", cmd_echo, help_echo, 0},          /* One or more commands, */
7566         { NULL },                                     /* terminated by NULL, */
7567     };
7568
7569
7570 void __attribute__((constructor))
7571 echo_init(void) /* Register the command set. */
7572 {
7573     register_extension(command_table);
7574 }
7575
7576 /*
7577  * This function is called if the shared object is unloaded.
7578  * If desired, perform any cleanups here.
7579  */
7580 void __attribute__((destructor))
7581 echo_fini(void) { }
7582
7583
7584 /*
7585  * Arguments are passed to the command functions in the global args[argcnt]
7586  * array. See getopt(3) for info on dash arguments. Check out defs.h and
7587  * other crash commands for usage of the myriad of utility routines available
7588  * to accomplish what your task.
7589  */
7590 void
7591 cmd_echo(void)
7592 {
7593     int c;
7594
7595     while ((c = getopt(argcnt, args, "")) != EOF) {
7596         switch(c)
7597         {
7598             default:
7599                 argerrs++;
7600                 break;
7601         }
7602     }
7603
7604     if (argerrs)
7605         cmd_usage(pc->curcmd, SYNOPSIS);
7606
7607     while (args[optind])
7608         fprintf(fp, "%s ", args[optind++]);
7609
7610     fprintf(fp, "\n");
7611 }
7612
7613 /*
7614  * The optional help data is simply an array of strings in a defined format.
7615  * For example, the "help echo" command will use the help_echo[] string
7616  * array below to create a help page that looks like this:
7617  *
7618  * NAME
7619  *     echo - echoes back its arguments
7620  *
7621  * SYNOPSIS
7622  *     echo arg ...
7623  *
7624  * DESCRIPTION
7625  *     This command simply echoes back its arguments.
7626  *
7627  * EXAMPLE
7628  *     Echo back all command arguments:
7629  *
7630  *         crash> echo hello, world
7631  *         hello, world
7632  *
7633  */
7634
7635 char *help_echo[] = {
7636     "echo",          /* command name */
7637     "echoes back its arguments", /* short description */

```

```

7638         "arg ...",                               /* argument synopsis, or " " if none */
7639
7640         " This command simply echoes back its arguments.",
7641         "\nEXAMPLE",
7642         " Echo back all command arguments:\n",
7643         " crash> echo hello, world",
7644         " hello, world",
7645         NULL
7646     };
7647
7648
7649
7650 NAME
7651     mach - machine specific data
7652
7653 SYNOPSIS
7654     mach [-m | -c -[xd] | -o]
7655
7656 DESCRIPTION
7657     This command displays data specific to a machine type.
7658
7659     -m Display the physical memory map (x86, x86_64 and ia64 only).
7660     -c Display each cpu's cpuinfo structure (x86, x86_64 and ia64 only).
7661         Display each cpu's x8664_pda structure (x86_64 only),
7662         Display the hwrpb_struct, and each cpu's percpu_struct (alpha only).
7663     -x override default output format with hexadecimal format.
7664     -d override default output format with decimal format.
7665     -o Display the OPAL console log (ppc64 only).
7666
7667 EXAMPLES
7668     crash> mach
7669             MACHINE TYPE: i686
7670             MEMORY SIZE: 512 MB
7671             CPUS: 2
7672             HYPERVISOR: KVM
7673             PROCESSOR SPEED: 1993 Mhz
7674             HZ: 100
7675             PAGE SIZE: 4096
7676             KERNEL VIRTUAL BASE: c0000000
7677             KERNEL VMALLOC BASE: e0800000
7678             KERNEL STACK SIZE: 8192
7679
7680 Display the system physical memory map:
7681
7682     crash> mach -m
7683             PHYSICAL ADDRESS RANGE          TYPE
7684     0000000000000000 - 000000000000a0000  E820_RAM
7685     000000000000f0000 - 00000000000100000  E820_RESERVED
7686     00000000000100000 - 0000000001ff75000  E820_RAM
7687     0000000001ff75000 - 0000000001ff77000  E820_NVS
7688     0000000001ff77000 - 0000000001ff98000  E820_ACPI
7689     0000000001ff98000 - 00000000020000000  E820_RESERVED
7690     00000000fec00000 - 00000000fec900000  E820_RESERVED
7691     00000000fee00000 - 00000000fee100000  E820_RESERVED
7692     00000000ffb00000 - 00000000100000000  E820_RESERVED
7693
7694 Display the OPAL console log:
7695
7696     crash> mach -o
7697     [ 65.219056911,5] SkiBoot skiboot-5.4.0-218-ge0225cc-df9a248 starting...
7698     [ 65.219065872,5] initial console log level: memory 7, driver 5
7699     [ 65.219068917,6] CPU: P8 generation processor(max 8 threads/core)
7700     [ 65.219071681,7] CPU: Boot CPU PIR is 0x0060 PVR is 0x004d0200
7701     [ 65.219074685,7] CPU: Initial max PIR set to 0x1fff
7702     [ 65.219607955,5] FDT: Parsing fdt @0xff000000
7703     [ 494.026291523,7] BT: seq 0x25 netfn 0x0a cmd 0x48: Message sent to host
7704     [ 494.027636927,7] BT: seq 0x25 netfn 0x0a cmd 0x48: IPMI MSG done
7705
7706
7707 NAME
7708     runq - run queue
7709
7710 SYNOPSIS

```

```

7711     runq [-t] [-T] [-m] [-g] [-c cpu(s)]
7712
7713 DESCRIPTION
7714     With no argument, this command displays the tasks on the run queues
7715     of each cpu.
7716
7717     -t   Display the timestamp information of each cpu's runqueue, which is the
7718           rq.clock, rq.most_recent_timestamp or rq.timestamp_last_tick value,
7719           whichever applies; following each cpu timestamp is the last_run or
7720           timestamp value of the active task on that cpu, whichever applies,
7721           along with the task identification.
7722     -T   Display the time lag of each CPU relative to the most recent runqueue
7723           timestamp.
7724     -m   Display the amount of time that the active task on each cpu has been
7725           running, expressed in a format consisting of days, hours, minutes,
7726           seconds and milliseconds.
7727     -g   Display tasks hierarchically by task_group. The task_group line shows
7728           the task_group address, the cfs_rq or rt_rq address, the task_group
7729           name (if any), and whether the task_group is throttled.
7730     -c cpu restrict the output to the run queue data of one or more CPUs,
7731           which can be specified using the format "3", "1,8,9", "1-23",
7732           or "1,8,9-14".
7733

```

EXAMPLES

Display the tasks on an O(1) scheduler run queue:

```

7737 crash> runq
7738 CPU 0 RUNQUEUE: ffff880001cdb460
7739     CURRENT: PID: 2739   TASK: ffff8800320fa7e0  COMMAND: "bash"
7740     ACTIVE PRIO_ARRAY: ffff880001cdb4d8
7741         [115] PID: 2739   TASK: ffff8800320fa7e0  COMMAND: "bash"
7742             PID: 1776   TASK: ffff88003217d820  COMMAND: "syslogd"
7743     EXPIRED PRIO_ARRAY: ffff880001cddb8
7744         [no tasks queued]
7745
7746 CPU 1 RUNQUEUE: ffff880001ce3460
7747     CURRENT: PID: 1779   TASK: ffff88003207a860  COMMAND: "klogd"
7748     ACTIVE PRIO_ARRAY: ffff880001ce34d8
7749         [115] PID: 1779   TASK: ffff88003207a860  COMMAND: "klogd"
7750     EXPIRED PRIO_ARRAY: ffff880001ce3db8
7751         [no tasks queued]

```

Display the tasks on a CFS run queue:

```

7755 crash> runq
7756 CPU 0 RUNQUEUE: ffff8800090436c0
7757     CURRENT: PID: 588    TASK: ffff88007e4877a0  COMMAND: "udev"
7758     RT PRIO_ARRAY: ffff8800090437c8
7759         [no tasks queued]
7760     CFS RB_ROOT: ffff880009043740
7761         [118] PID: 2110   TASK: ffff88007d470860  COMMAND: "check-cdrom.sh"
7762         [118] PID: 2109   TASK: ffff88007f1247a0  COMMAND: "check-cdrom.sh"
7763         [118] PID: 2114   TASK: ffff88007f20e080  COMMAND: "udev"
7764
7765 CPU 1 RUNQUEUE: ffff88000905b6c0
7766     CURRENT: PID: 2113   TASK: ffff88007e8ac140  COMMAND: "udev"
7767     RT PRIO_ARRAY: ffff88000905b7c8
7768         [no tasks queued]
7769     CFS RB_ROOT: ffff88000905b740
7770         [118] PID: 2092   TASK: ffff88007d7a4760  COMMAND: "MAKEDEV"
7771         [118] PID: 1983   TASK: ffff88007e59f140  COMMAND: "udev"
7772         [118] PID: 2064   TASK: ffff88007e40f7a0  COMMAND: "udev"
7773         [115] PID: 2111   TASK: ffff88007e4278a0  COMMAND: "kthread"

```

Display run queue timestamp data:

```

7777 crash> runq -t
7778 CPU 0: 2680990637359
7779     2680986653330  PID: 28228  TASK: ffff880037ca2ac0  COMMAND: "loop"
7780 CPU 1: 2680940618478
7781     2680940618478  PID: 28167  TASK: ffff880078130040  COMMAND: "bash"
7782 CPU 2: 2680990763425
7783     2680986785772  PID: 28227  TASK: ffff8800787780c0  COMMAND: "loop"

```

```

7784 CPU 3: 2680990954469
7785 2680986059540 PID: 28226 TASK: ffff880078778b00 COMMAND: "loop"
7786
7787 Display the amount of time the active task on each cpu has been running:
7788
7789 crash> runq -m
7790 CPU 0: [0 00:00:00.014] PID: 5275 TASK: f5dbcaa0 COMMAND: "sh"
7791 CPU 1: [0 00:00:00.002] PID: 5203 TASK: f5c7baa0 COMMAND: "cat"
7792 CPU 2: [0 00:00:00.014] PID: 7971 TASK: f5c6c550 COMMAND: "khelper"
7793 CPU 3: [0 00:00:00.002] PID: 0 TASK: f4ccd000 COMMAND: "swapper"
7794
7795 Display tasks hierarchically by task_group:
7796
7797 crash> runq -g
7798 CPU 0
7799 CURRENT: PID: 14734 TASK: ffff88010626f500 COMMAND: "sh"
7800 ROOT_TASK_GROUP: ffffffff81ed93e0 RT_RQ: ffff880028216808
7801 [ 0] TASK_GROUP: ffff88022c6bbc00 RT_RQ: ffff880139fc9800 (THROTTLED)
7802 [ 0] PID: 14750 TASK: ffff88013a4dd540 COMMAND: "rtloop99"
7803 [ 1] PID: 14748 TASK: ffff88013bbca040 COMMAND: "rtloop98"
7804 [ 1] TASK_GROUP: ffff88012b0fb400 RT_RQ: ffff880089029000
7805 [ 1] PID: 14752 TASK: ffff880088abf500 COMMAND: "rtloop98"
7806 [ 54] PID: 14749 TASK: ffff880037a4e080 COMMAND: "rtloop45"
7807 [ 98] PID: 14746 TASK: ffff88012678c080 COMMAND: "rtloop1"
7808 ROOT_TASK_GROUP: ffffffff81ed93e0 CFS_RQ: ffff8800282166e8
7809 [120] PID: 14740 TASK: ffff88013b1e6080 COMMAND: "sh"
7810 [120] PID: 14738 TASK: ffff88012678d540 COMMAND: "sh"
7811 [120] PID: 14734 TASK: ffff88010626f500 COMMAND: "sh" [CURRENT]
7812 TASK_GROUP: ffff884052bc9800 CFS_RQ: ffff8831e4alb000 (THROTTLED)
7813 [120] PID: 14732 TASK: ffff88013bbcb500 COMMAND: "sh"
7814 [120] PID: 14728 TASK: ffff8800b3496080 COMMAND: "sh"
7815 [120] PID: 14730 TASK: ffff880037833540 COMMAND: "sh"
7816 TASK_GROUP: ffff884058f1d000 CFS_RQ: ffff88120a101600 (THROTTLED)
7817 [120] PID: 14726 TASK: ffff880138d42aa0 COMMAND: "sh"
7818 ...
7819
7820 Display tasks hierarchically by task_group for cpu 3 only:
7821
7822 crash> runq -g -c3
7823 CPU 3
7824 CURRENT: PID: 2948 TASK: ffff88022af2a100 COMMAND: "bash"
7825 INIT_TASK_GROUP: ffffffff81e1a780 RT_RQ: ffff880028216148
7826 [no tasks queued]
7827 INIT_TASK_GROUP: ffffffff81e1a780 CFS_RQ: ffff880028216028
7828 [120] PID: 2948 TASK: ffff88022af2a100 COMMAND: "bash" [CURRENT]
7829 TASK_GROUP: ffff88012b880800 CFS_RQ: ffff88012c5d1000 <libvirt>
7830 TASK_GROUP: ffff88012c078000 CFS_RQ: ffff88012c663e00 <qemu>
7831 TASK_GROUP: ffff88022c7f4c00 CFS_RQ: ffff88012bb56000 <guest2>
7832 TASK_GROUP: ffff88022b621400 CFS_RQ: ffff88012b012000 <vcpu0>
7833 [120] PID: 3248 TASK: ffff88012a9d4100 COMMAND: "qemu-kvm"
7834
7835
7836 NAME
7837 trace - show or dump the tracing info
7838
7839 SYNOPSIS
7840 trace [ <show [-c <cpulist>] [-f [no]<flagname>]> ] | <dump [-sm] <dest-dir>> ] |
7841 <dump -t <trace.dat> ]
7842
7843 DESCRIPTION
7844 trace
7845 shows the current tracer and other informations.
7846
7847 trace show
7848 shows all events with readability text(sorted by timestamp)
7849
7850 trace report
7851 the same as "trace show"
7852
7853 trace dump [-sm] <dest-dir>
7854 dump ring_buffers to dest-dir. Then you can parse it
7855 by other tracing tools. The dirs and files are generated
the same as debugfs/tracing.

```

```
7856     -m: also dump metadata of ftrace.  
7857     -s: also dump symbols of the kernel.  
7858 trace dump -t [output-file-name]  
7859     dump ring_buffers and all meta data to a file that can  
7860     be parsed by trace-cmd. Default output file name is "trace.dat".  
7861  
7862
```